

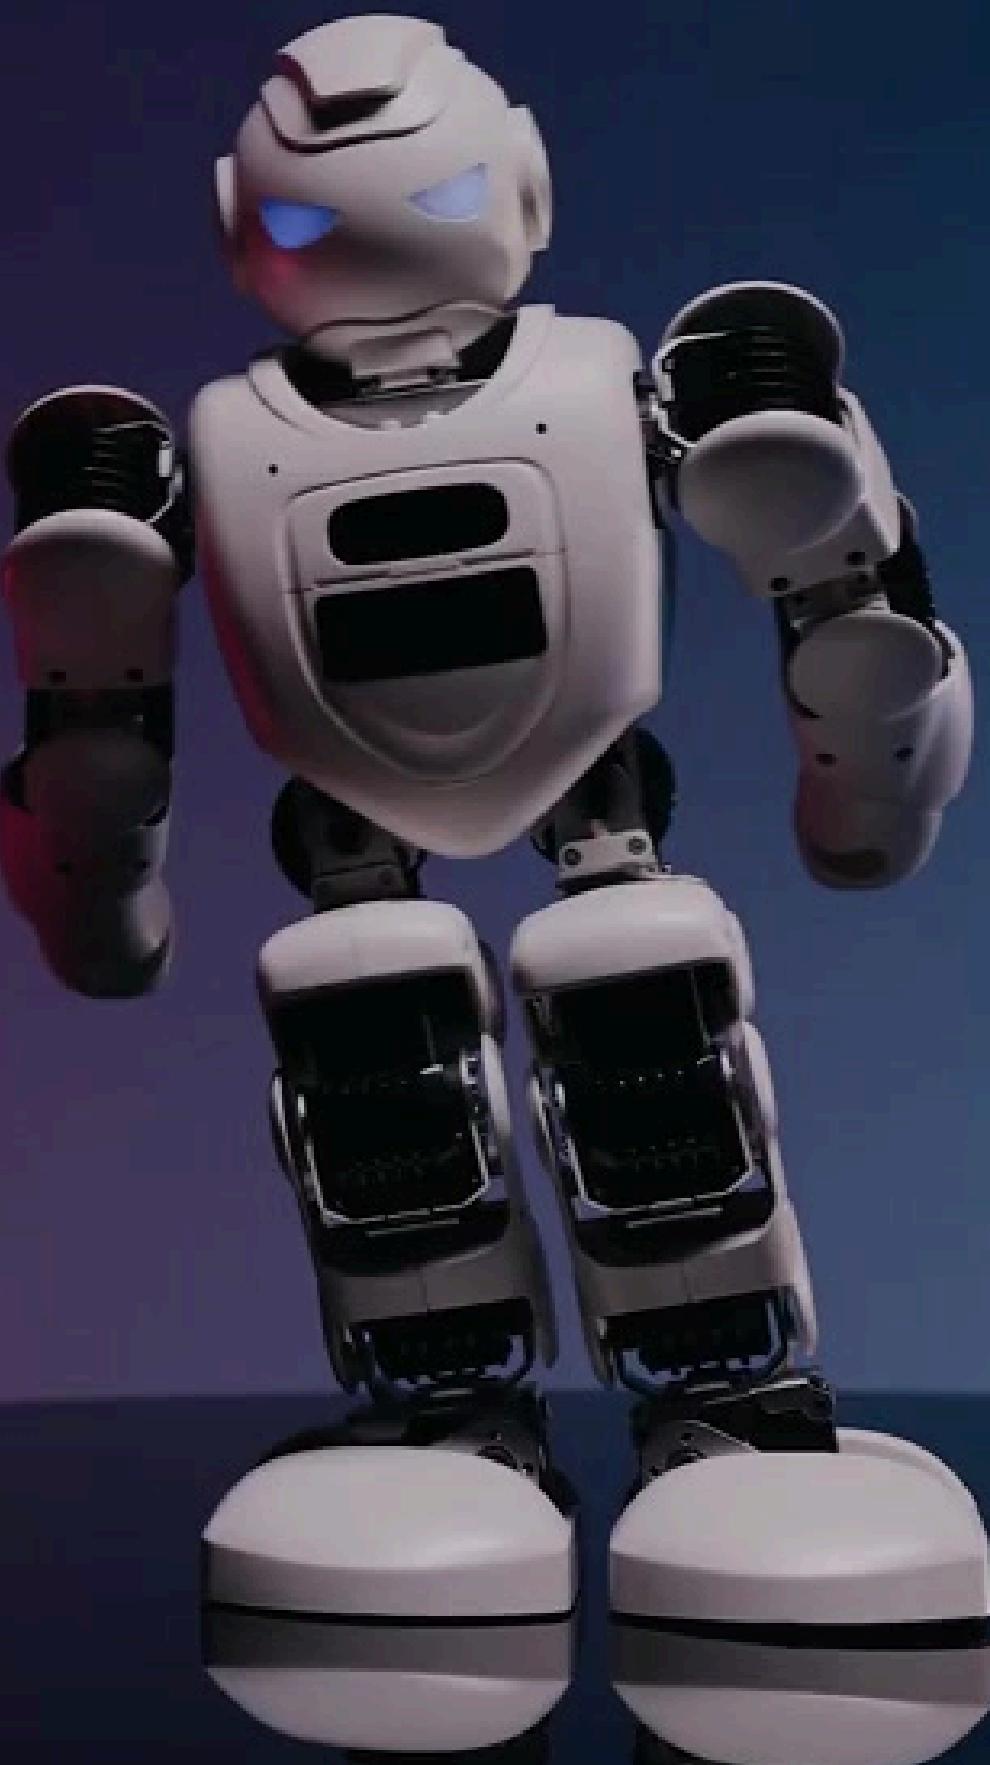


WEEK 1

DEEP LEARNING FOR COMPUTER VISION

UNLIMITED

Presented by **Asst. Prof. Dr. Tuchsanai Ploysuwan**



github.com/Tuchsanai/DL-FOR-COMPUTER-VISION

Tuchsanai / DL-FOR-COMPUTER-VISION

Type to search

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

DL-FOR-COMPUTER-VISION Public

Unpin Unwatch 1 Fork 9 Star 4

main Go to file + <> Code

Tuchsanai 00 d346e79 · 2 hours ago 259 Commits

| File | Commit | Time |
|----------------|------------------|---------------|
| week01 | 00 | 2 hours ago |
| week02 | zz | last year |
| week03_1 | z | last year |
| week03_2 | d | last year |
| week04 | g | 10 months ago |
| week05 | y | 9 months ago |
| week06 | aa | 9 months ago |
| week07 | zzz | 9 months ago |
| week08 | ss | 9 months ago |
| week09 | l | 8 months ago |
| week10 | dd | 9 months ago |
| week11 | d | 9 months ago |
| .gitattributes | Initial commit | 2 years ago |
| .gitignore | 11 | 7 months ago |
| README.md | Update README.md | last year |

README

DL FOR COMPUTER VISION

Document



line



Course Learning Outcomes (CLOs)

Upon completion of this course, students are expected to be able to

| CLO ID | CLO Description |
|--------|---|
| CLO-1 | ผู้เรียนเข้าใจหลักการพื้นฐานของงาน Deep learning สำหรับงาน Computer Vision |
| CLO-2 | ผู้เรียนเข้าใจหลักการพื้นฐานในการใช้ Convolutional neural network (CNN) สำหรับงาน Image Classifications |
| CLO-3 | ผู้เรียนสามารถอินพุตและเลือกใช้ Modern deep learning architectures ที่เหมาะสมสำหรับงาน Computer Vision ชนิดต่างๆ เช่น VGG, ResNet, and Efficientnet |
| CLO-4 | ผู้เรียนเข้าใจหลักการการเลือกใช้งาน Loss function, activation functions และ stochastic optimization ได้อย่างเหมาะสม |
| CLO-5 | ผู้เรียนเข้าใจหลักการในการใช้ Transfer learning และการ fine tuning ที่เหมาะสม |
| CLO-6 | ผู้เรียนเข้าใจหลักการพื้นฐานในการใช้ Deep learning สำหรับการทำ Image Segmentation ในงาน Computer Vision |
| CLO-7 | ผู้เรียนสามารถเข้าใจหลักการของ Object detection and Object Tracking ในงาน Computer Vision |
| CLO-8 | ผู้เรียนสามารถเข้าใจหลักการของ Generative Adversarial Network |

| ลำดับที่ | หัวข้อ | วิธีการเรียน | ผลการเรียนรู้ | หมายเหตุ |
|----------|--|--------------|---------------------|-----------|
| 1 | The basic concept of deep learning and their applications in computer vision | บรรยาย | CLO-1 | Chapter 1 |
| 2 | Introduction to Images and Videos Basics with Python and OpenCV | บรรยาย | CLO-1 | Chapter 1 |
| 3 | Advance in OpenCV | บรรยาย | CLO-1 | Chapter 1 |
| 4 | Introduction to Pytorch | บรรยาย | CLO-2, CLO-3, CLO-4 | Chapter 2 |
| 5 | Fundamental Neural network | บรรยาย | CLO-2 | Chapter 2 |
| 6 | Stochastic optimization methods | บรรยาย | CLO-4 | Chapter 2 |
| 7 | Convolutional neural network for image classification | บรรยาย | CLO-2, CLO-3 | Chapter 3 |
| 9 | Modern architectures such as VGG, ResNet, and Efficientnet | บรรยาย | CLO-3, | Chapter 3 |
| 10 | Transfer learning and fine tuning | บรรยาย | CLO-5 | Chapter 3 |
| 11 | Deep learning for Segmentation for images | บรรยาย | CLO-6 | Chapter 4 |
| 12 | Deep learning Object detection and Object Tracking | บรรยาย | CLO-7 | Chapter 4 |
| 13 | GAN architecture (Generative Adversarial Network) for photo-realistic images | บรรยาย | CLO-8 | Chapter 5 |
| 14 | Generative image from text CLIP model | บรรยาย | CLO-8 | Chapter 5 |
| 15 | Discussions | บรรยาย | | |

| การให้คะแนน | |
|-------------------|----|
| Final Examination | 50 |
| HOMEWORK, Kaggle | 30 |
| Participation | 20 |

APPROACHING (ALMOST) ANY MACHINE LEARNING PROBLEM



ABHISHEK THAKUR

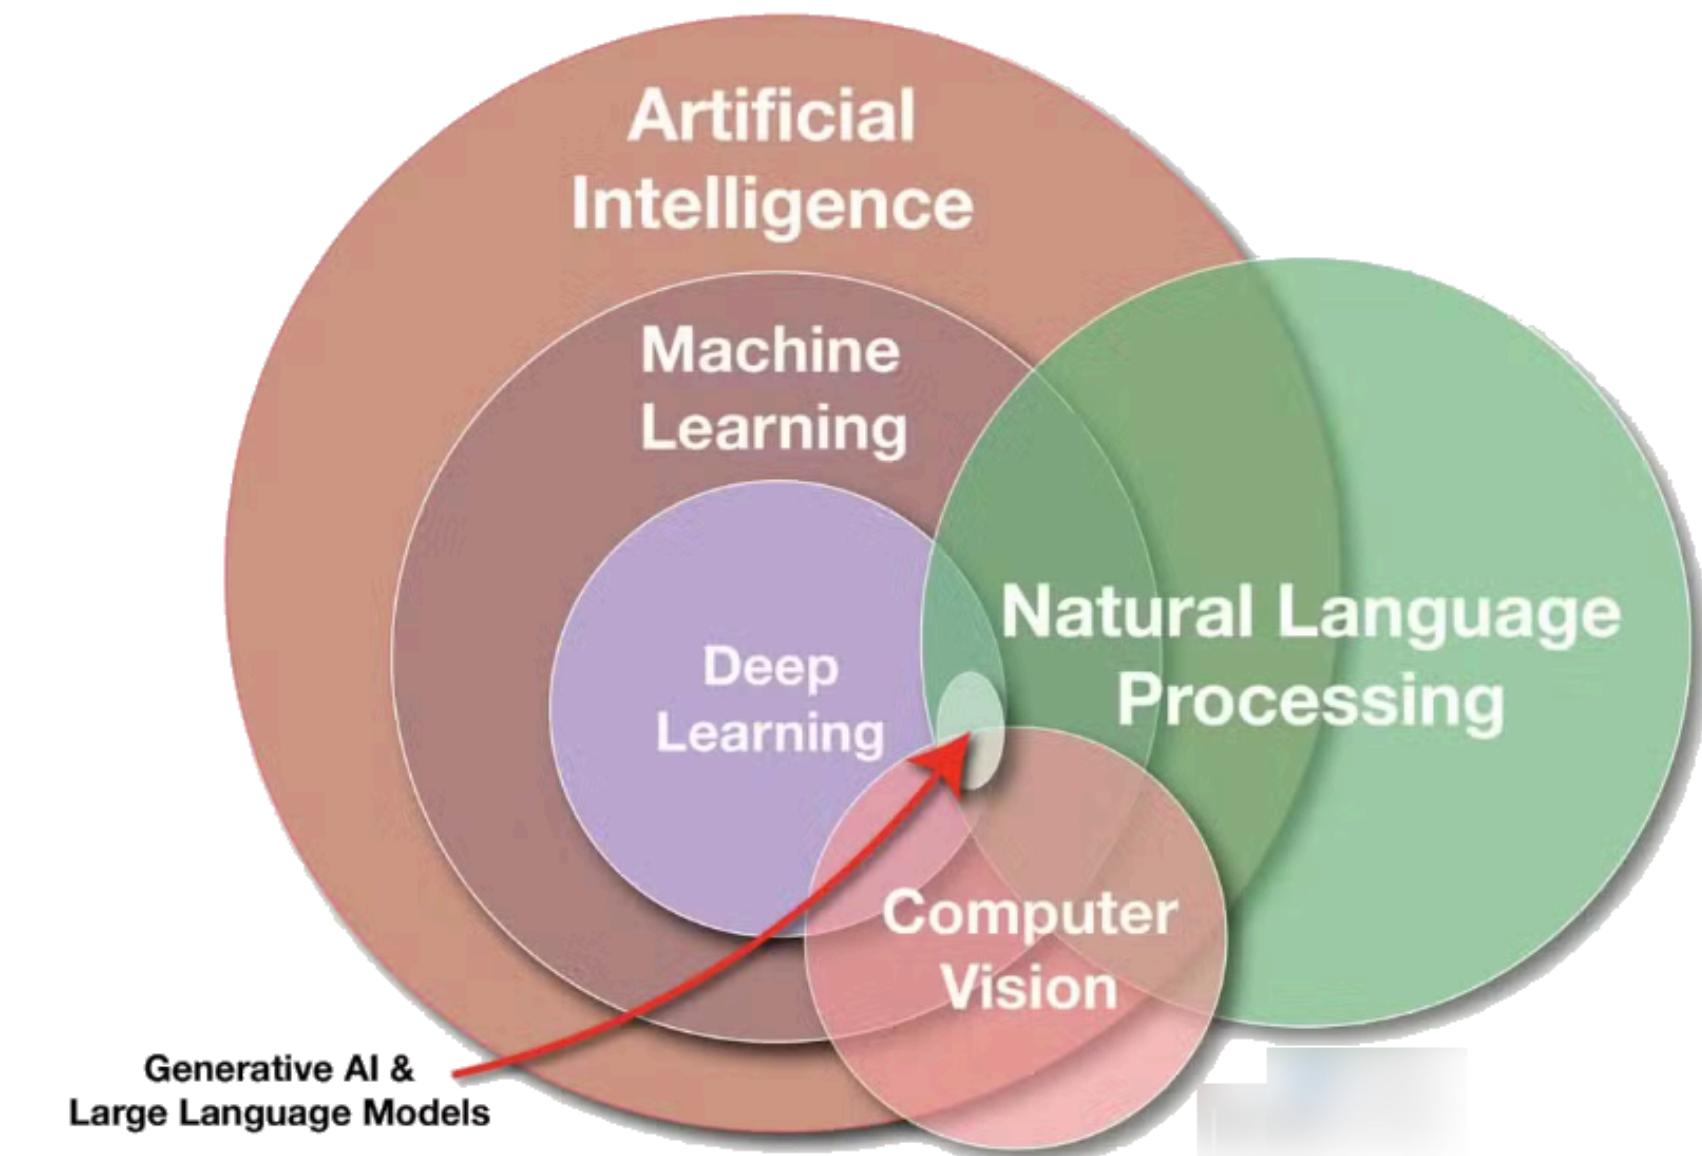
Books

<https://github.com/abhishekrthakur/approachingalmost>

Is Computer Vision Artificial Intelligence?

- Relation between Artificial Intelligence, Machine Learning and Deep Learning, Computer Vision.

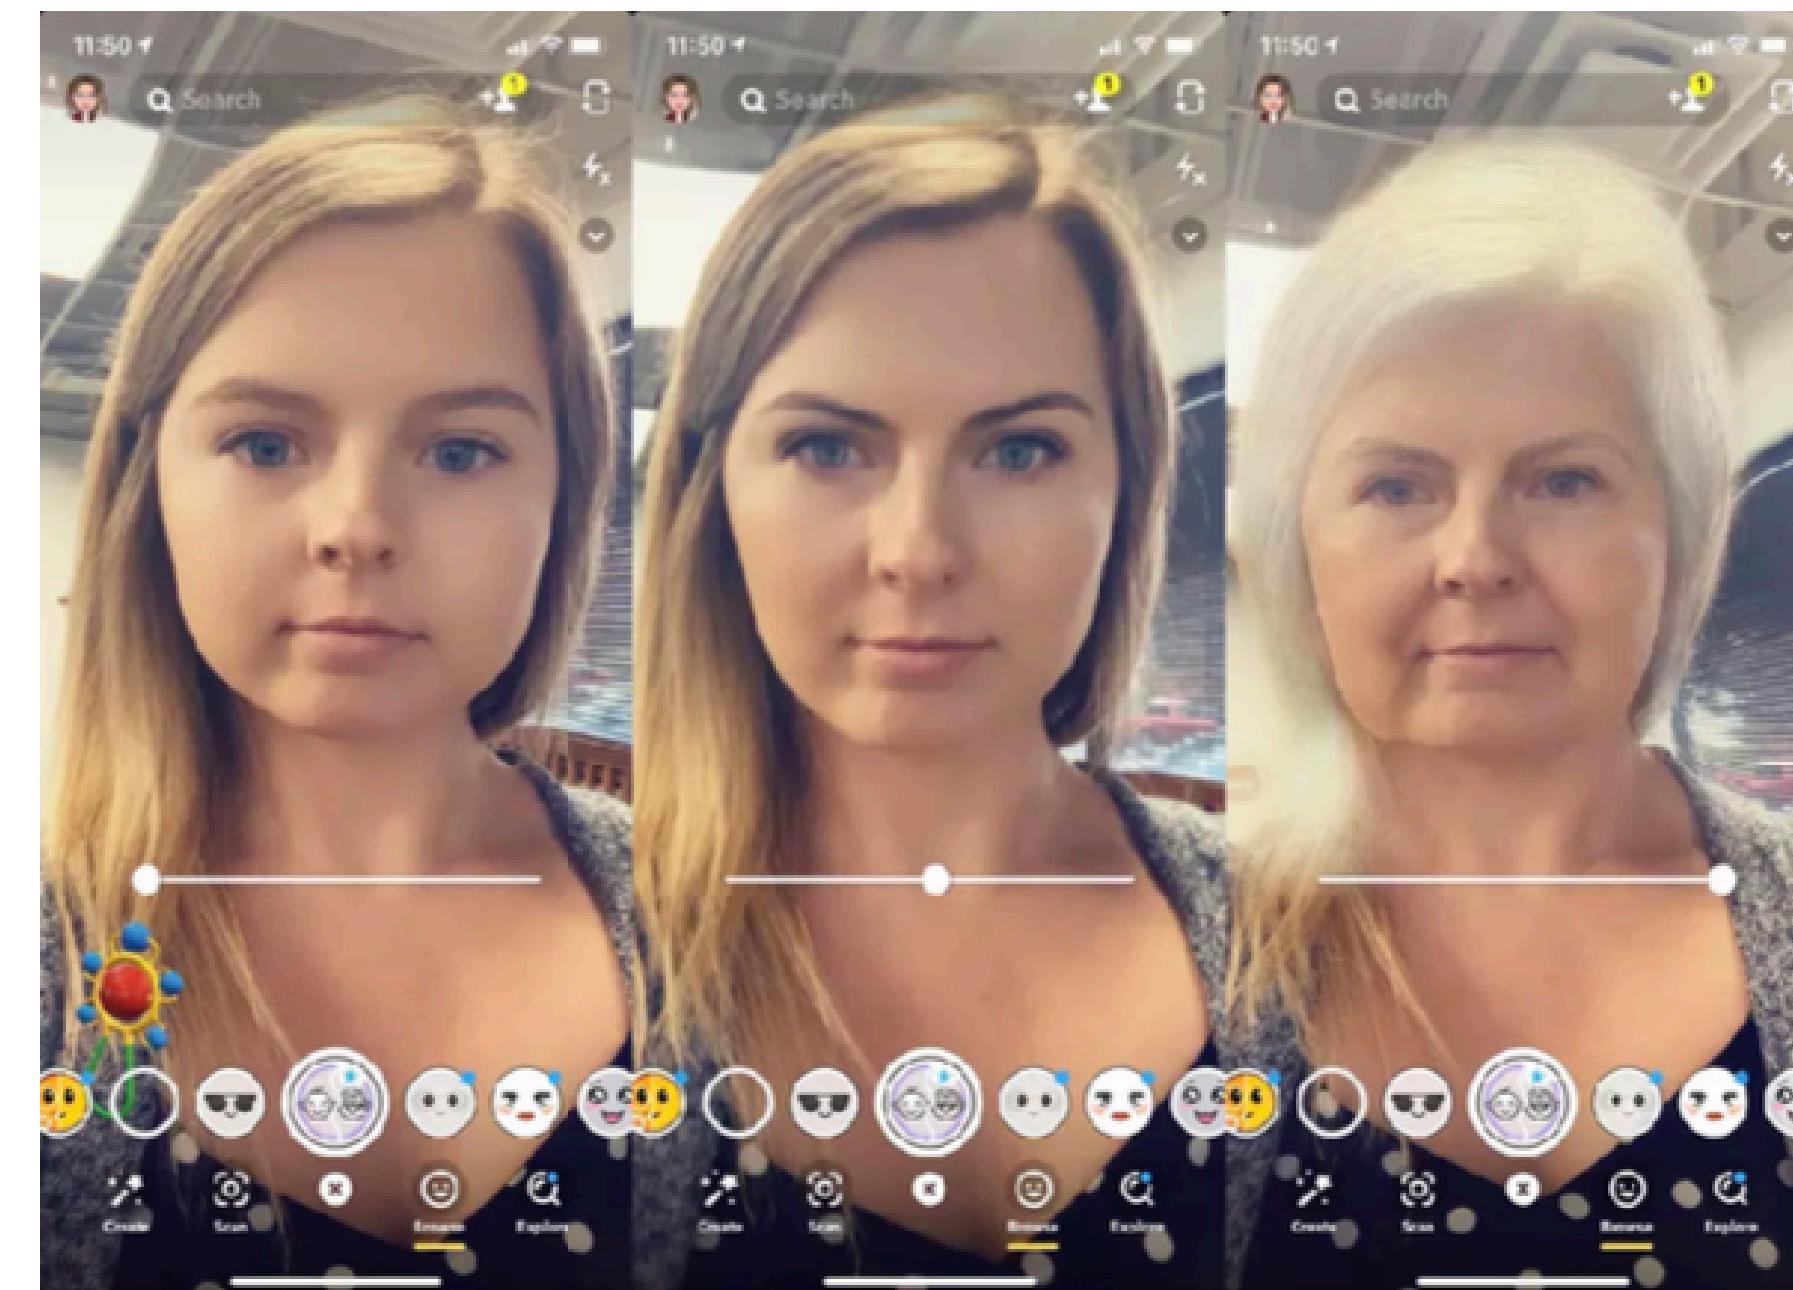
https://www.researchgate.net/figure/Relation-between-Artificial-Intelligence-Machine-Learning-and-Deep-Learning-Computer_fig1_342978934



So what can Computer Vision do?

You might be familiar with these...

- **Snapchat and Instagram filters**
- Optical Character Recognition (OCR)
- Licence Plate Reading
- Self-driving cars
- Sporting Analysis
- Facial Recognition

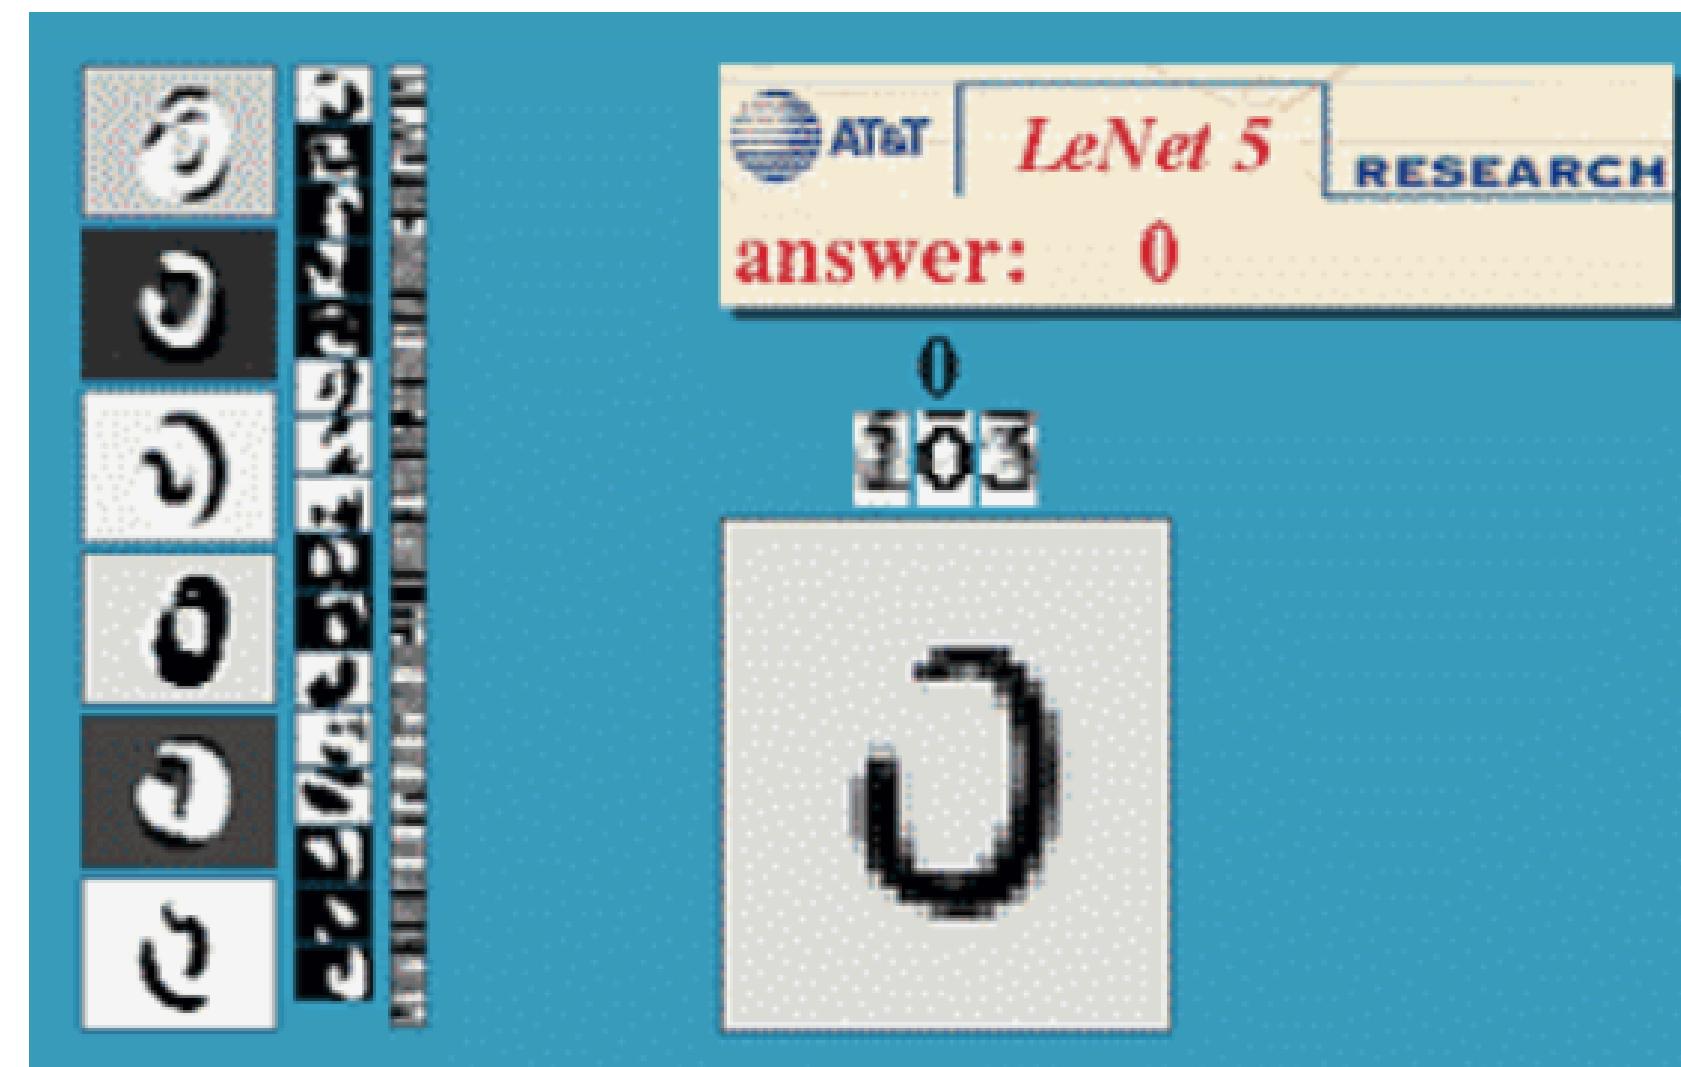


Source - Cnet - Snapchat's Time Machine AR lens creepily shows what you'll look like old

So what can Computer Vision do?

You might be familiar with these...

- Snapchat and Instagram filters
- **Optical Character Recognition (OCR)**
- Licence Plate Reading
- Self-driving cars
- Sporting Analysis
- Facial Recognition



Source -AT&T's LeNet OCR for Handwritten Digits

So what can Computer Vision do?

- Snapchat and Instagram filters
- Optical Character Recognition (OCR)
- **Licence Plate Reading**
- Self-driving cars
- Sporting Analysis
- Facial Recognition

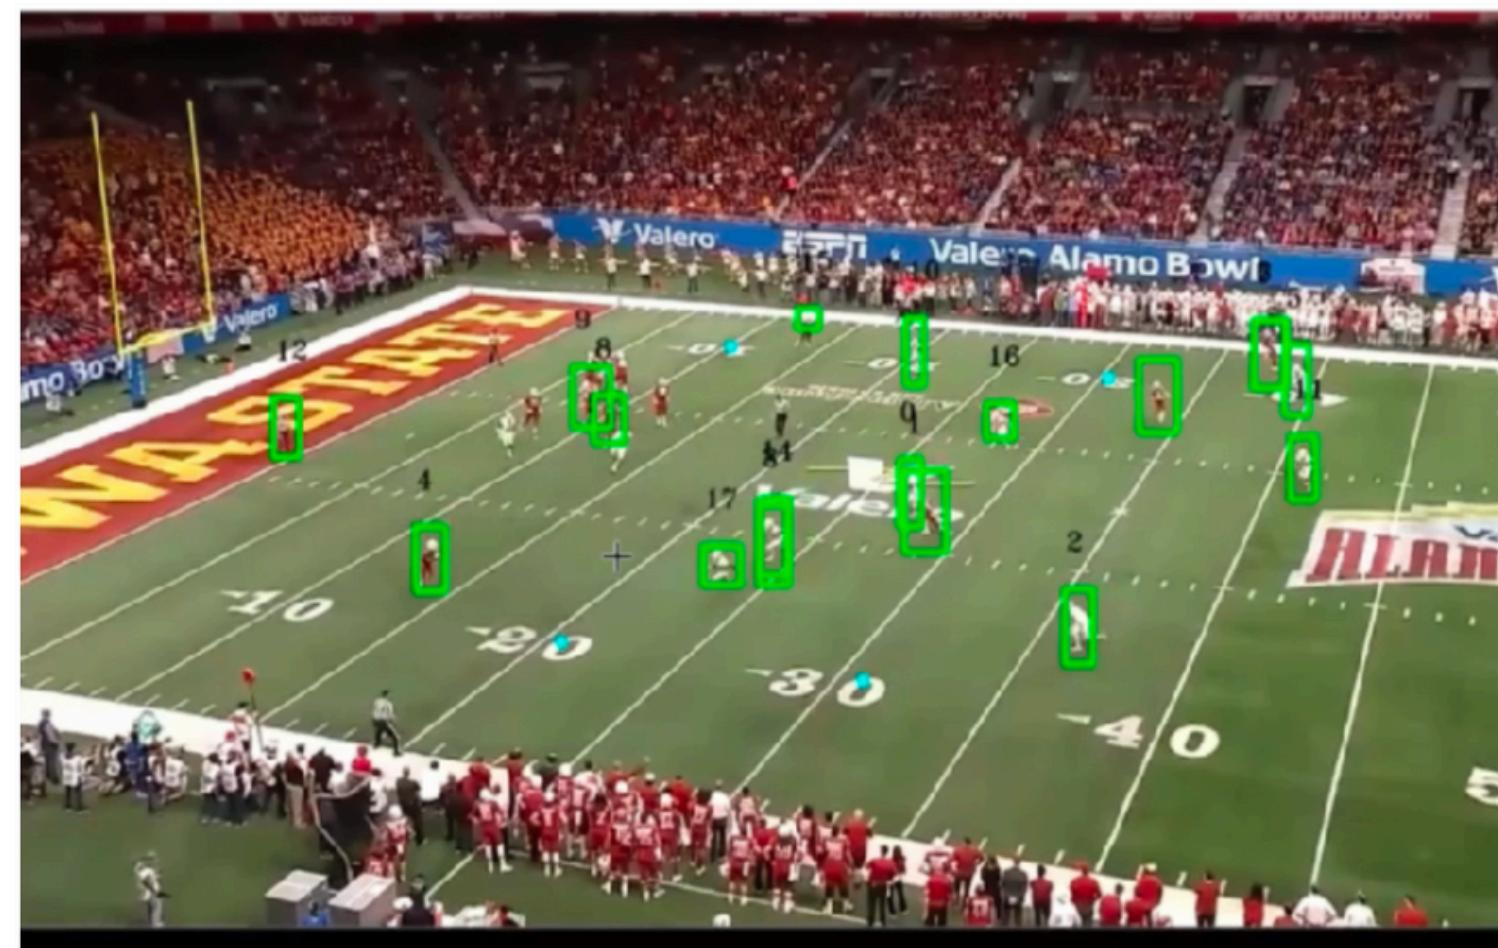
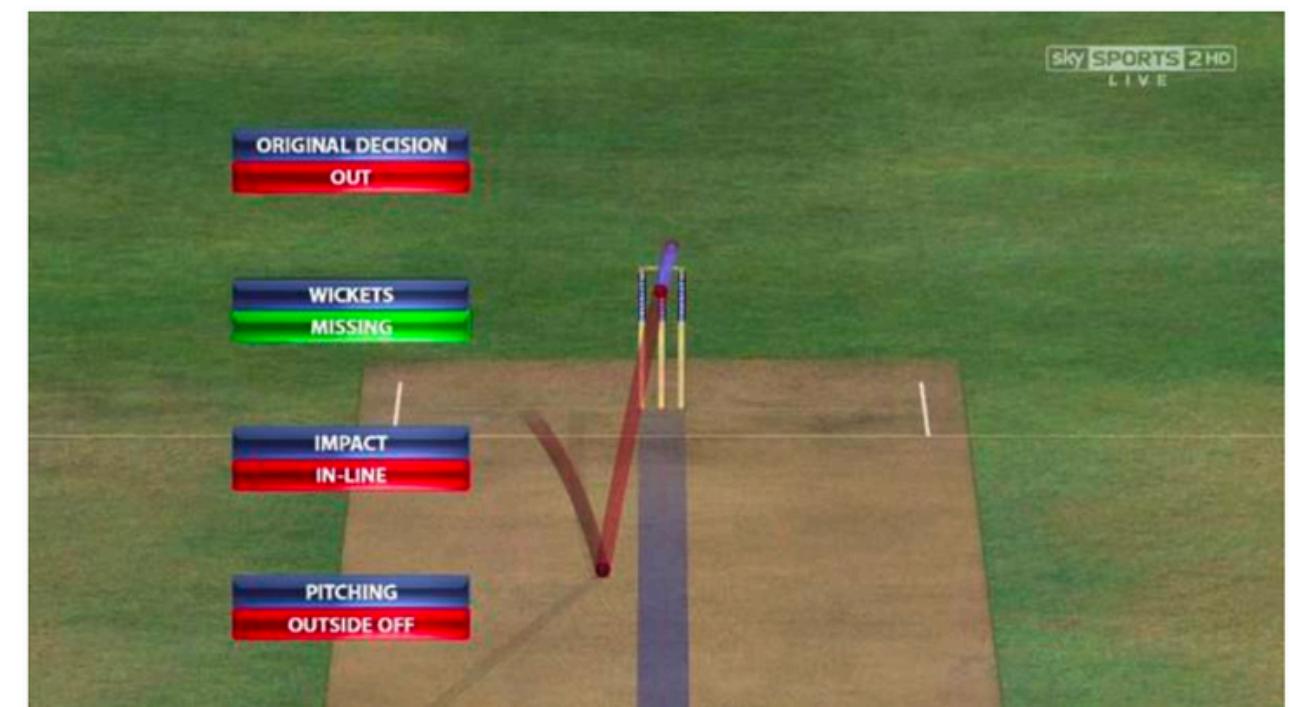


So what can Computer Vision do?

You might be familiar with these...

- Snapchat and Instagram filters
- Optical Character Recognition (OCR)
- Licence Plate Reading
- Self-driving cars
- **Sporting Analysis**
- Facial Recognition

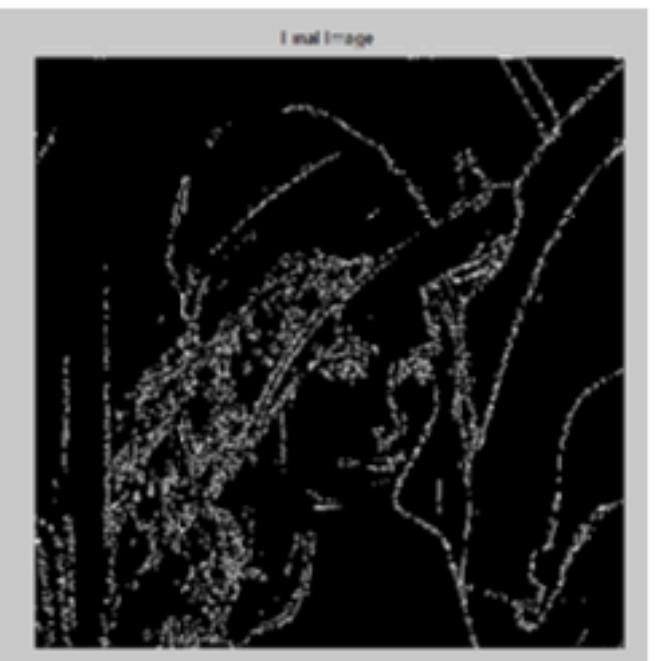
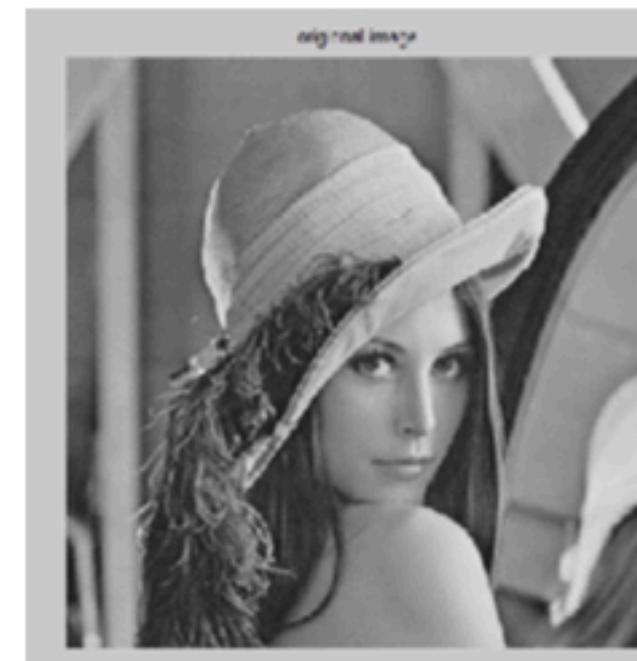
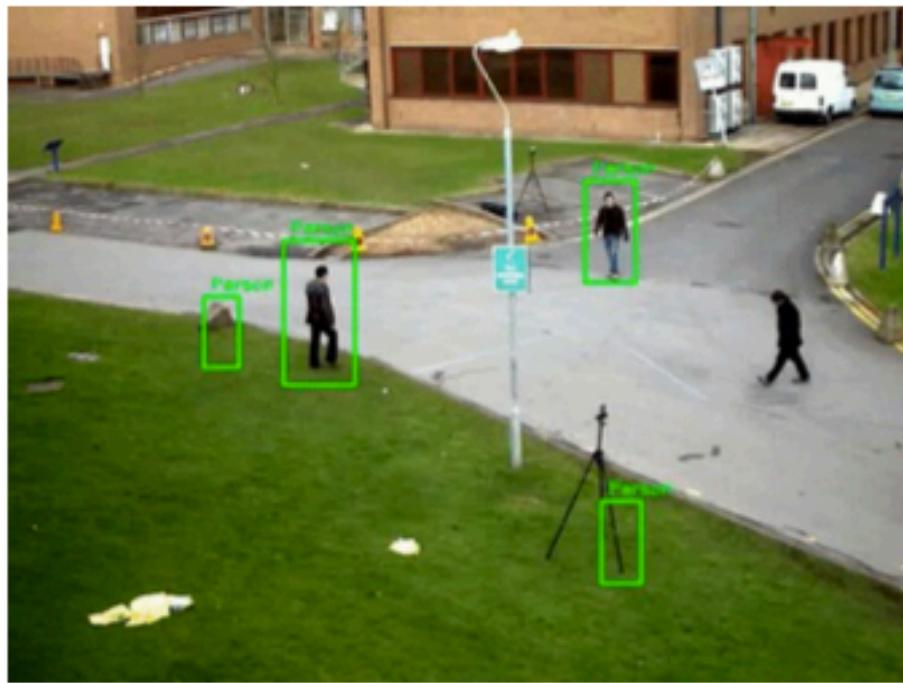
HawkEye in Cricket



Source -Roboflow - AI Coach

Classical Computer Vision?

- What is meant by **Classical Computer Vision?**
- It encompasses Computer Vision algorithms that **do not** involve Machine Learning
- Before the advent of Machine Learning and Deep Learning, Computer Vision was a deeply explored field and many useful algorithms were developed for things like **feature extraction, OCR, Segmentation and simple transformations.**
- **OpenCV** is the Classical Computer Vision library of choice!



Deep Learning Computer Vision

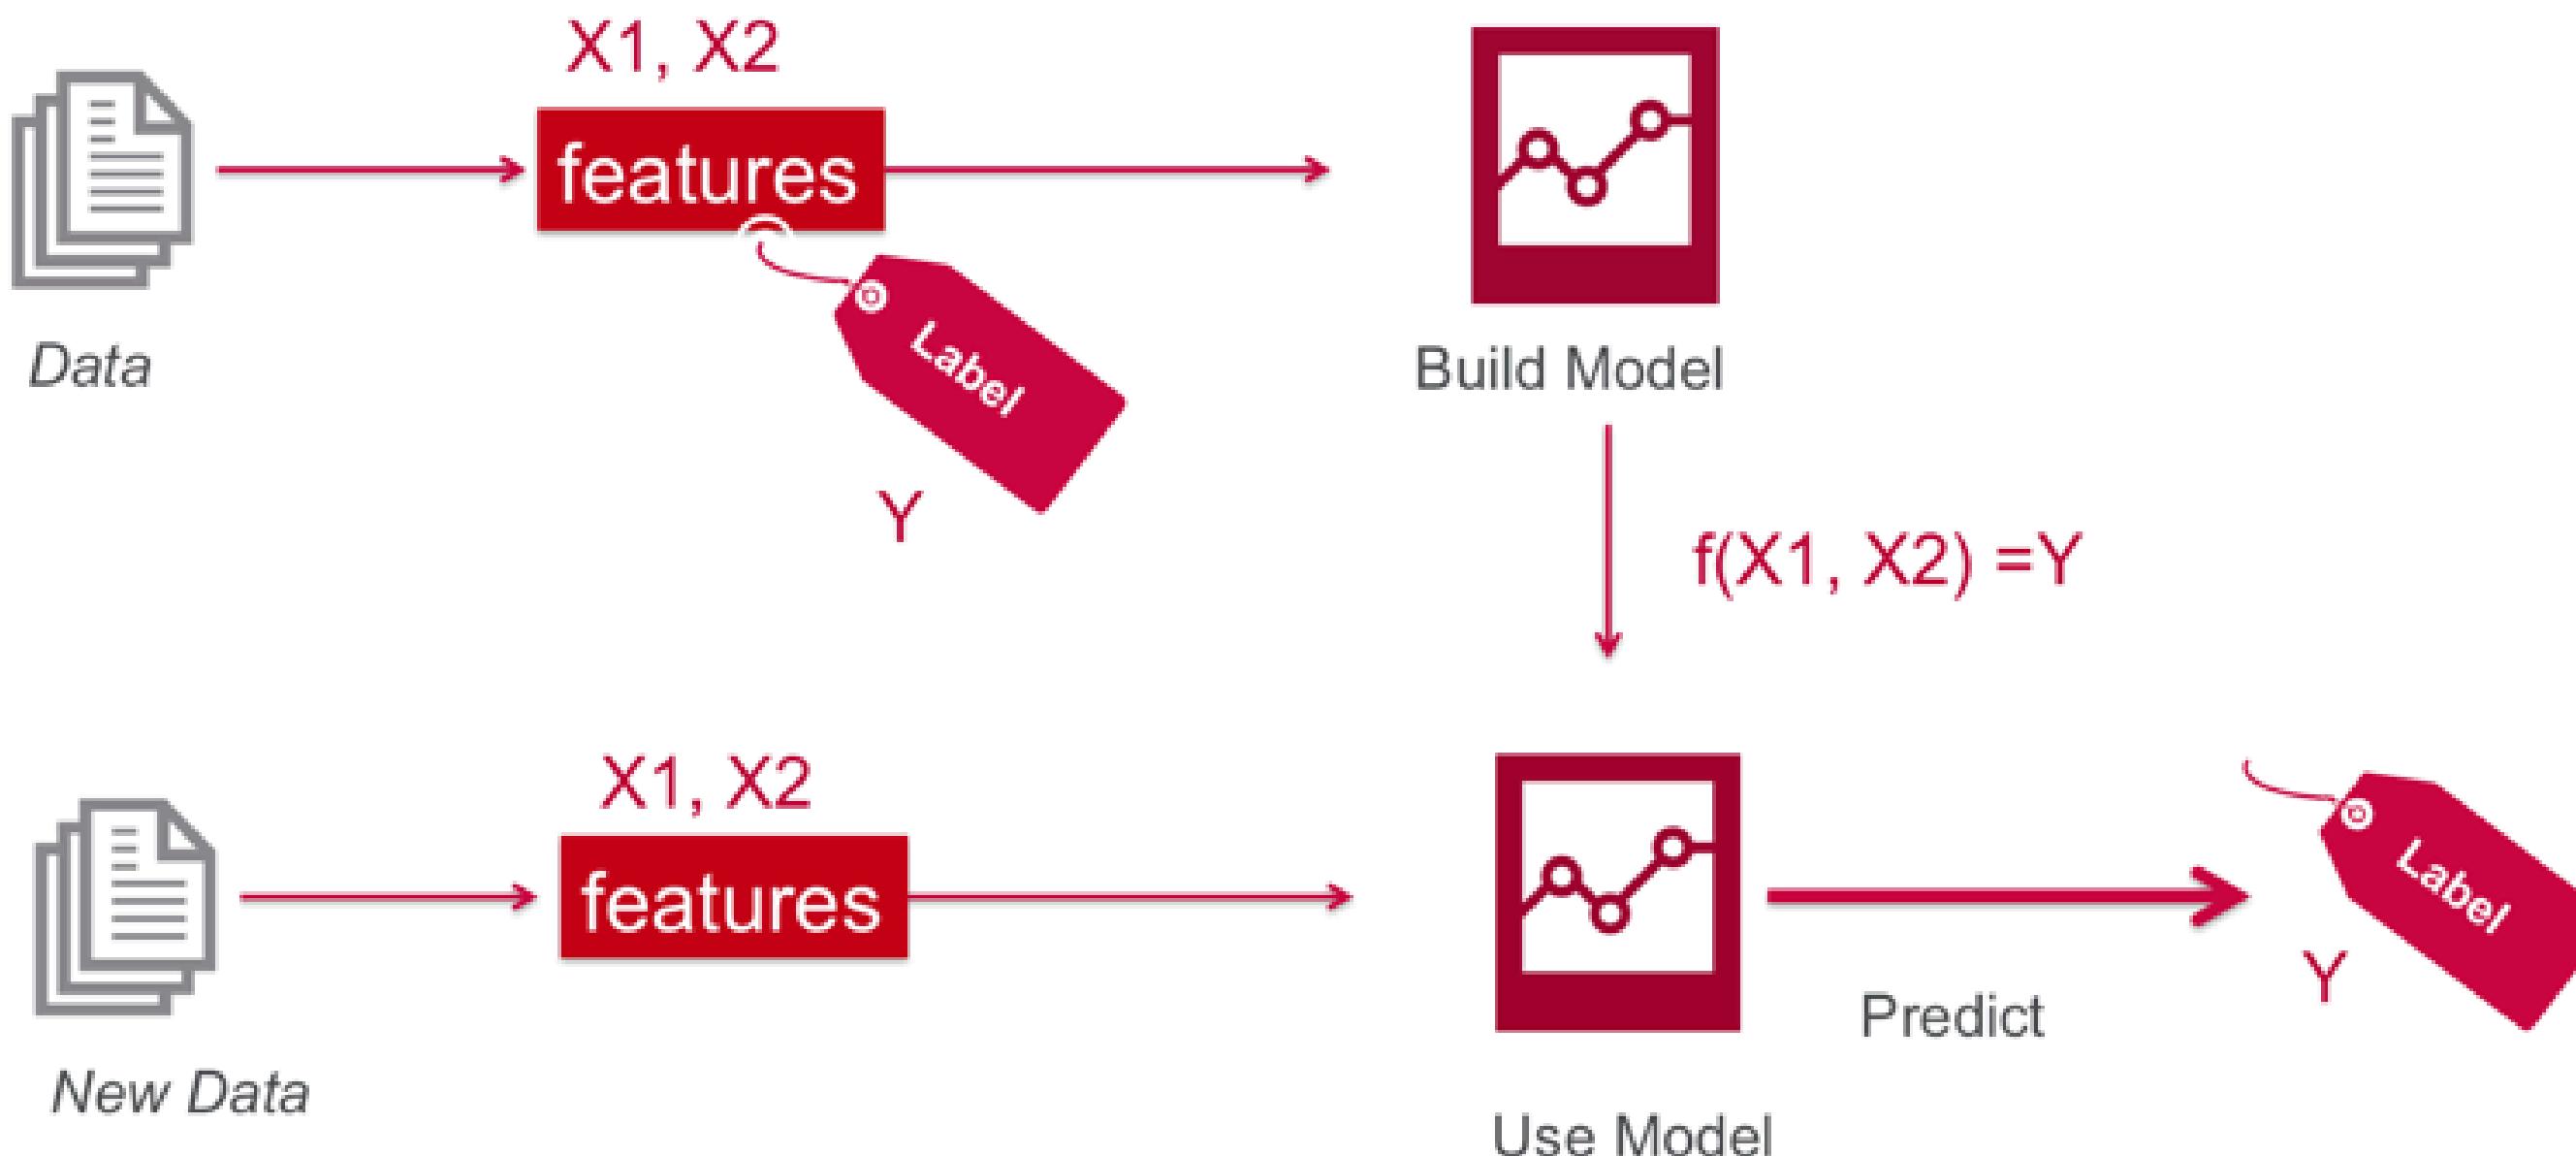
- Deep Learning was used in Computer Vision since the 1990s, however due to the computational requirements and intricate design, it remained on the sidelines for decades.
- Until the mid 2010s...which brought **two important building blocks** together.
 - **Accessible GPU processing** (NVIDIA's CUDA)



Deep Learning CV vs Classical CV

| Deep Learning | Classical Computer Vision |
|---|---|
| Adapts to new images well (assuming it's similar to the data it was trained on) | Small changes can have big negative impacts |
| Requires Models to be trained | Doesn't require training and can be used once coded |
| Model weights learn to adapt to varying image conditions | Relies on hardcode features and parameters |
| Requires GPU hardware (most times) | Can be run on CPU |

Supervised Learning



Classification v.s. Regression

Classification

- Predict which category an item belongs to based on labeled examples of known items
- Use in classify a new data

Regression

- Predict a continuous value target.
- Linear regression predicts a numeric value
- Logistic regression predicts a probability

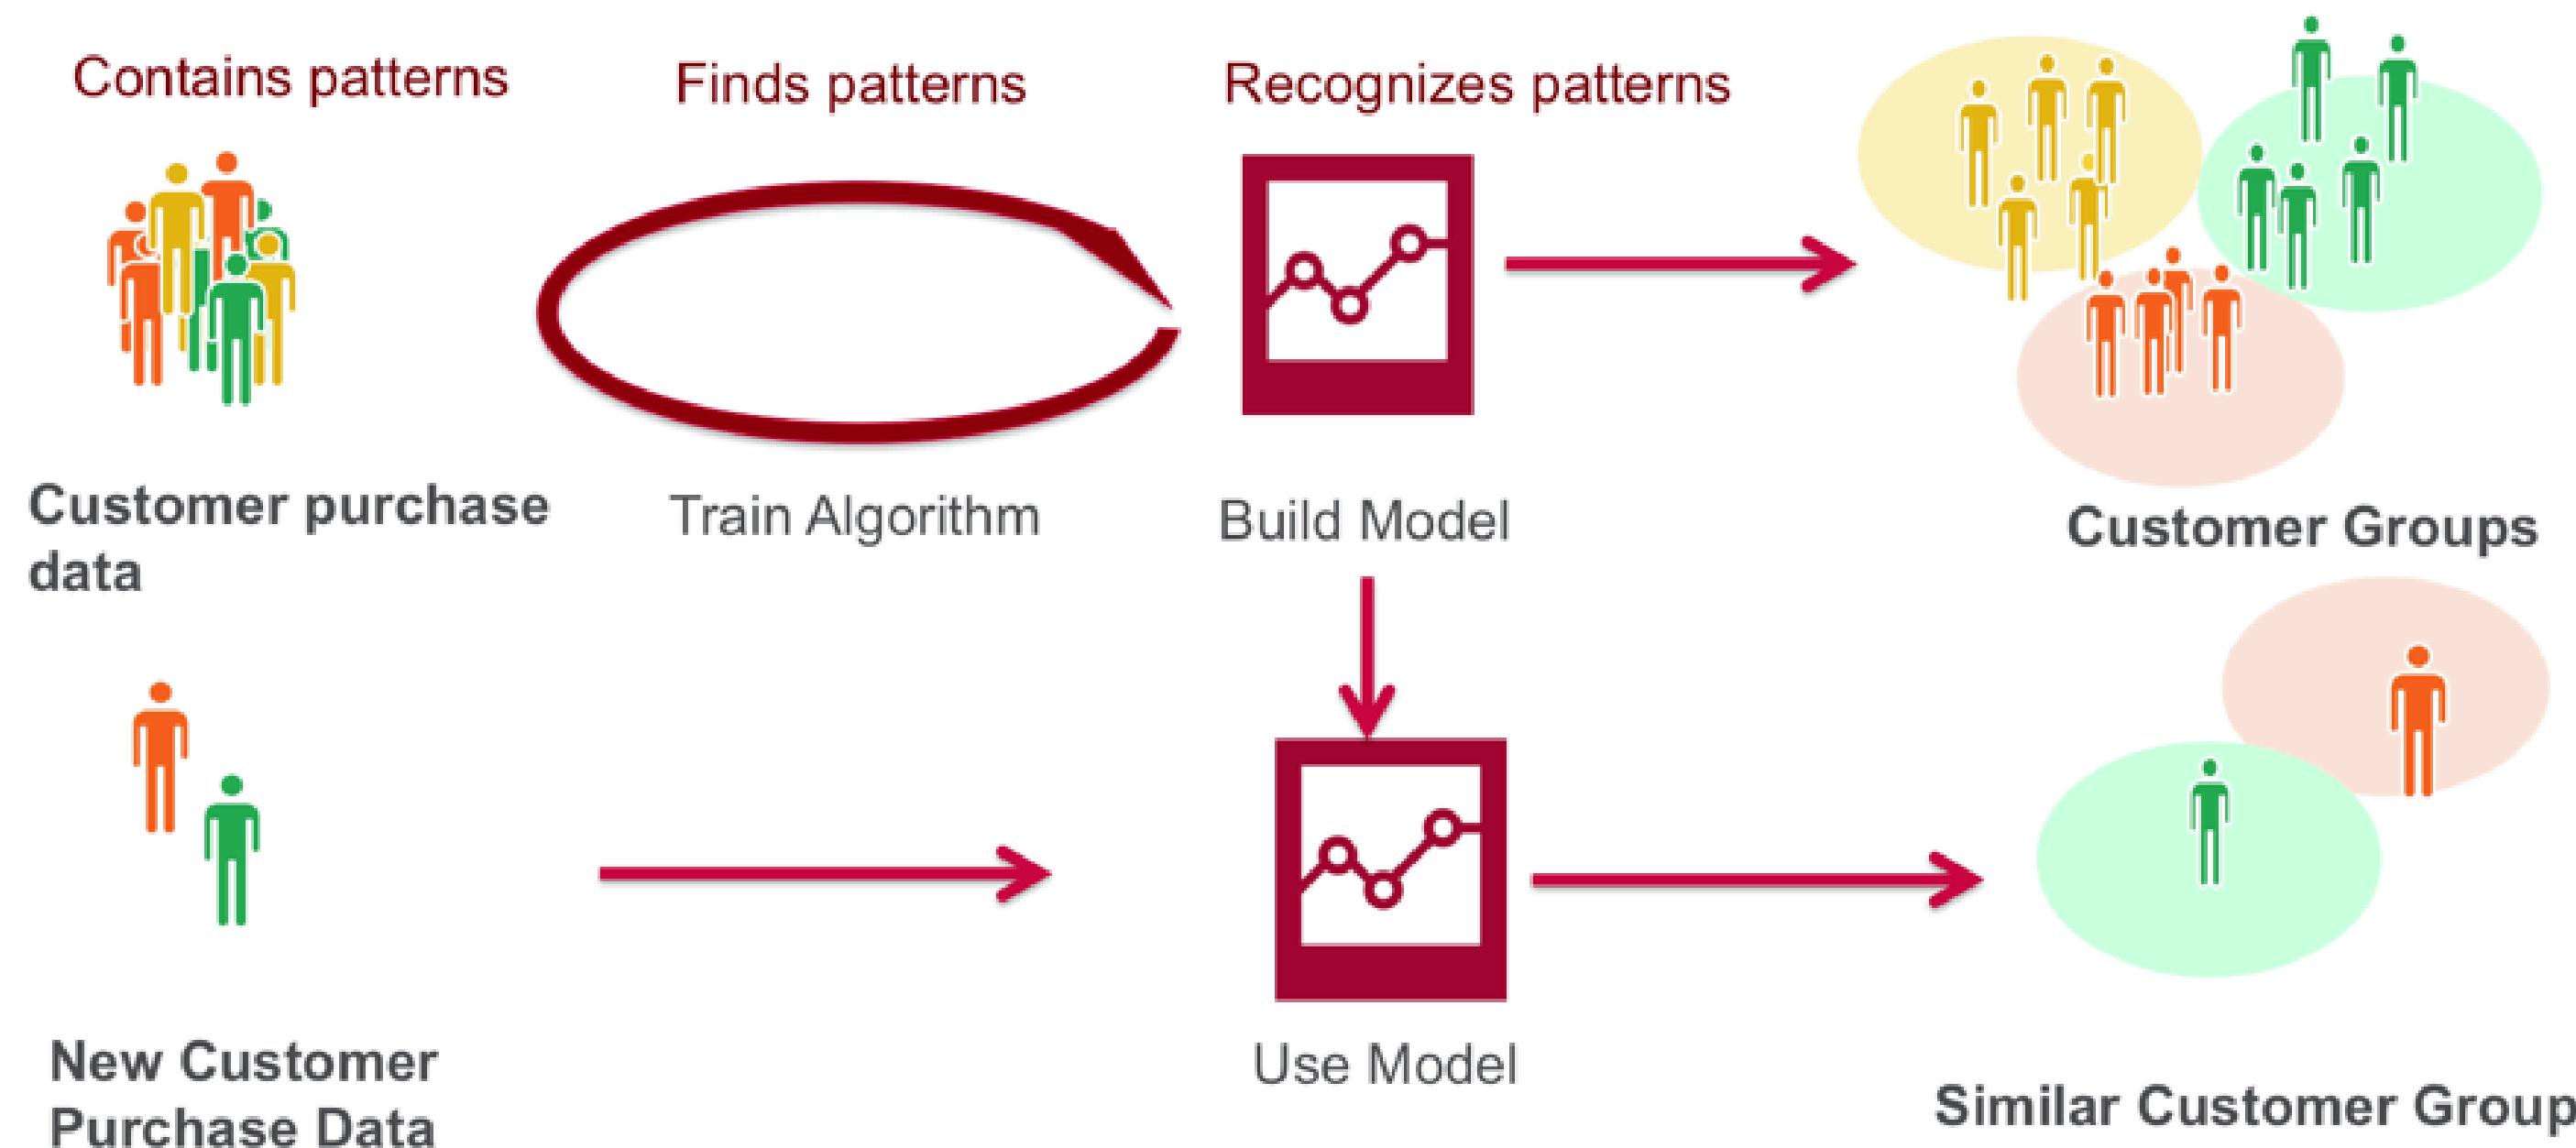
Examples of classification

- Credit card fraud detection (fraud, not fraud)
- Credit card application (good credit, bad credit)
- Email spam detection (spam, not spam)
- Text sentiment analysis (happy, not happy)
- Predicting patient risk (high risk patient, low risk patient)
- Classifying a tumor as malignant or not

Examples of regression

- Predicting the price of the house.
- Predicting age of a person
- Predicting the stock price for tomorrow.
- Predicting next year GDP.

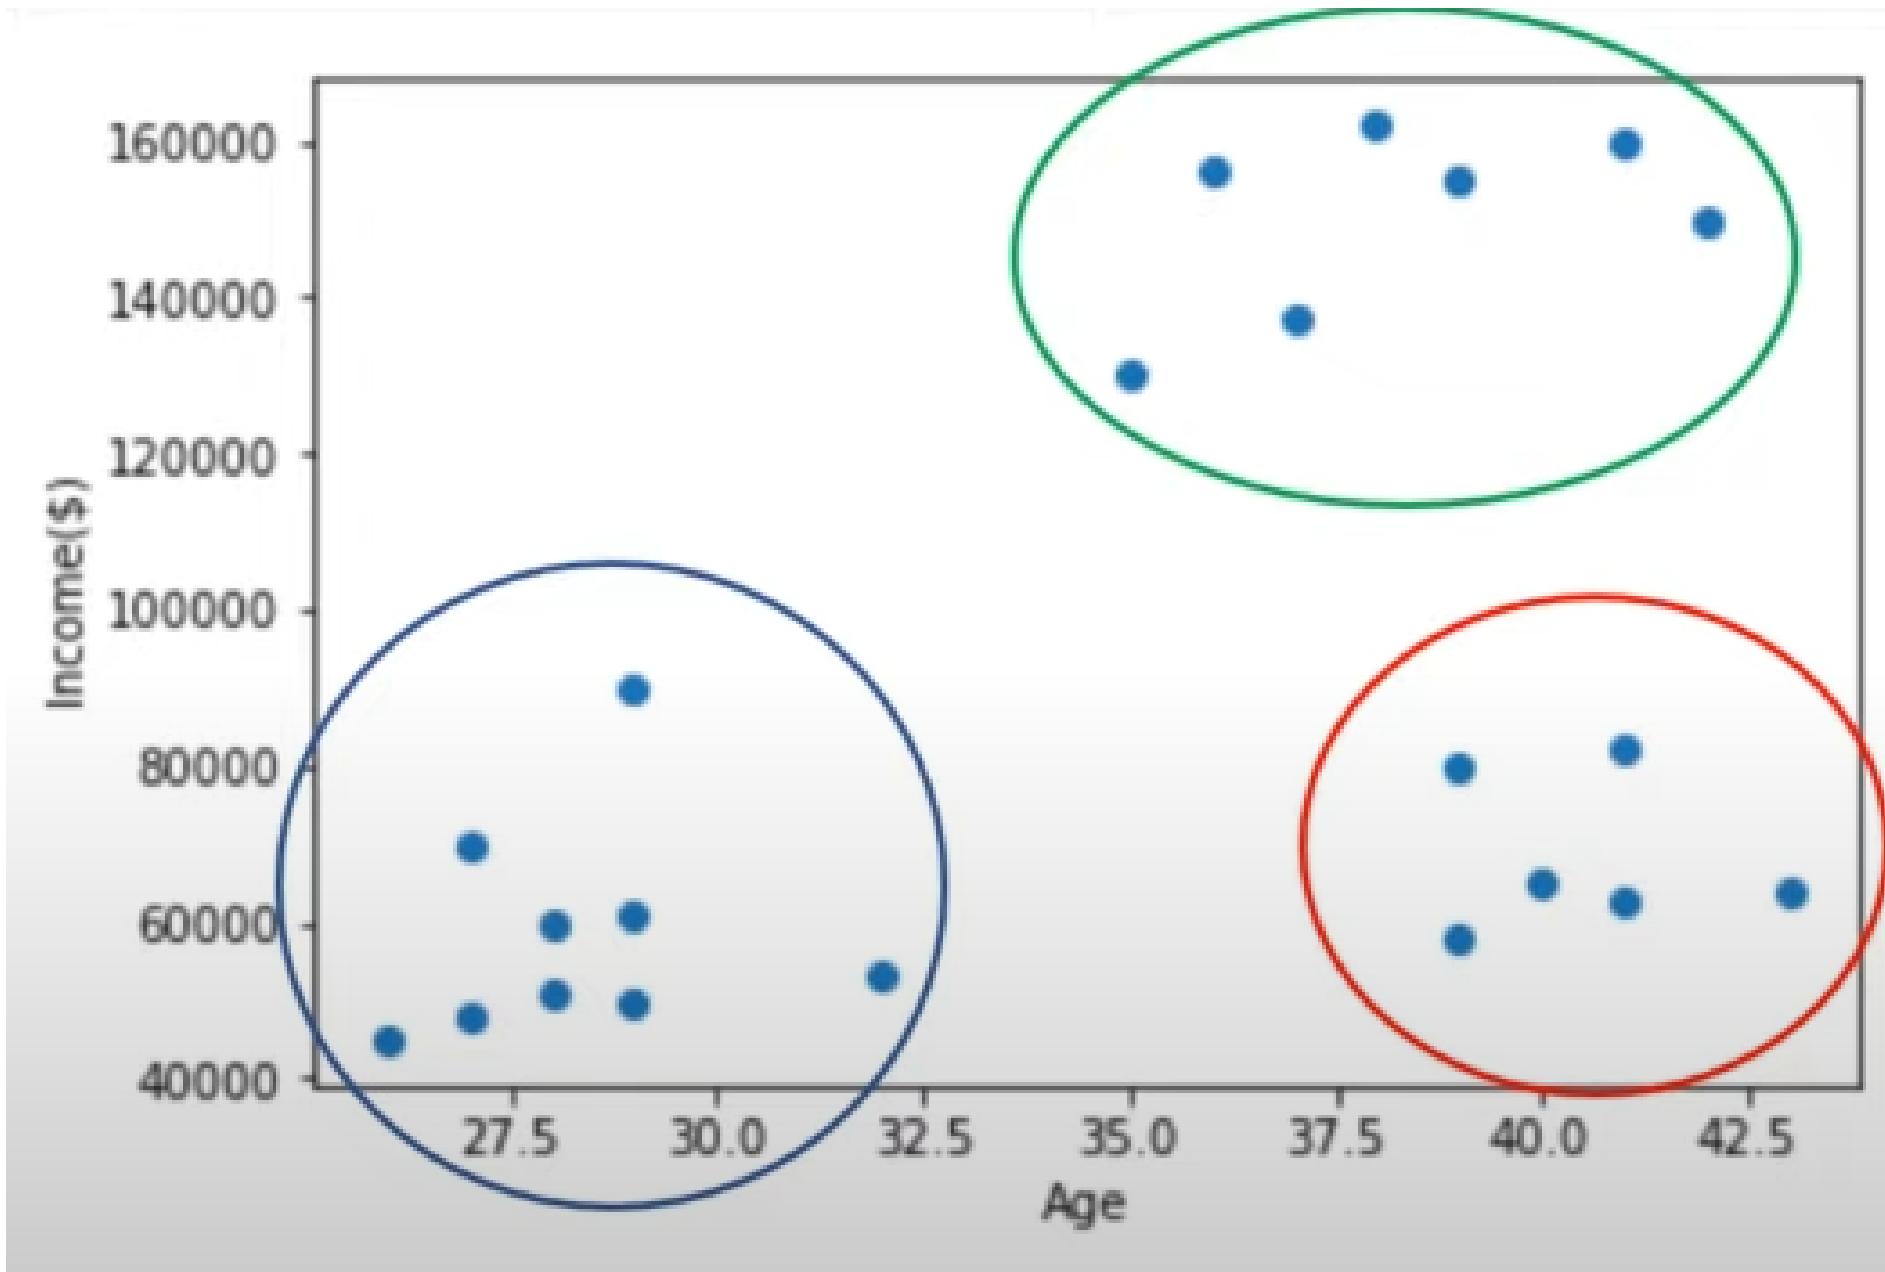
Unsupervised Learning | Self Supervised



Examples of clustering

- Search results grouping
- Grouping similar customers
- Grouping similar patients
- Text categorization
- Network Security Anomaly detection

Unsupervised learning



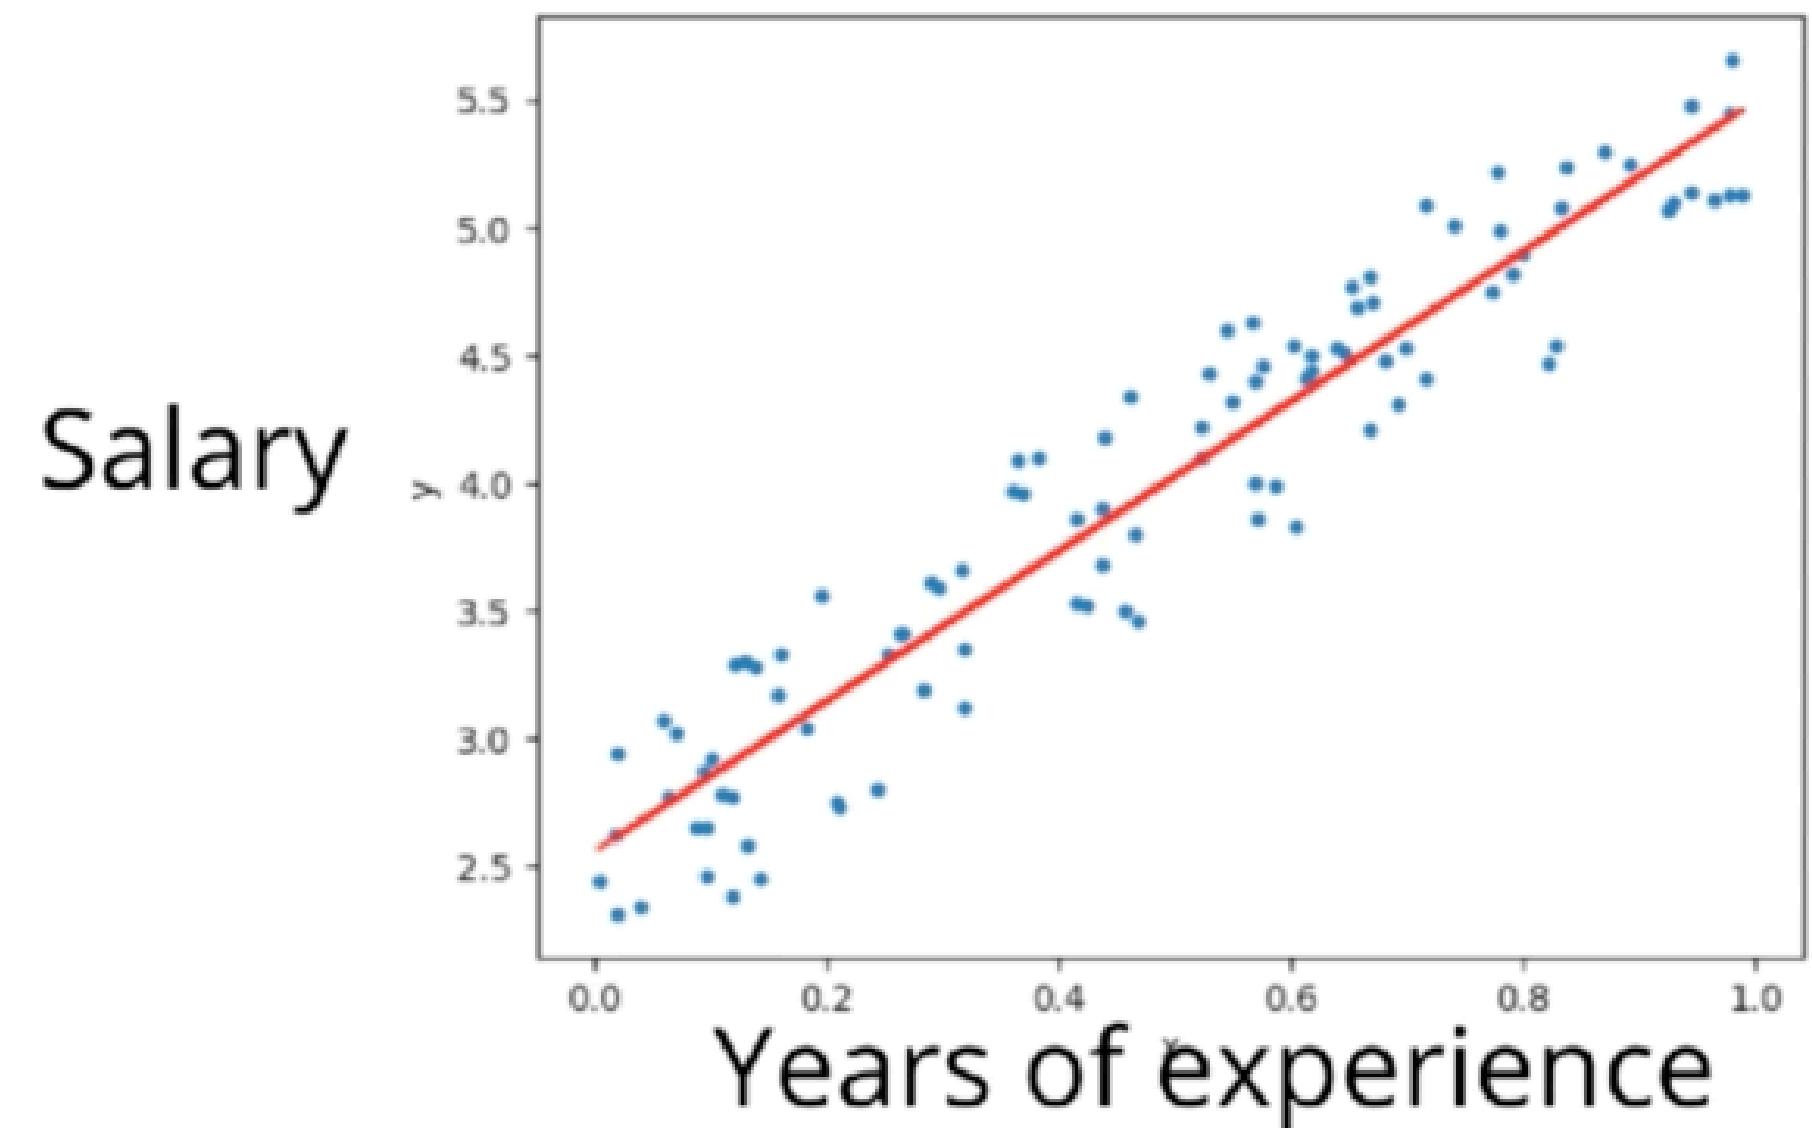
| | A | B | C |
|----|---------|-----|------------|
| 1 | Name | Age | Income(\$) |
| 2 | Rob | 27 | 70000 |
| 3 | Michael | 29 | 90000 |
| 4 | Mohan | 29 | 61000 |
| 5 | Ismail | 28 | 60000 |
| 6 | Kory | 42 | 150000 |
| 7 | Gautam | 39 | 155000 |
| 8 | David | 41 | 160000 |
| 9 | Andrea | 38 | 162000 |
| 10 | Brad | 36 | 156000 |

Example: Regression

- Fit a line or curve
- This is why grumpy old statisticians like to say: “machine learning is just glorified curve-fitting”

$$\hat{y} = mx + b$$

(sometimes we also use “a” for the slope)



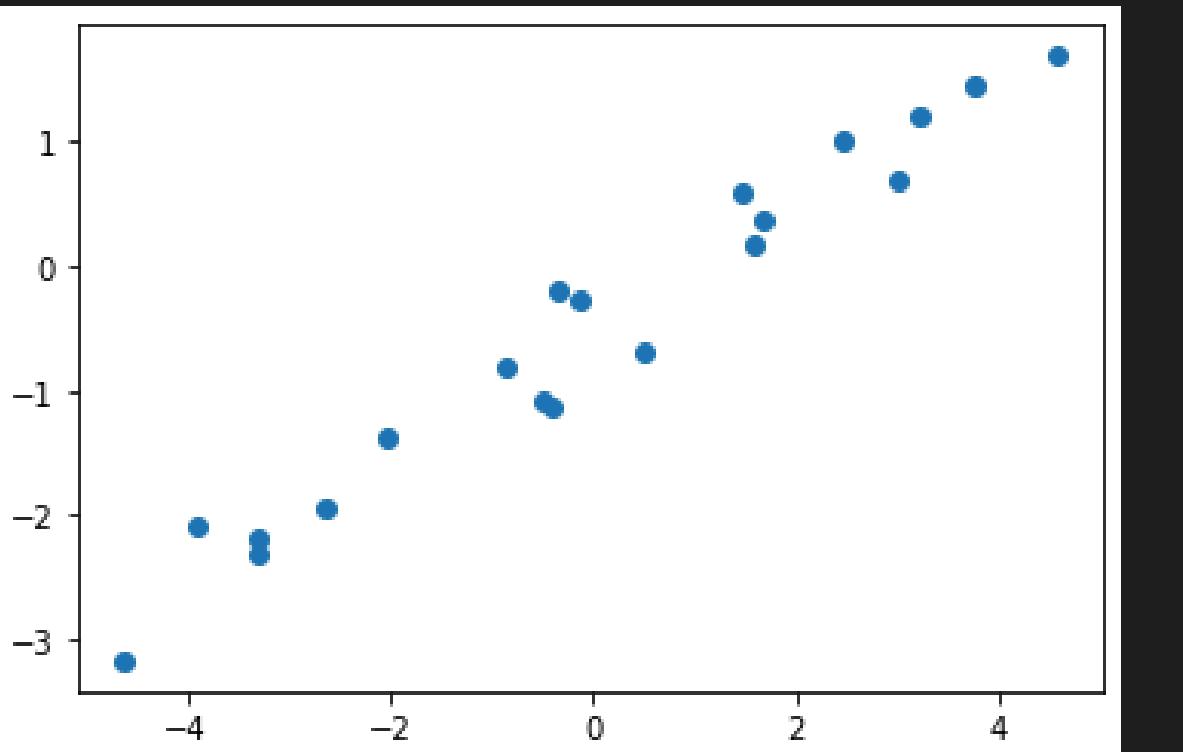
```
import matplotlib.pyplot as plt  
import numpy as np  
✓ 0.4s
```

```
# generate 20 data points  
N = 20  
  
#random data on x-axis  
x= np.random.rand(N)*10-5  
  
y = 0.5*x -1+ np.random.rand(N)
```

```
✓ 0.2s
```

```
plt.scatter(x,y)  
✓ 0.1s
```

```
<matplotlib.collections.PathCollection at 0x7fc1b147a410>
```



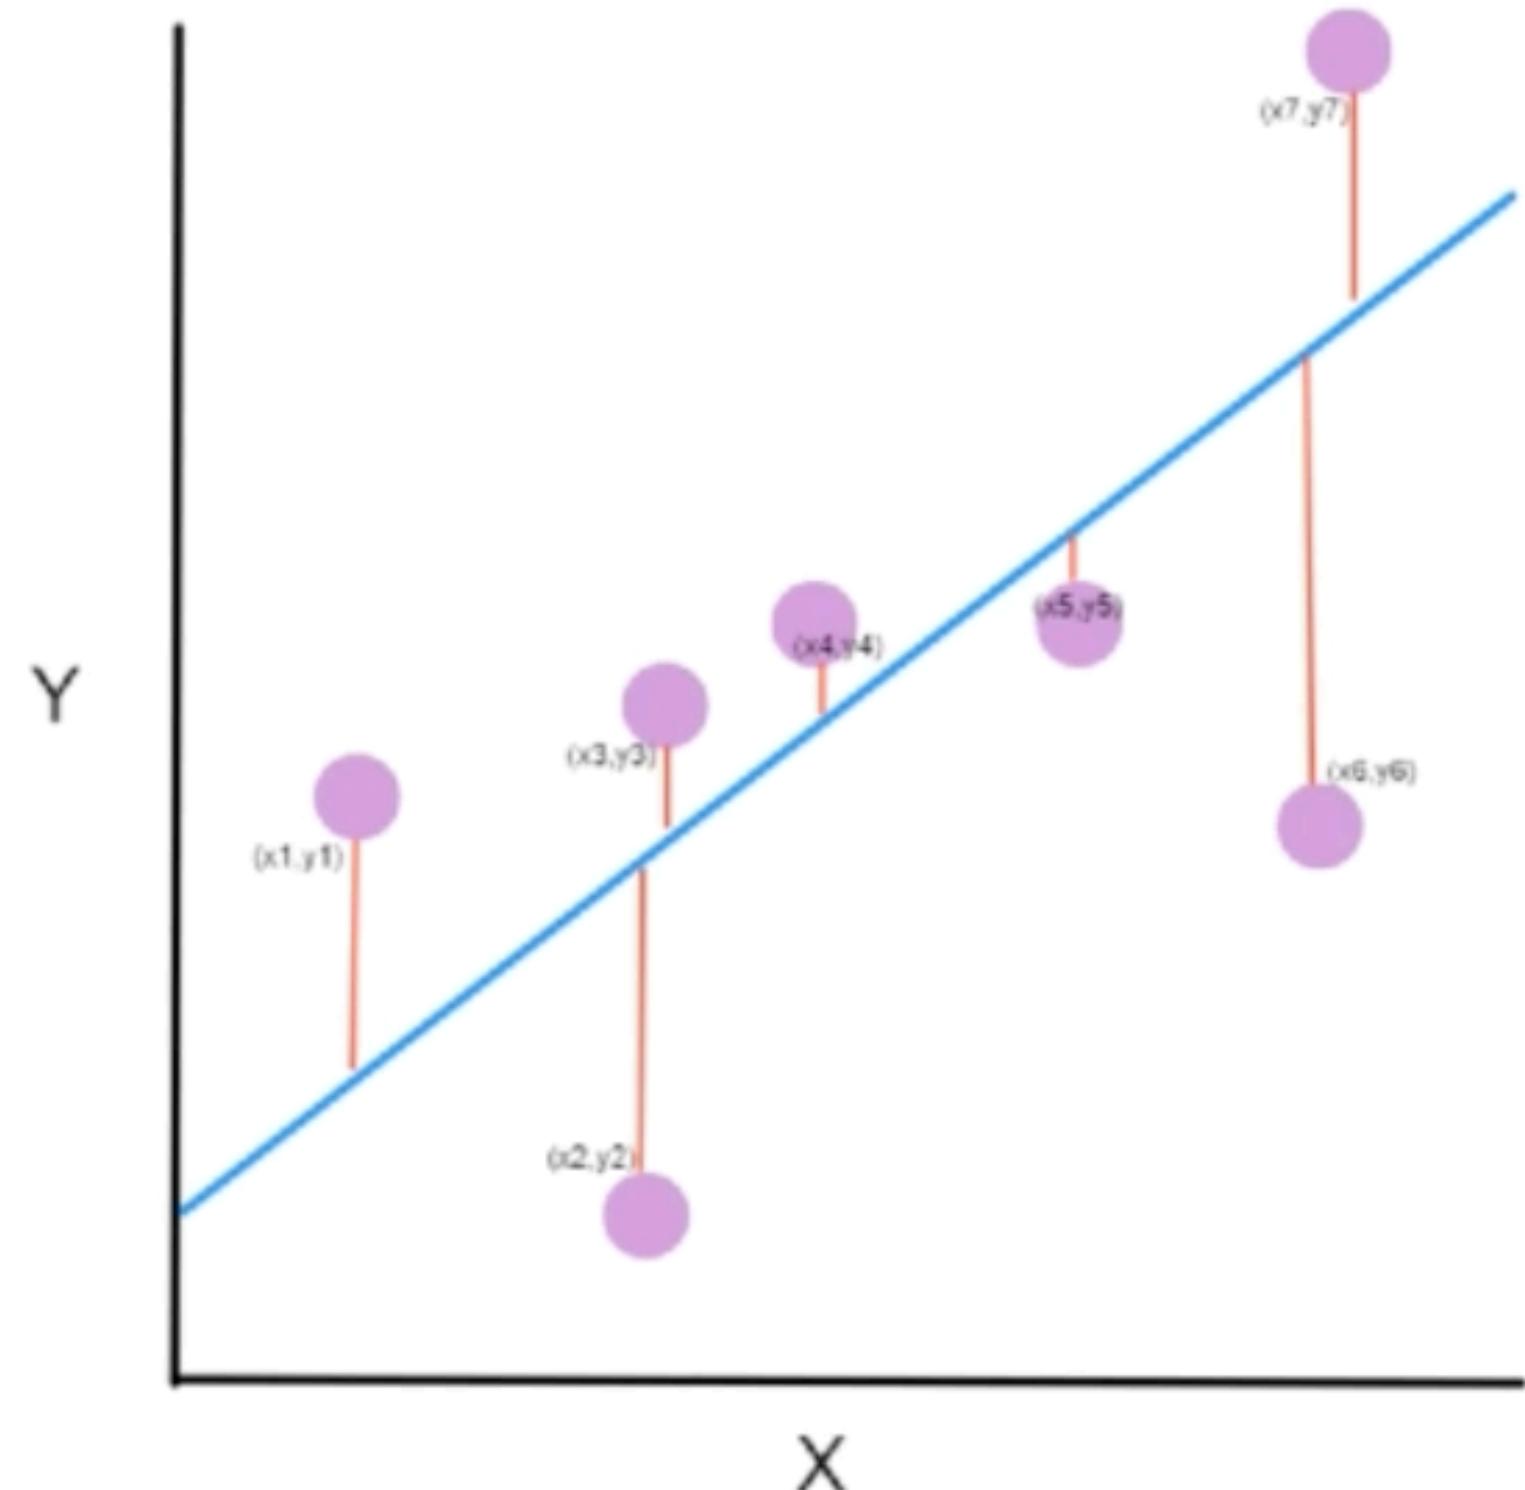
$$\hat{y} = mx + b$$

The Loss

- Data = $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- N = number of samples in the dataset
- The line **cannot** perfectly pass through all the data points

MSE = Mean Squared Error

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$



Applying the MSE to find slope / intercept

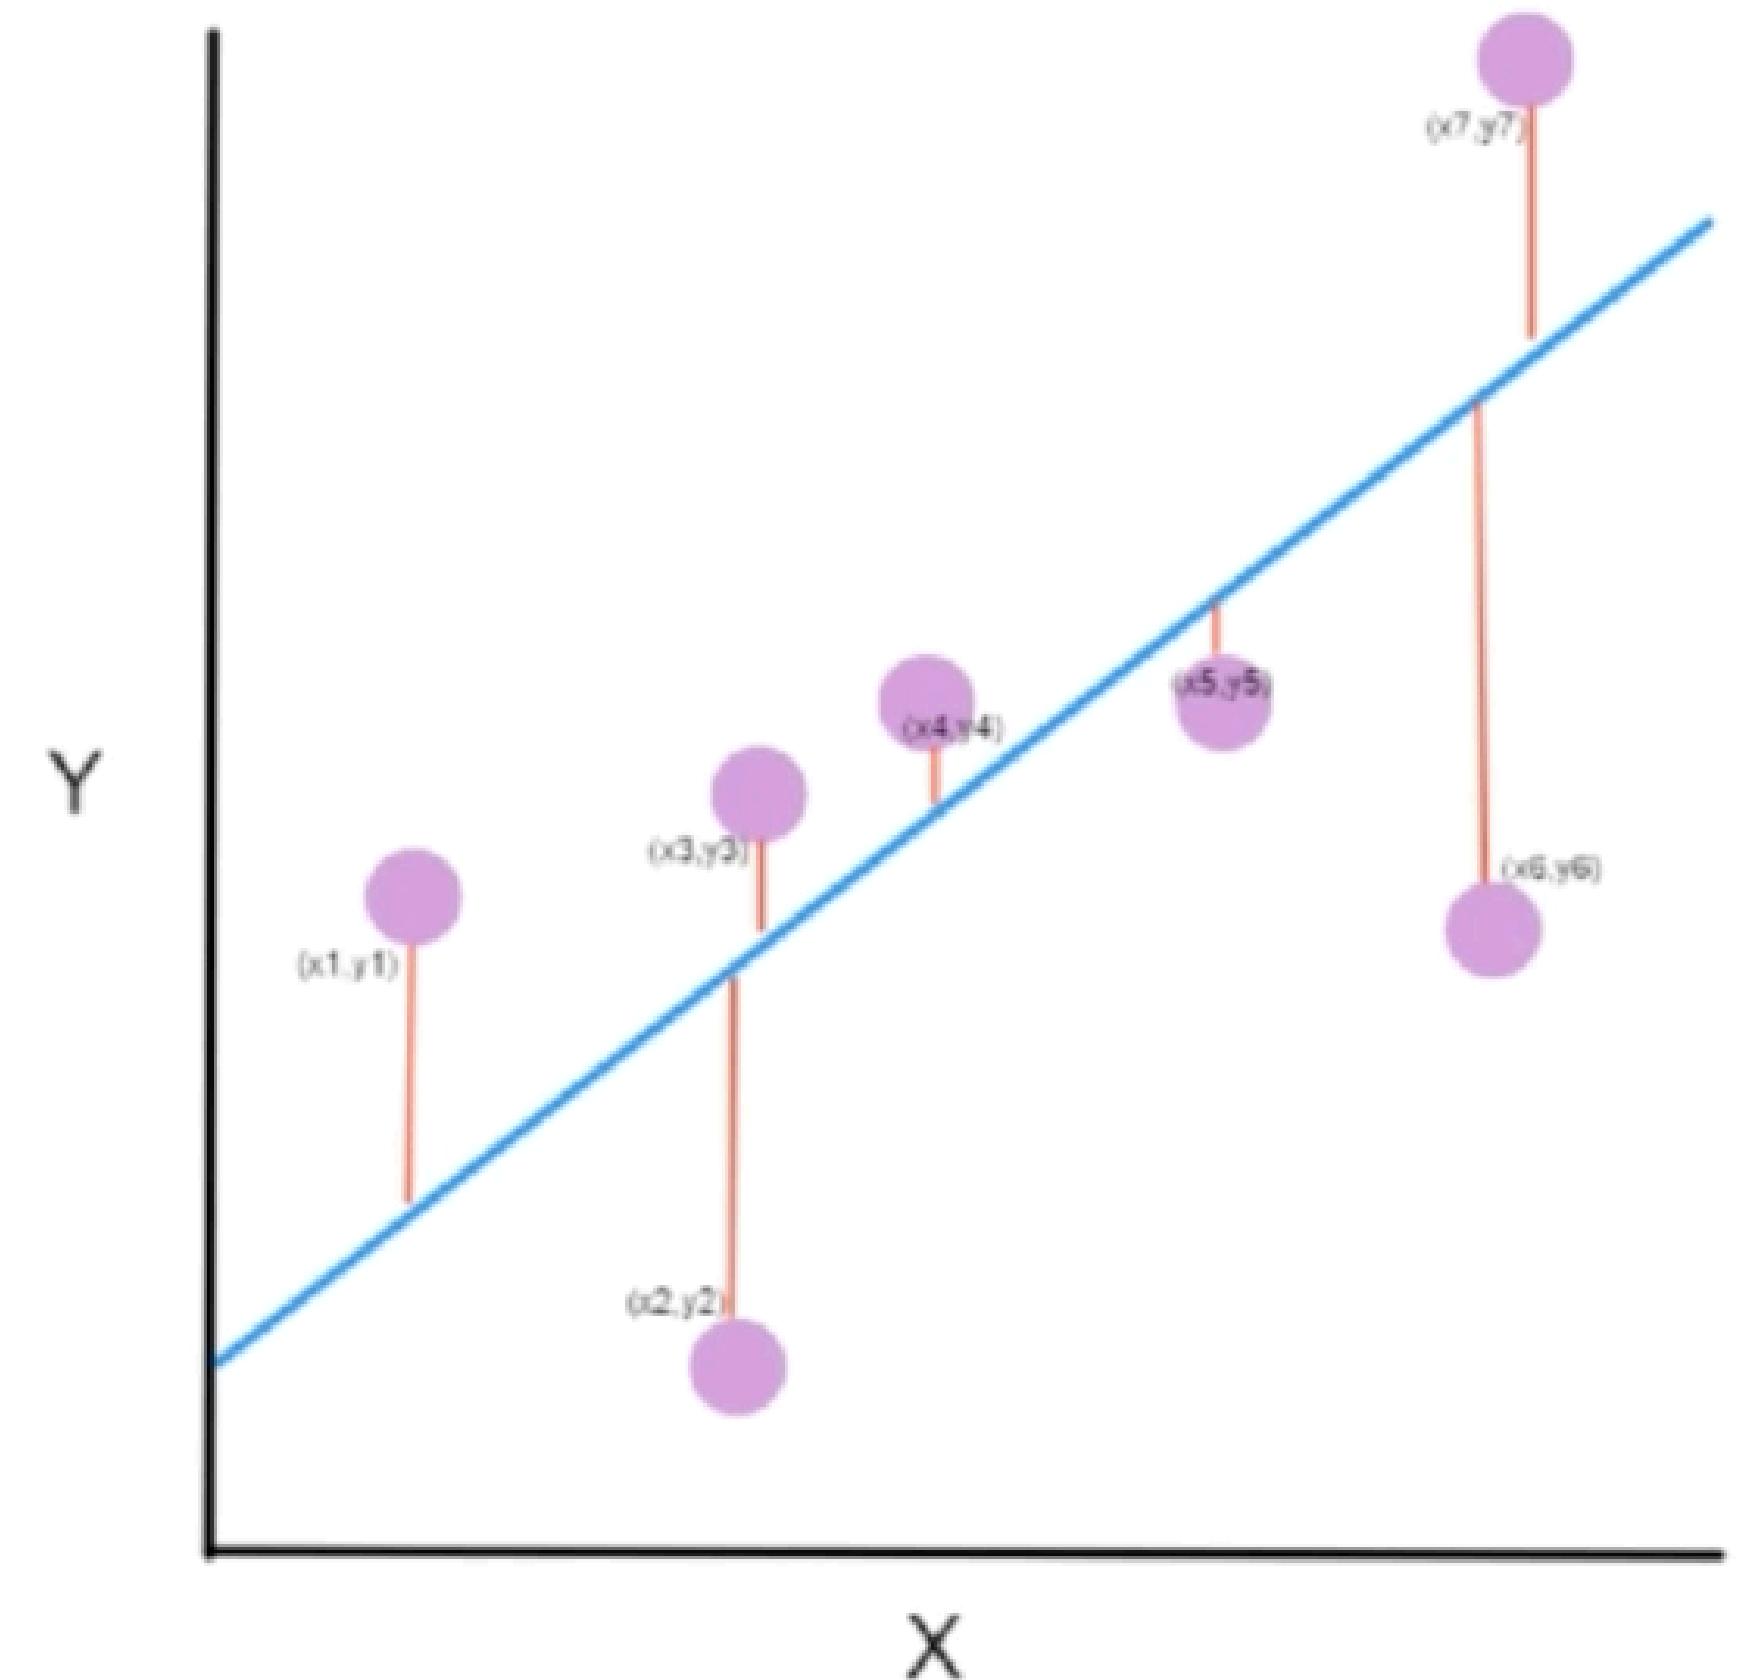
- Plug in the expression for the predictions (\hat{y}_i)
- Quiz: What are the **variables** in this expression?

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

Applying the MSE to find slope / intercept

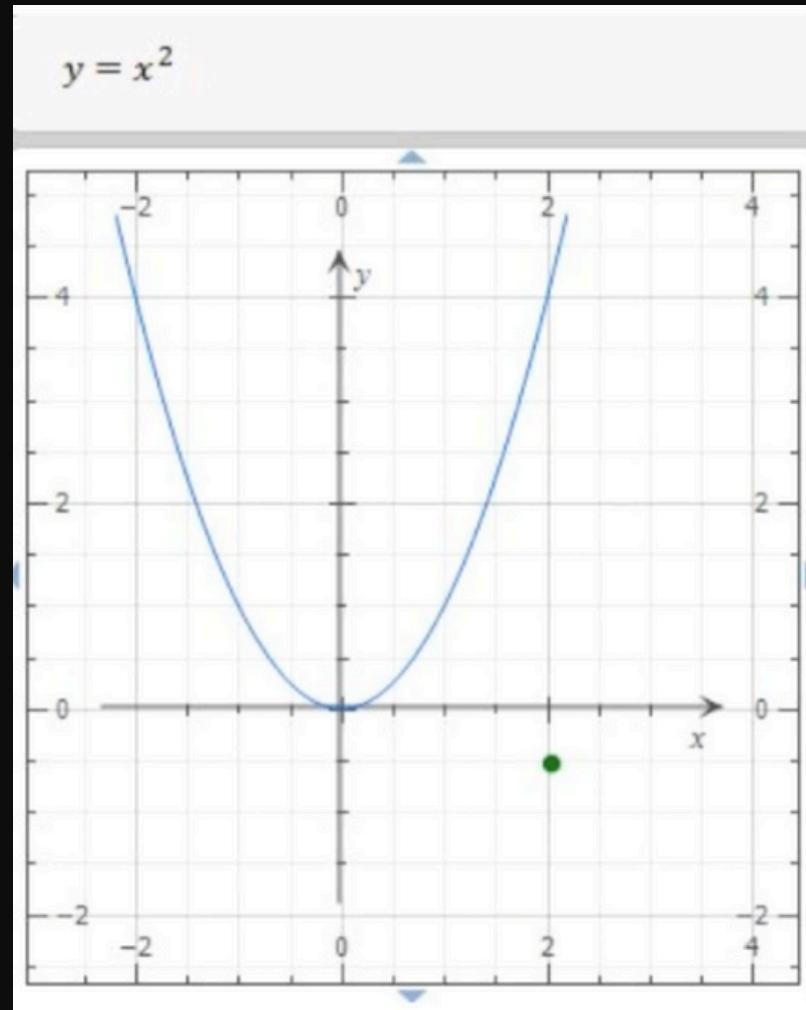
- Note: L (loss) = MSE
- We want to *minimize the loss* with respect to the parameters (m, b)

$$m^*, b^* = \arg \min_{m, b} L$$



First Solution

Minimize $f(x) = x^2$



- Use calculus!
- $df / dx = 2x = 0$
- Solve for x : $x = 0$

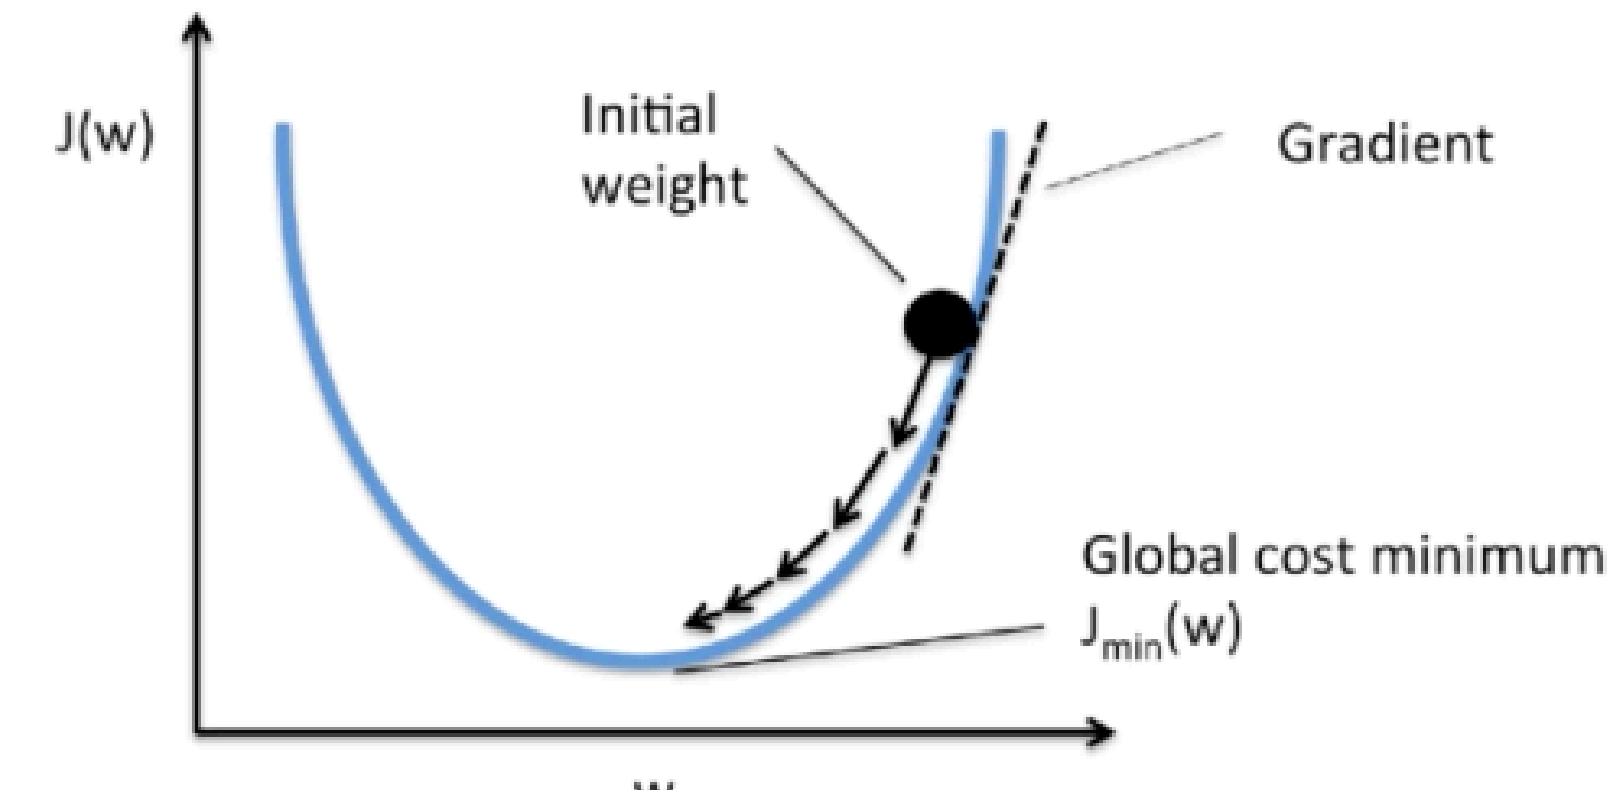
Second Solution : Gradient Descent

- Find the derivatives and set them to 0, solve for the parameters
- Yields 2 equations and 2 unknowns: can solve for (m, b)
- Try it as an exercise!: Should be in terms of $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$

$$\partial L/\partial m = 0, \partial L/\partial b = 0$$

Gradient Descent to find hidden parameters

```
# gradient descent  
# pseudocode  
 $\theta$  = (m, b) = random()  
for i in range(n_epochs):  
     $\theta$  =  $\theta$  -  $\eta \nabla_{\theta} L$ 
```

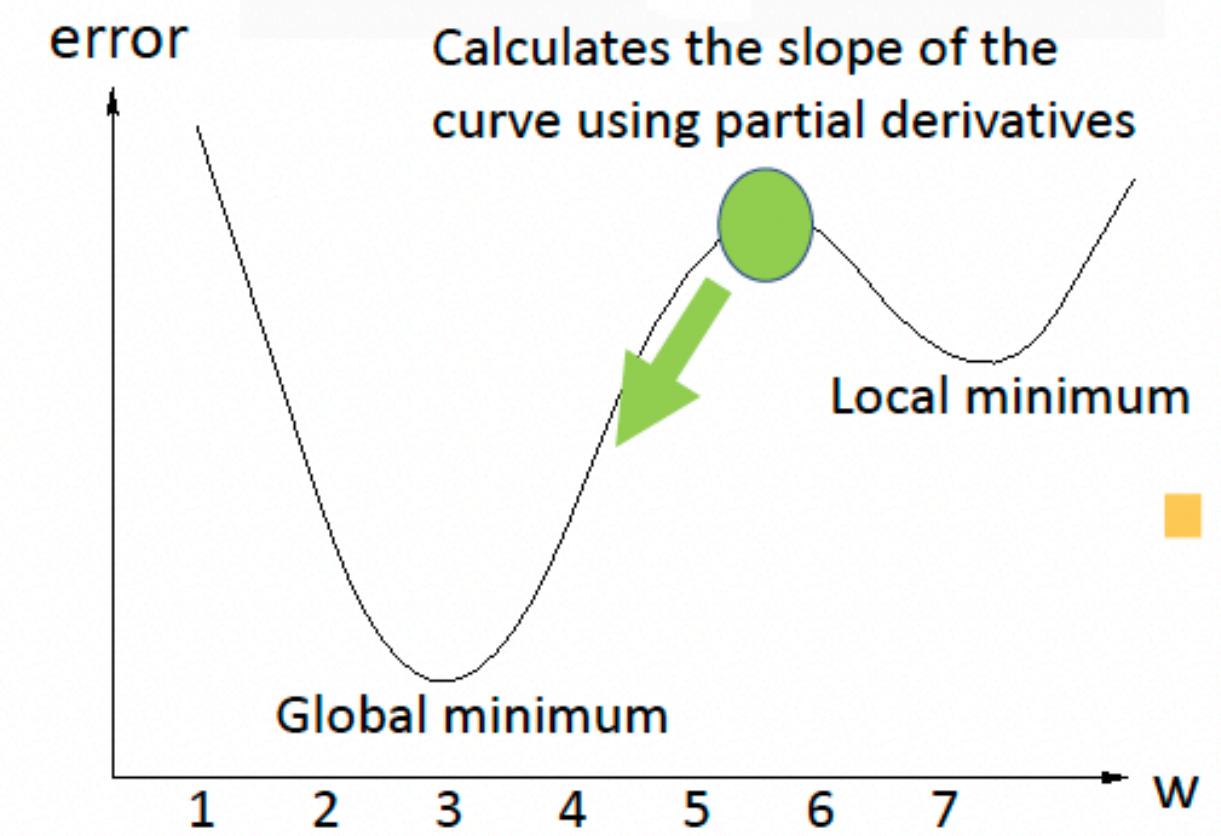
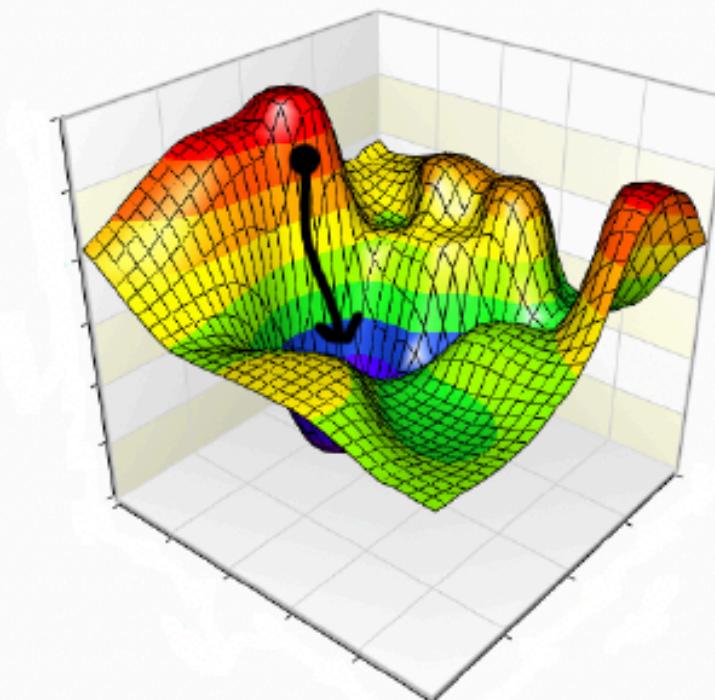
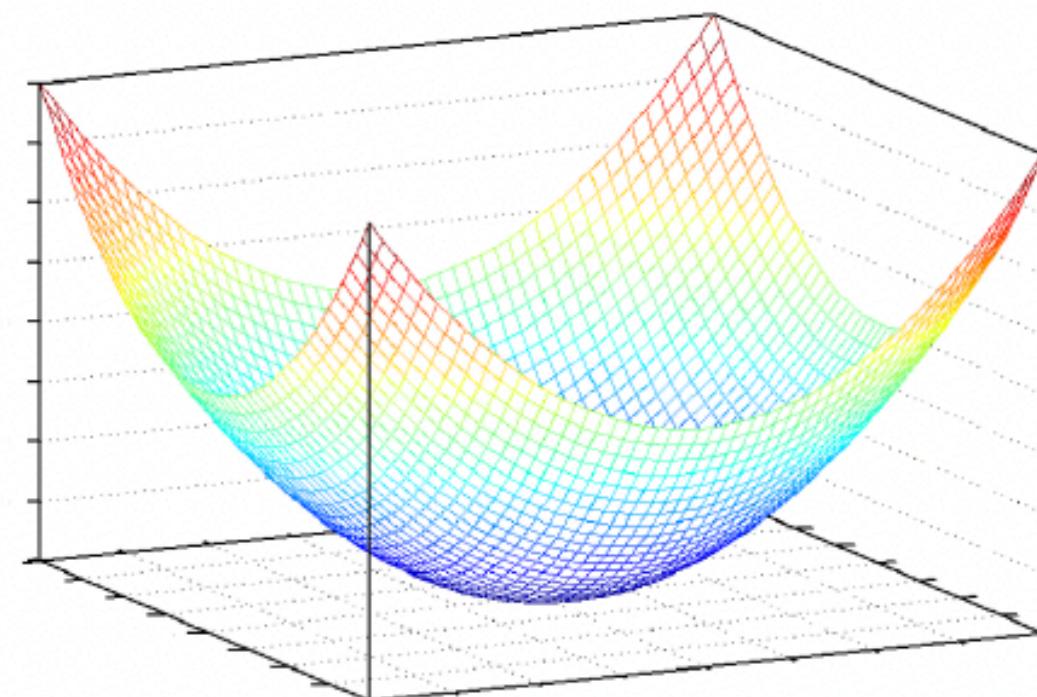


$$L = MSE = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

GRADIENT DESCENT

$$\min C(w_1, w_2 \dots w_n)$$

Calculate the partial derivative to move to the gradient direction



$$\frac{\partial}{\partial m} = \frac{2}{n} \sum_{i=1}^n -x_i(y_i - (mx_i + b))$$

$$\frac{\partial}{\partial b} = \frac{2}{n} \sum_{i=1}^n -(y_i - (mx_i + b))$$

```
def Gradient(m,b) :  
    grad_m = 0; grad_b = 0;  
  
    for i in range(N) :  
        grad_m += -X[i]*(Y[i] - (m*X[i]+b))  
        grad_b += -(Y[i] - (m*X[i]+b))  
  
    return grad_m, grad_b
```

```
Gradient(-1,-1)
```

✓ 0.3s

```
(array([-209.28382035]), array([-10.12894605]))
```

```

# gradient descent
# pseudocode

θ = (m, b) = random()
for i in range(n_epochs):
    θ = θ - η∇_θ L

```

```

# Define the gradient function
def Gradient(m, b):
    grad_m = 0
    grad_b = 0
    for i in range(N):
        grad_m += -X[i] * (Y[i] - (m * X[i] + b))
        grad_b += -(Y[i] - (m * X[i] + b))
    return grad_m, grad_b

# Initialize m and b
m = 0
b = 0

# Set learning rate and number of iterations
learning_rate = 0.001
iterations = 1000

# Perform gradient descent
for _ in range(iterations):
    grad_m, grad_b = Gradient(m, b)
    m -= learning_rate * grad_m
    b -= learning_rate * grad_b

print("Estimated m:", m)
print("Estimated b:", b)

```

✓ 0.0s

```

Estimated m: 0.4860476543029485
Estimated b: -0.8921942796543206

```

```

# Gradient descent parameters
learning_rate = 0.001
iterations = 50000
m = 0 # Initial guess for m
b = 0 # Initial guess for b

# Lists to store parameters history
m_history = []
b_history = []
loss_history = []

# Gradient descent
for i in range(iterations):
    # Calculate gradients
    grad_m, grad_b = Gradient(m, b, X, Y)

    # Update parameters
    m = m - learning_rate * grad_m
    b = b - learning_rate * grad_b

    # Store parameters
    m_history.append(m)
    b_history.append(b)

    # Calculate and store loss
    loss = np.mean((Y - (m*X + b))**2)
    loss_history.append(loss)

print(f"Final parameters: m = {m:.4f}, b = {b:.4f}")
print(f"True parameters: m = 0.5000, b = -1.0000")

# Plotting
plt.figure(figsize=(15, 5))

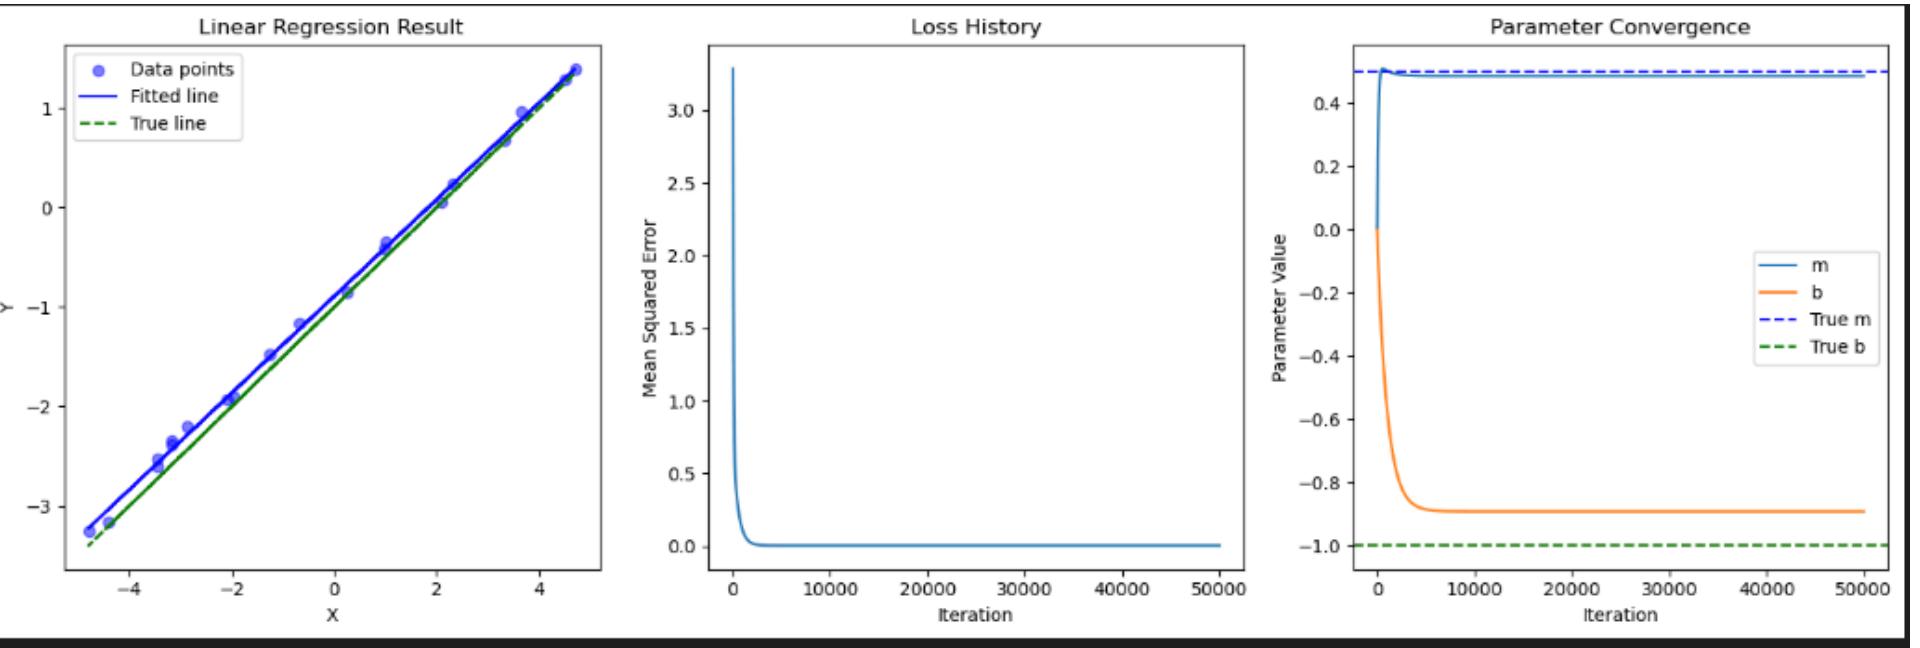
# Plot 1: Data and regression line
plt.subplot(131)
plt.scatter(X, Y, color='blue', alpha=0.5, label='Data points')
plt.plot(X, m*X + b, color='blue', label='Fitted line')
plt.plot(X, 0.5*X - 1, '--', color='green', label='True line')
plt.xlabel('X')
plt.ylabel('Y')
plt.title('Linear Regression Result')
plt.legend()

# Plot 2: Loss history
plt.subplot(132)
plt.plot(loss_history)
plt.xlabel('Iteration')
plt.ylabel('Mean Squared Error')
plt.title('Loss History')

# Plot 3: Parameter convergence
plt.subplot(133)
plt.plot(m_history, label='m')
plt.plot(b_history, label='b')
plt.axhline(y=0.5, color='b', linestyle='--', label='True m')
plt.axhline(y=-1, color='g', linestyle='--', label='True b')
plt.xlabel('Iteration')
plt.ylabel('Parameter Value')
plt.title('Parameter Convergence')
plt.legend()

plt.tight_layout()
plt.show()

```



Final parameters: $m = 0.4860$, $b = -0.8922$
 True parameters: $m = 0.5000$, $b = -1.0000$

Third Solution :Pytorch

SDG: Stochastic gradient
descent

$$\theta = \theta - \eta \nabla_{\theta} L$$

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

```
# Loss and optimizer
criterion = nn.MSELoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

```
# gradient descent
# pseudocode

θ = (m, b) = random()
for i in range(n_epochs):
    θ = θ - η∇_θ L
```

```
# Train the model
n_epochs = 30
for it in range(n_epochs):
    # zero the parameter gradients
    optimizer.zero_grad()

    # Forward pass
    outputs = model(inputs)
    loss = criterion(outputs, targets)

    # Backward and optimize
    loss.backward()
    optimizer.step()
```

```

# Convert data to PyTorch tensors
X_train = torch.tensor(X, dtype=torch.float32).view(-1, 1)
Y_train = torch.tensor(Y, dtype=torch.float32).view(-1, 1)

# Initialize m and b as PyTorch parameters
m = torch.randn(1, requires_grad=True, dtype=torch.float32) # Random initialization for m
b = torch.randn(1, requires_grad=True, dtype=torch.float32) # Random initialization for b

# Define optimizer
learning_rate = 0.01
optimizer = torch.optim.SGD([m, b], lr=learning_rate) # Optimizer for m and b

# Training loop
epochs = 1000
for epoch in range(epochs):

    optimizer.zero_grad() # Clear gradients

    # Forward pass: Calculate y_pred using y = mx + b
    y_pred = m * X_train + b
    # Calculate Mean Squared Error (MSE) Loss
    loss = torch.mean((y_pred - Y_train) ** 2)

    # Backward pass: Compute gradients
    loss.backward() # Calculate new gradients

    # Update parameters using optimizer1
    optimizer.step()

    # Print the loss every 100 epochs
    if (epoch + 1) % 100 == 0:
        print(f'Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}')

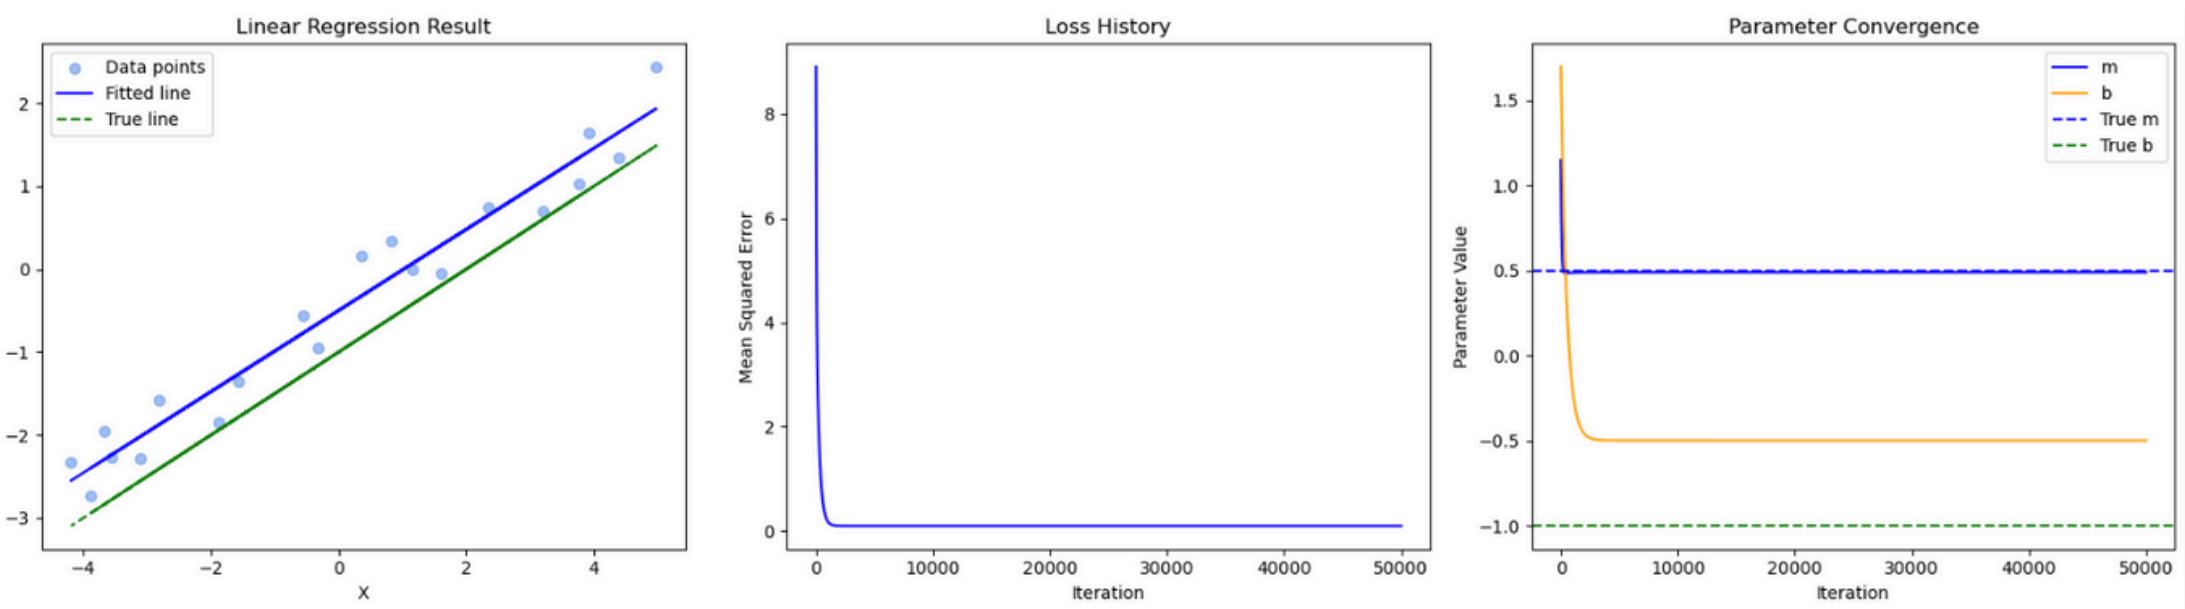
print("Estimated m:", m.item())
print("Estimated b:", b.item())

✓ 0.4s

Epoch [100/1000], Loss: 0.0752
Epoch [200/1000], Loss: 0.0240
Epoch [300/1000], Loss: 0.0230
Epoch [400/1000], Loss: 0.0230
Epoch [500/1000], Loss: 0.0230
Epoch [600/1000], Loss: 0.0230
Epoch [700/1000], Loss: 0.0230
Epoch [800/1000], Loss: 0.0230
Epoch [900/1000], Loss: 0.0230
Epoch [1000/1000], Loss: 0.0230
Estimated m: 0.48729196190834045
Estimated b: -0.7519025206565857

```

Pytorch Method 1

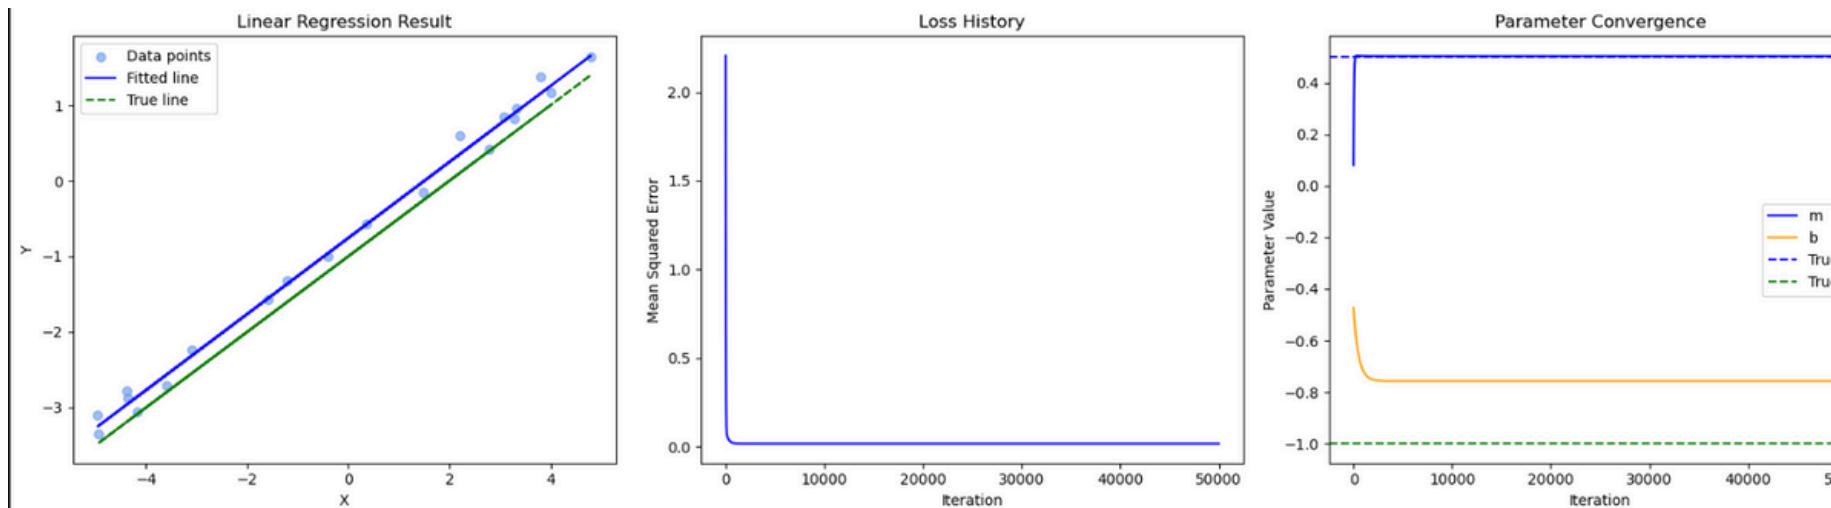


Pytorch Method 2

```
# Convert data to PyTorch tensors
X_train = torch.tensor(X, dtype=torch.float32).view(-1, 1)
Y_train = torch.tensor(Y, dtype=torch.float32).view(-1, 1)

# Define the model class
class LinearRegressionModel(nn.Module):
    def __init__(self):
        super(LinearRegressionModel, self).__init__()
        # Initialize m and b as parameters with random values
        self.m = nn.Parameter(torch.randn(1, dtype=torch.float32))
        self.b = nn.Parameter(torch.randn(1, dtype=torch.float32))

    def forward(self, x):
        # Apply the linear transformation y = mx + b
        return self.m * x + self.b
```



```
# Initialize model and optimizer
model = LinearRegressionModel()
learning_rate = 0.001
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

# Training loop with data recording for plots
epochs = 50000
loss_history = []
m_values, b_values = [], []

for epoch in range(epochs):
    # Forward pass
    y_pred = model(X_train)

    # Calculate Mean Squared Error (MSE) Loss
    loss = torch.mean((y_pred - Y_train) ** 2)
    loss_history.append(loss.item())
    m_values.append(model.m.item())
    b_values.append(model.b.item())

    # Backward pass and parameter update
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

Pytorch Method 3

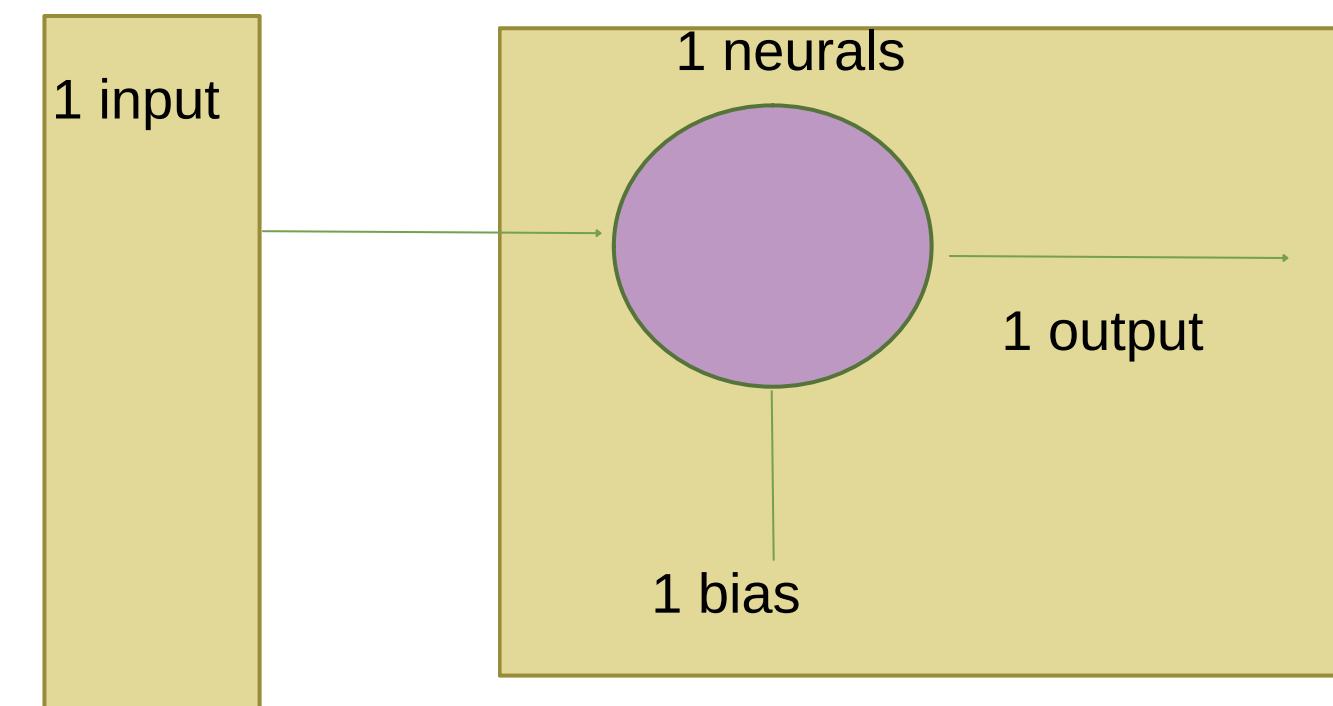
```
1 # Train the model
2 n_epochs = 30
3 losses = []
4 for it in range(n_epochs):
5     # zero the parameter gradients
6     optimizer.zero_grad()
7
8     # Forward pass
9     outputs = model(inputs)
10    loss = criterion(outputs, targets)
11
12    # keep the loss so we can plot it later
13    losses.append(loss.item())
14
15    # Backward and optimize
16    loss.backward()
17    optimizer.step()
18
19    print(f'Epoch {it+1}/{n_epochs}, Loss: {loss.item():.4f}')
```

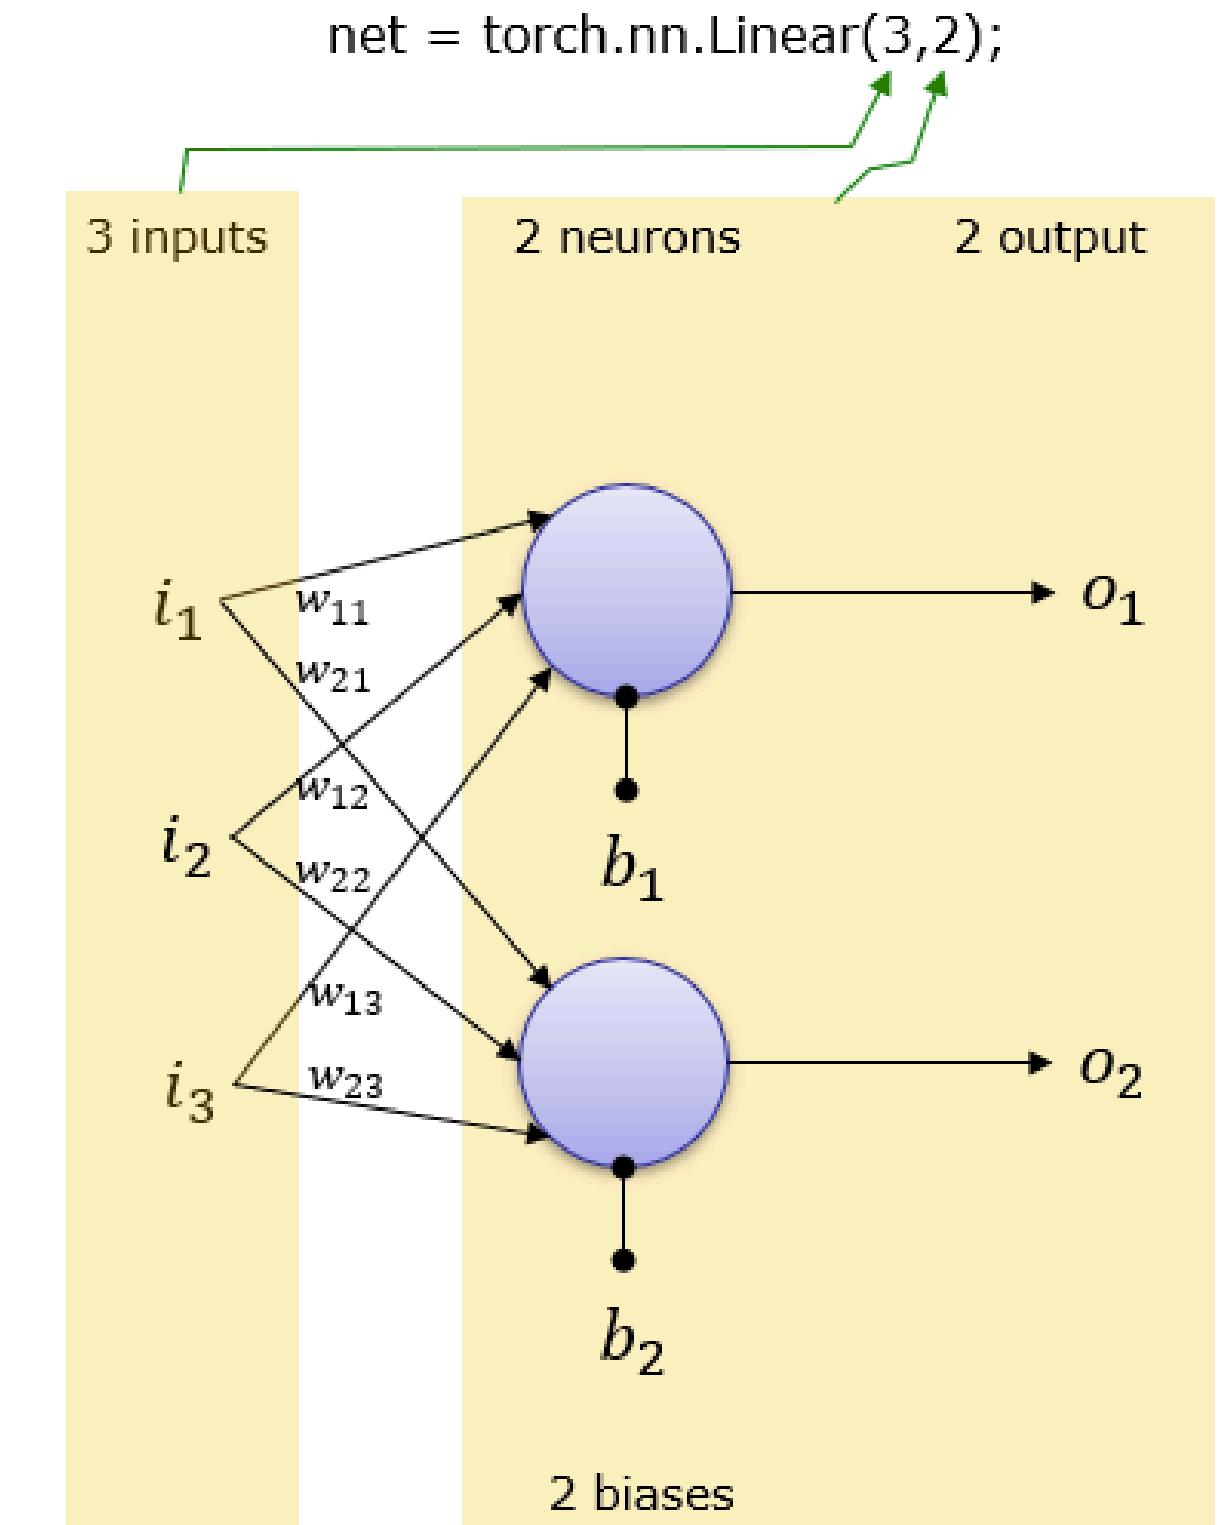
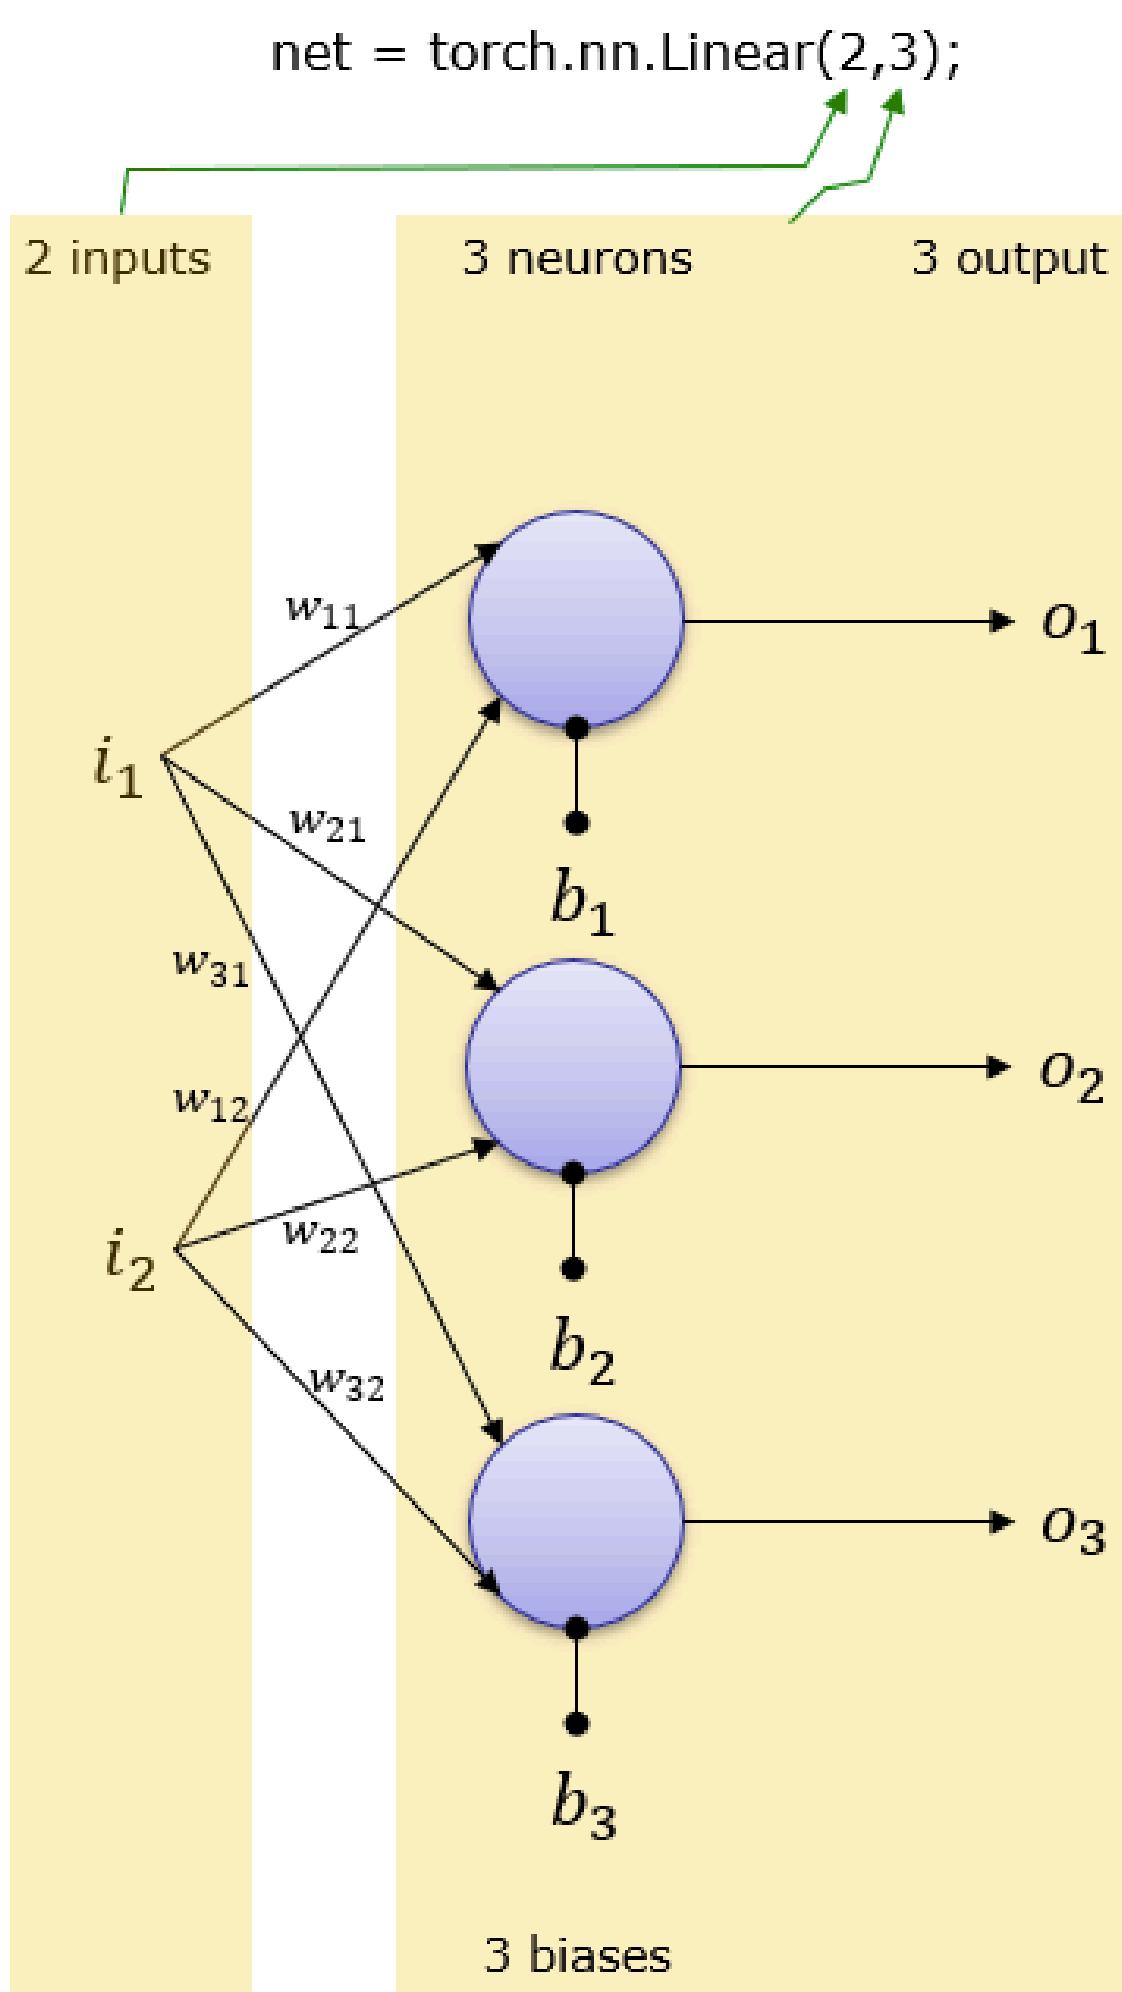
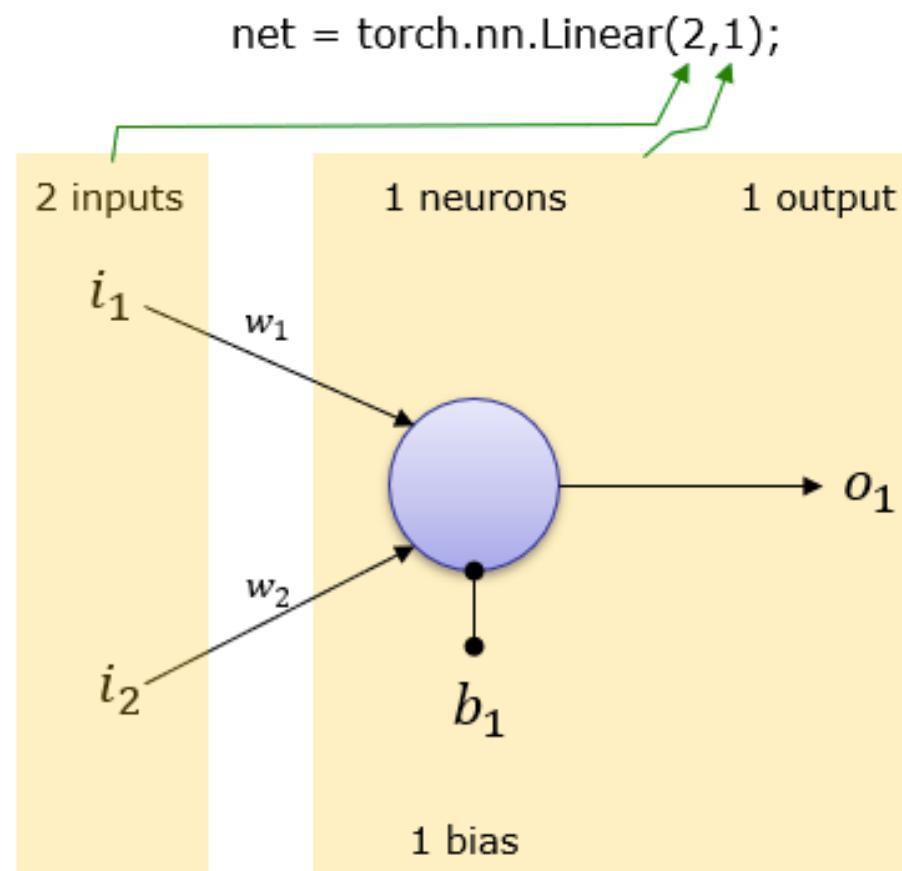
```
Epoch 1/30, Loss: 12.1352
Epoch 2/30, Loss: 2.8066
Epoch 3/30, Loss: 2.4554
Epoch 4/30, Loss: 2.1753
Epoch 5/30, Loss: 1.9483
Epoch 6/30, Loss: 1.7641
Epoch 7/30, Loss: 1.6148
```

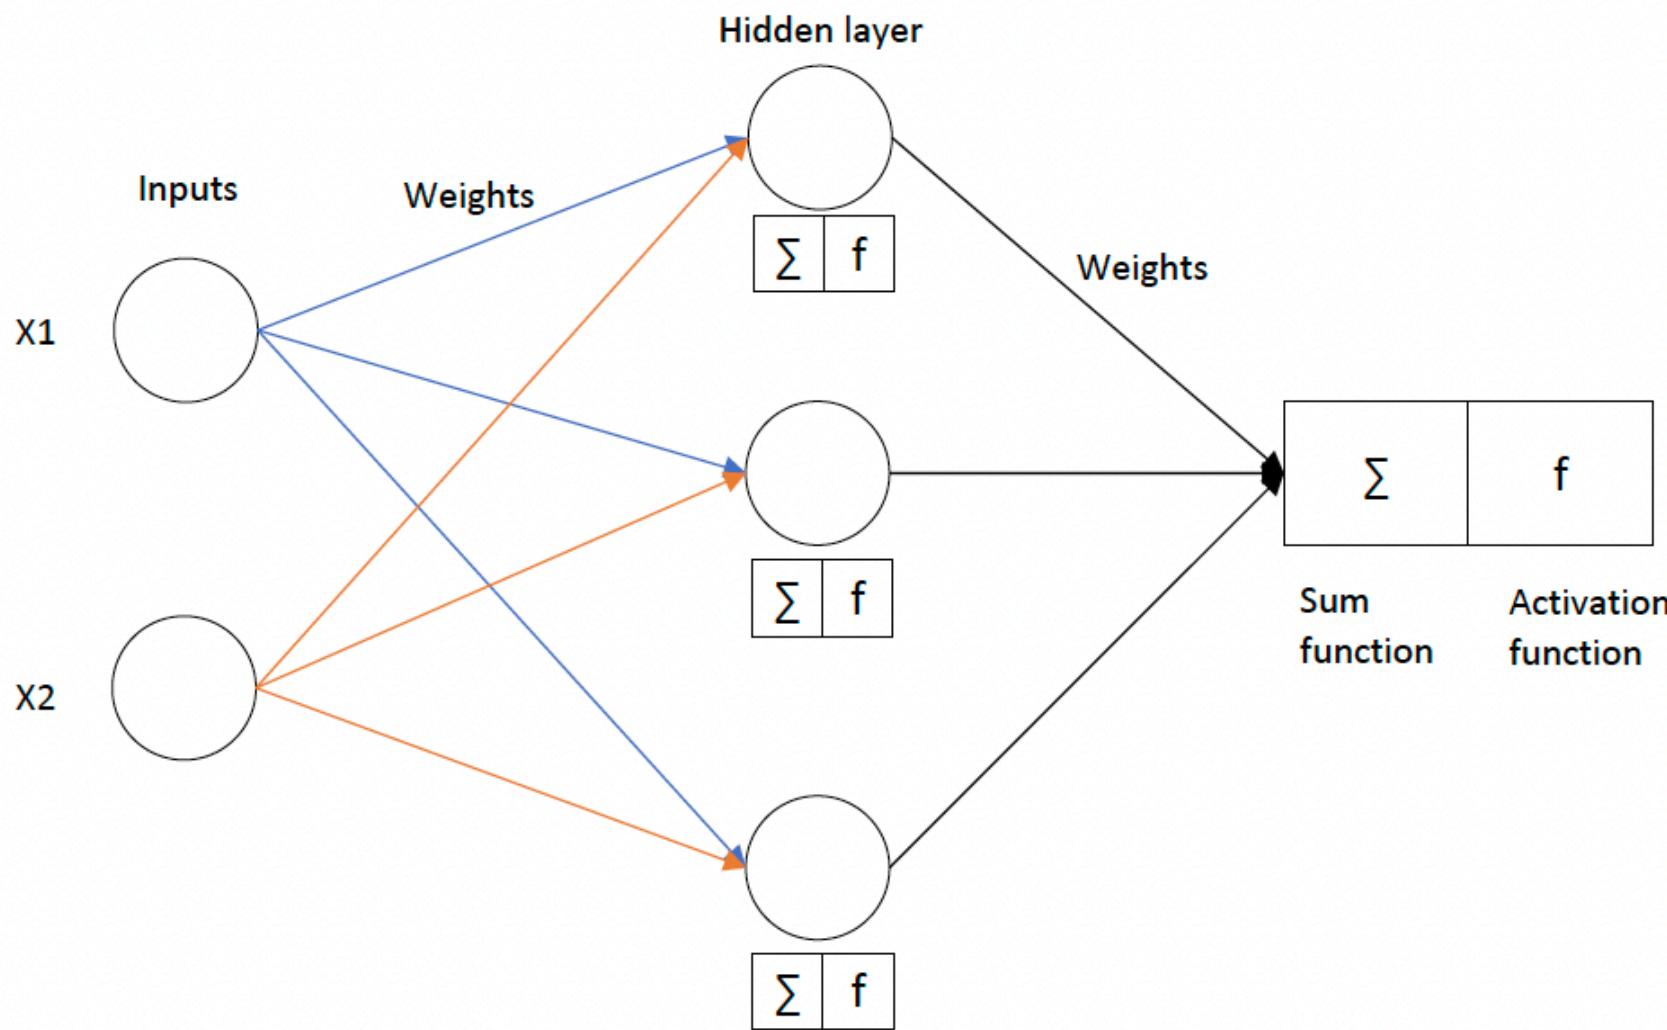
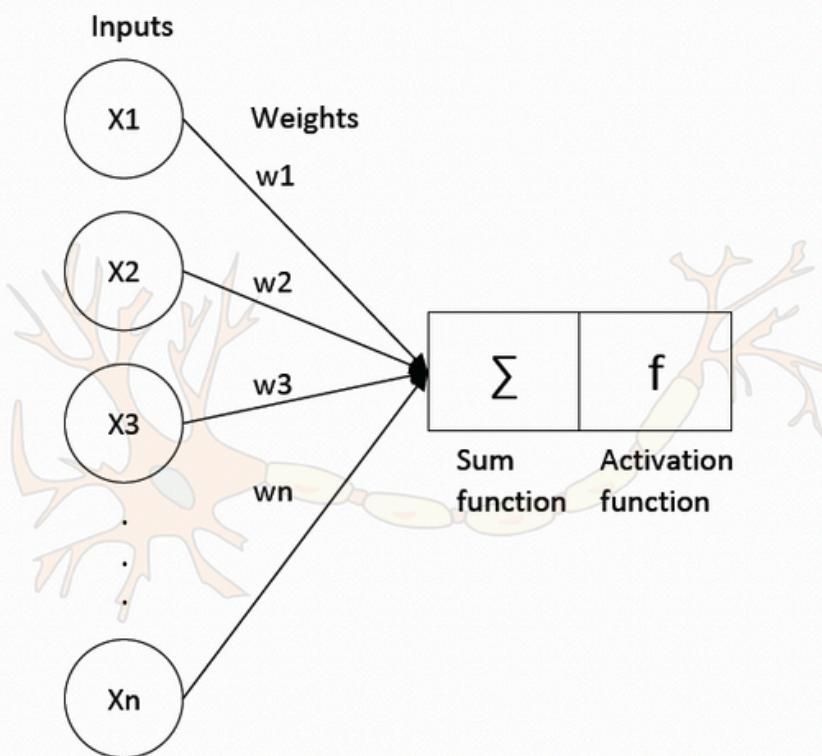
```
[ ] 1 # Create the linear regression model
[ ] 2 model = nn.Linear(1, 1)

[ ] 1 # Loss and optimizer
[ ] 2 criterion = nn.MSELoss()
[ ] 3 optimizer = torch.optim.SGD(model.parameters(), lr=0.05)

[ ] 1 # In ML we want our data to be of shape:
[ ] 2 # (num_samples x num_dimensions)
[ ] 3 X = X.reshape(N, 1)
[ ] 4 Y = Y.reshape(N, 1)
[ ] 5
[ ] 6 # PyTorch uses float32 by default
[ ] 7 # Numpy creates float64 by default
[ ] 8 inputs = torch.from_numpy(X.astype(np.float32))
[ ] 9 targets = torch.from_numpy(Y.astype(np.float32))
```





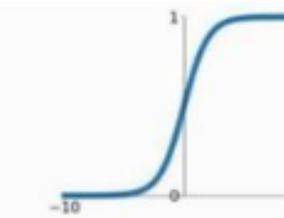


$$sum = \sum_{i=1}^n x_i * w_i$$

$$X_1 * w_1 + X_2 * w_2 + X_3 * w_3$$

Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



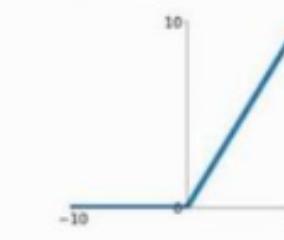
tanh

$$\tanh(x)$$



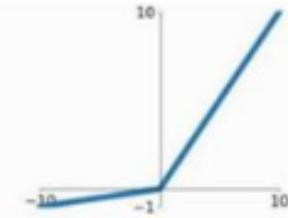
ReLU

$$\max(0, x)$$



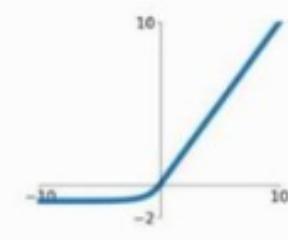
Leaky ReLU

$$\max(0.1x, x)$$



Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$



ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

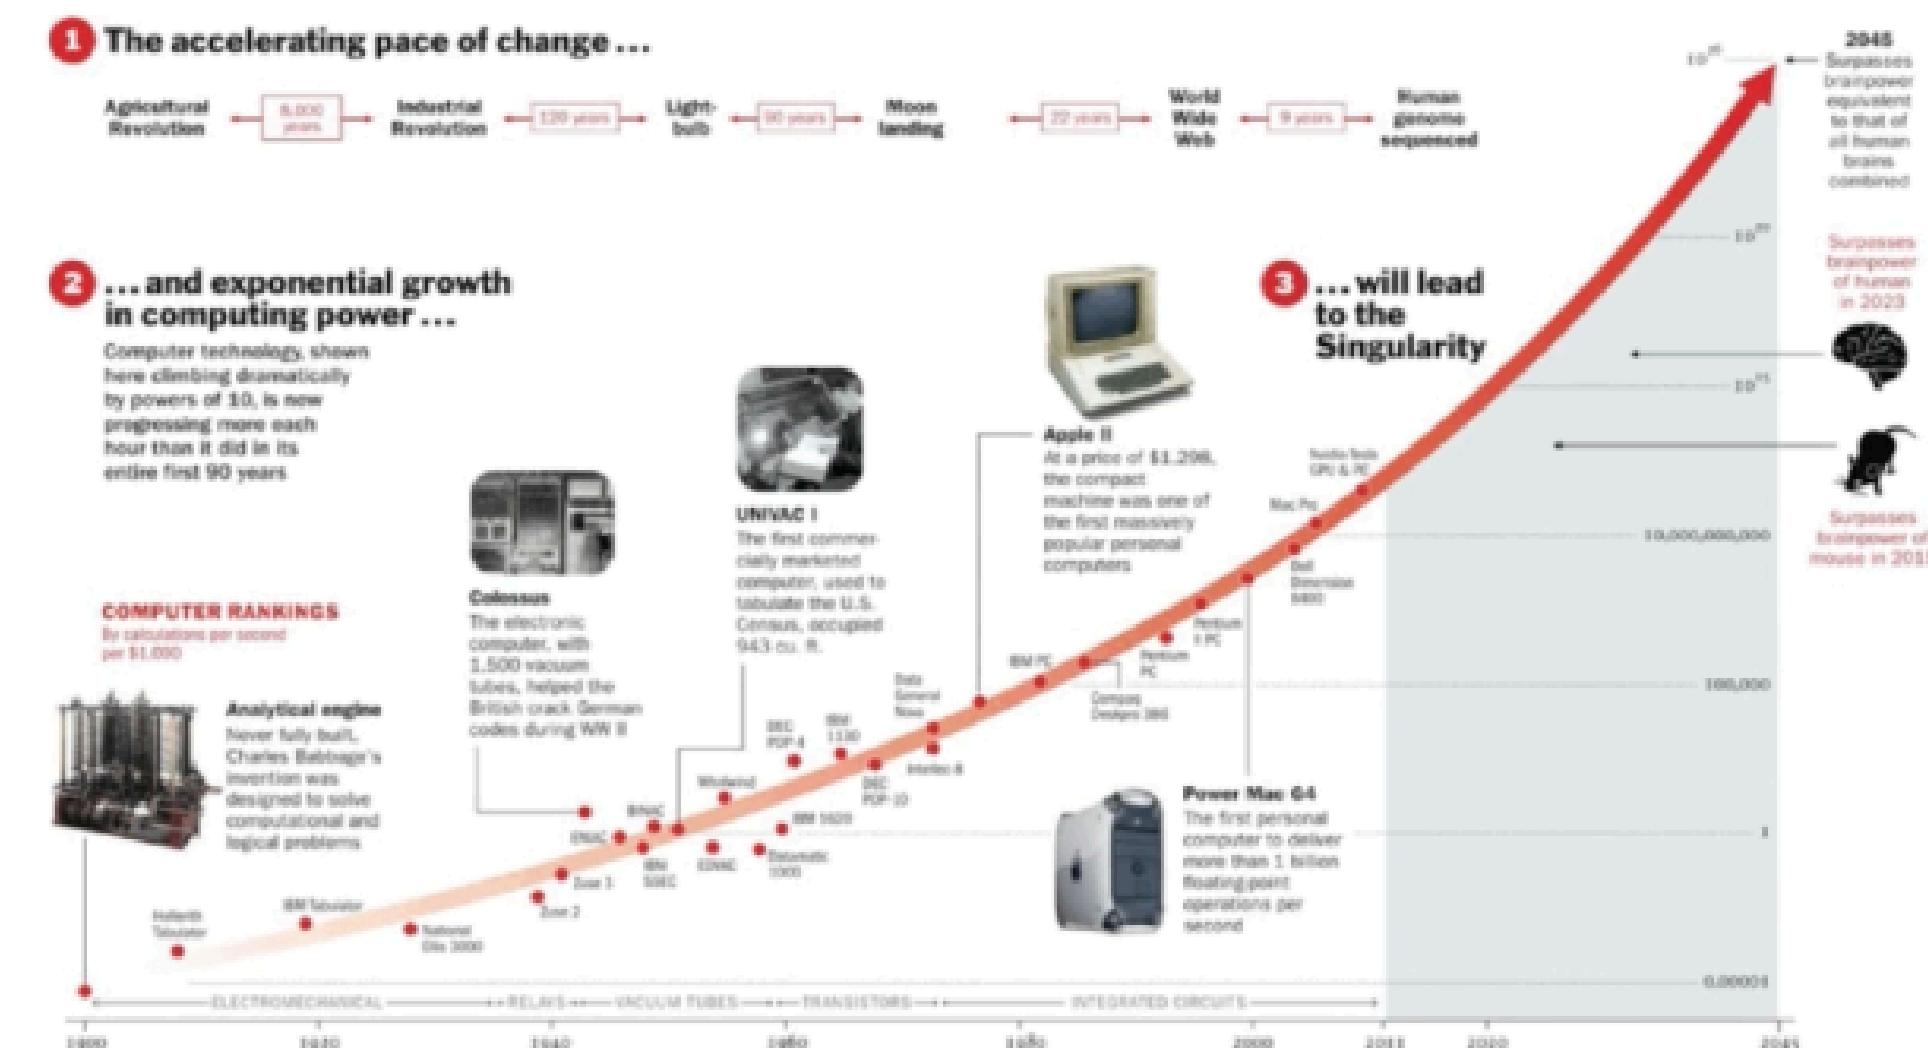
Home Work

Week 1

HOME WORK 2

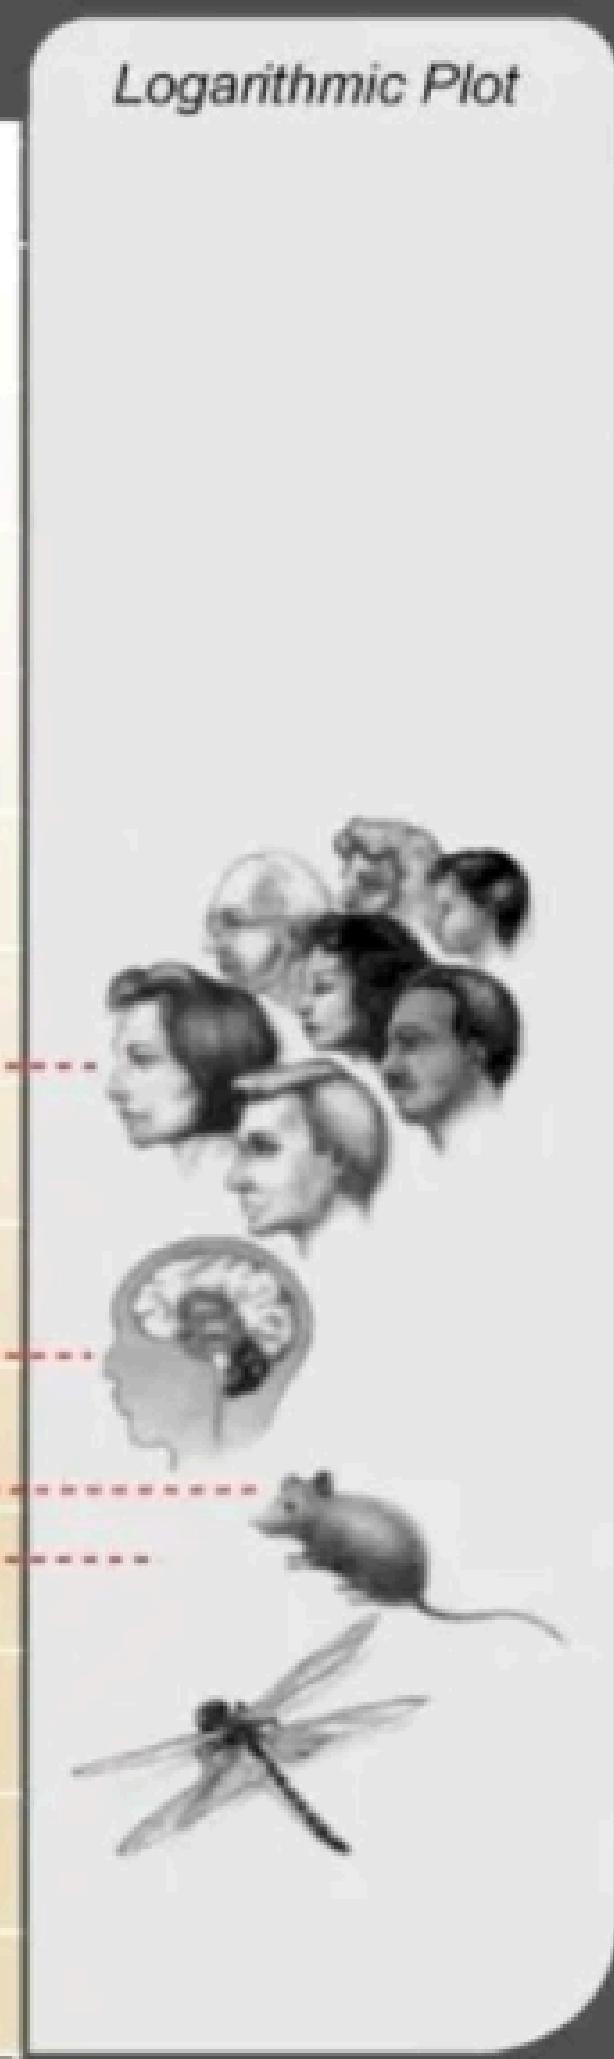
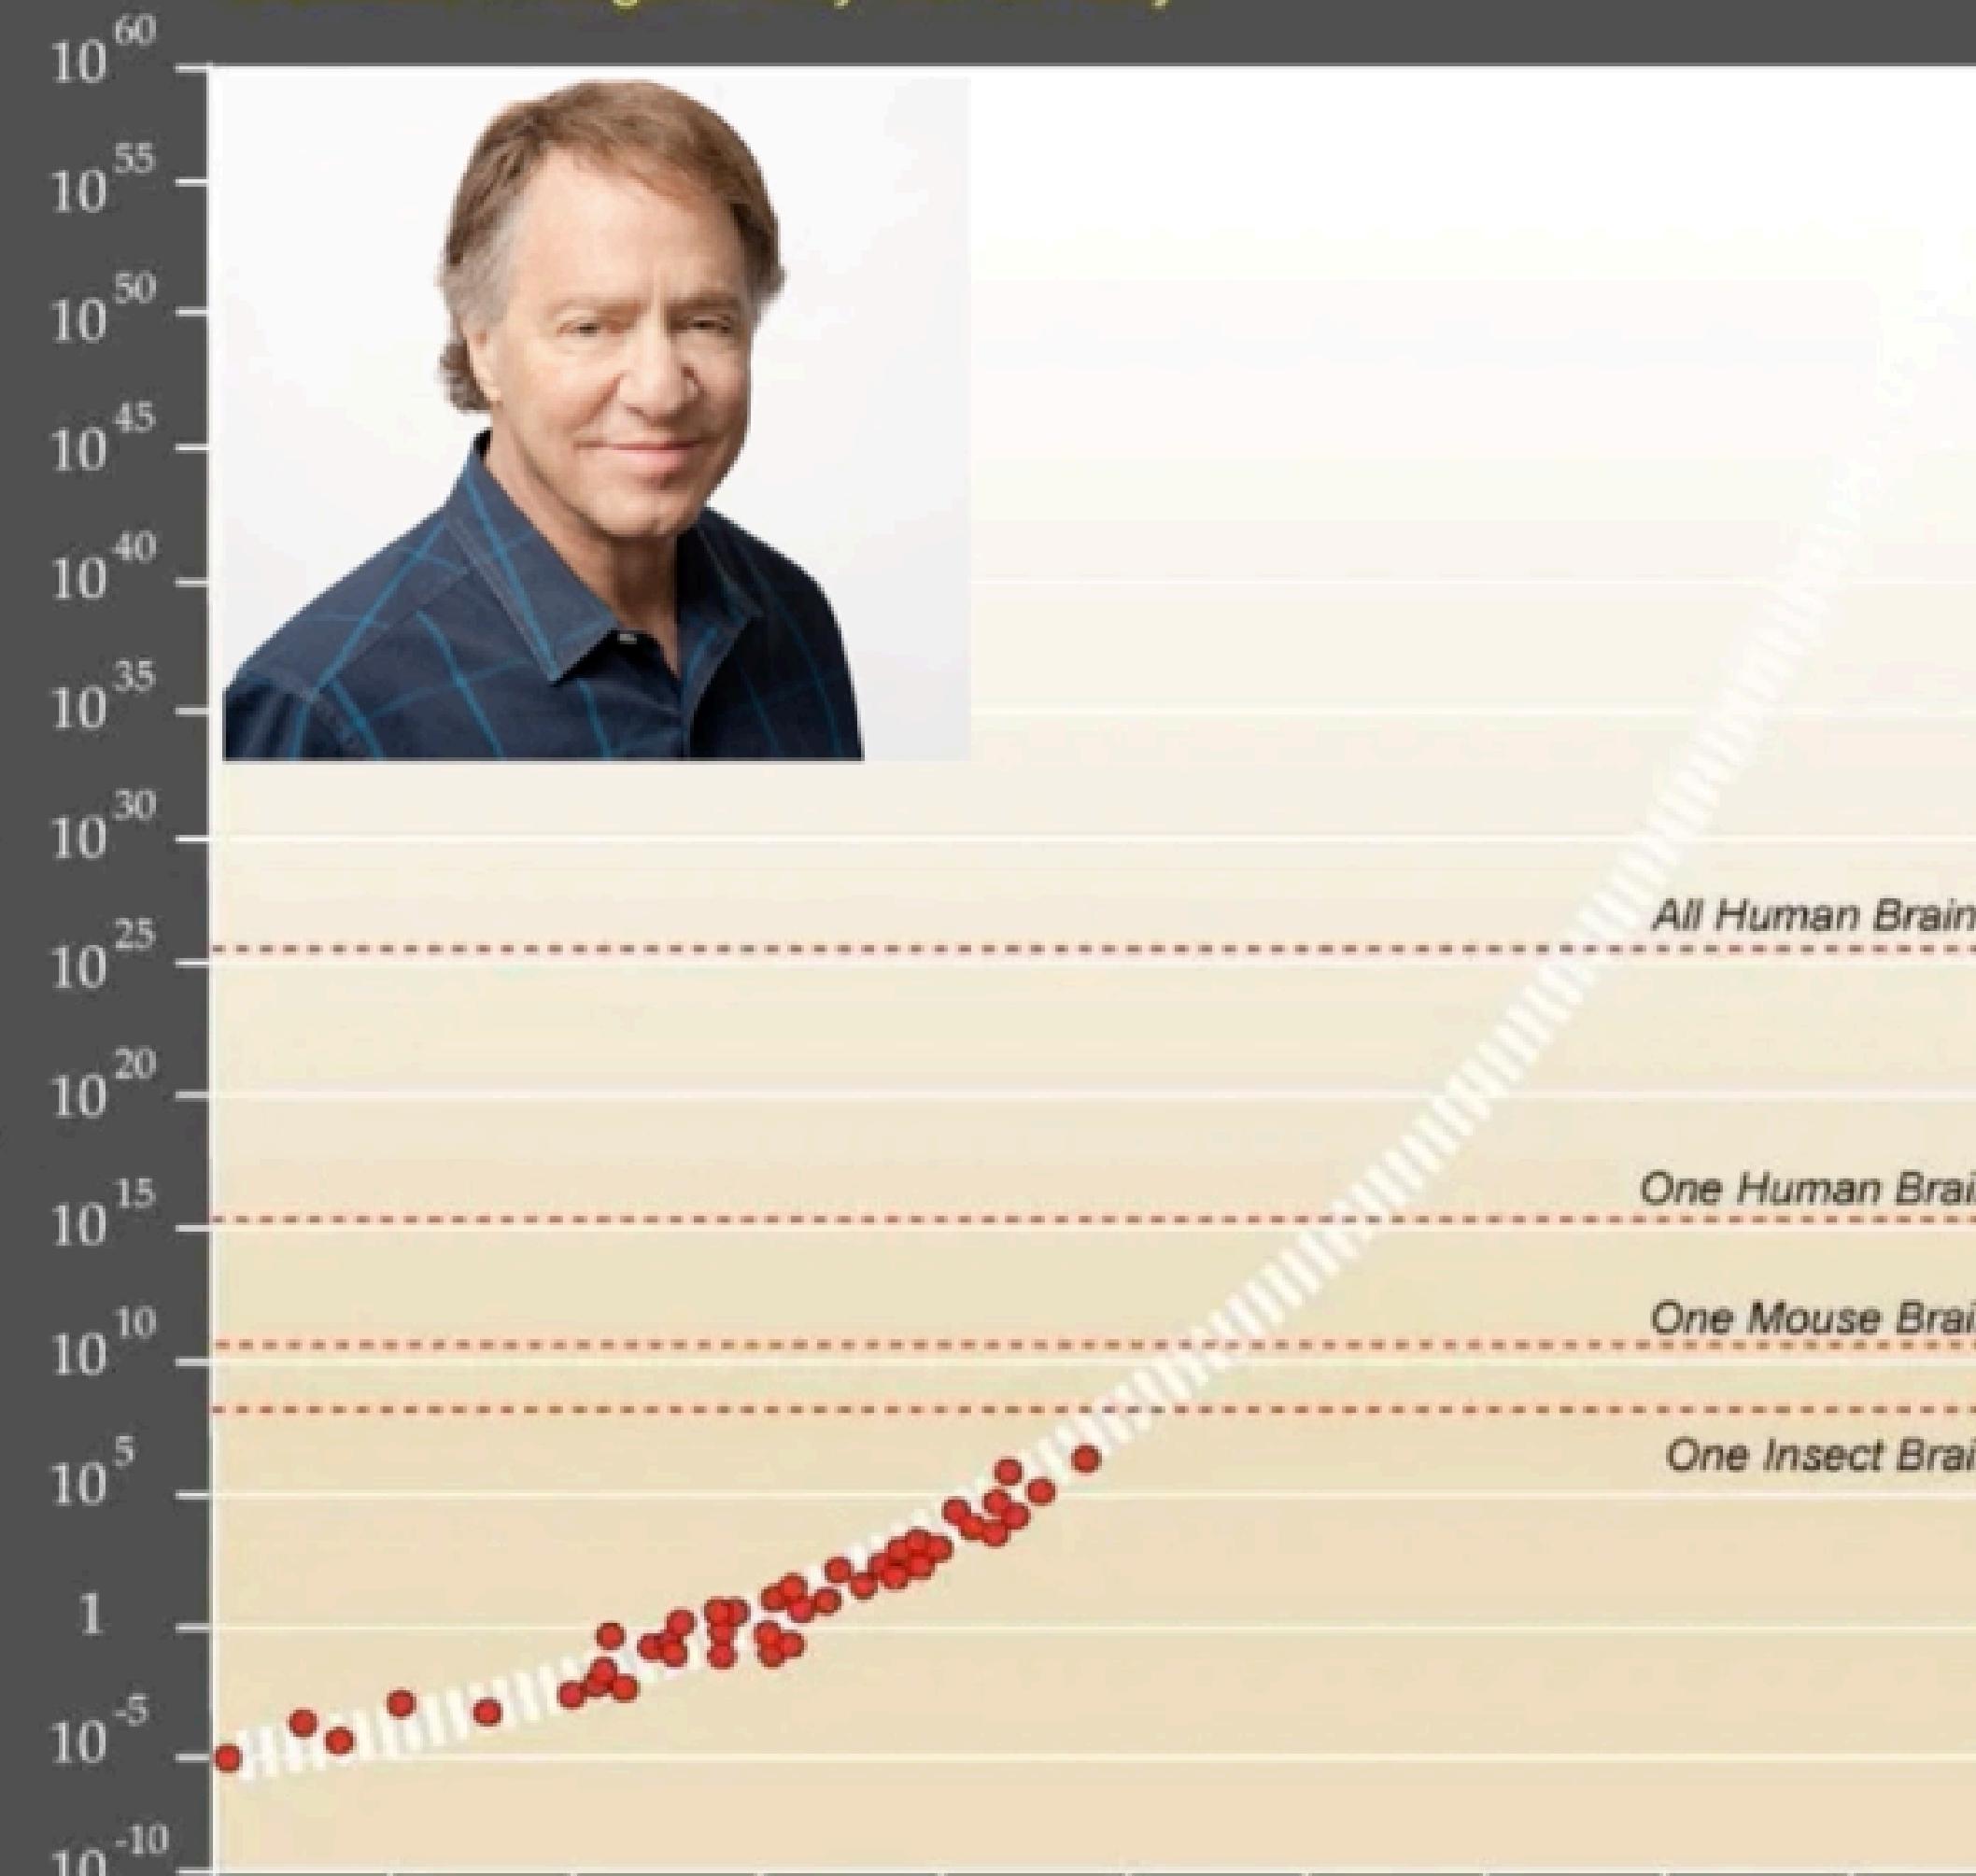
Real-World Dataset: Moore's Law

- The number of transistors per square inch on integrated circuits doubles approximately every 2 years



Exponential Growth of Computing

Twentieth through twenty first century



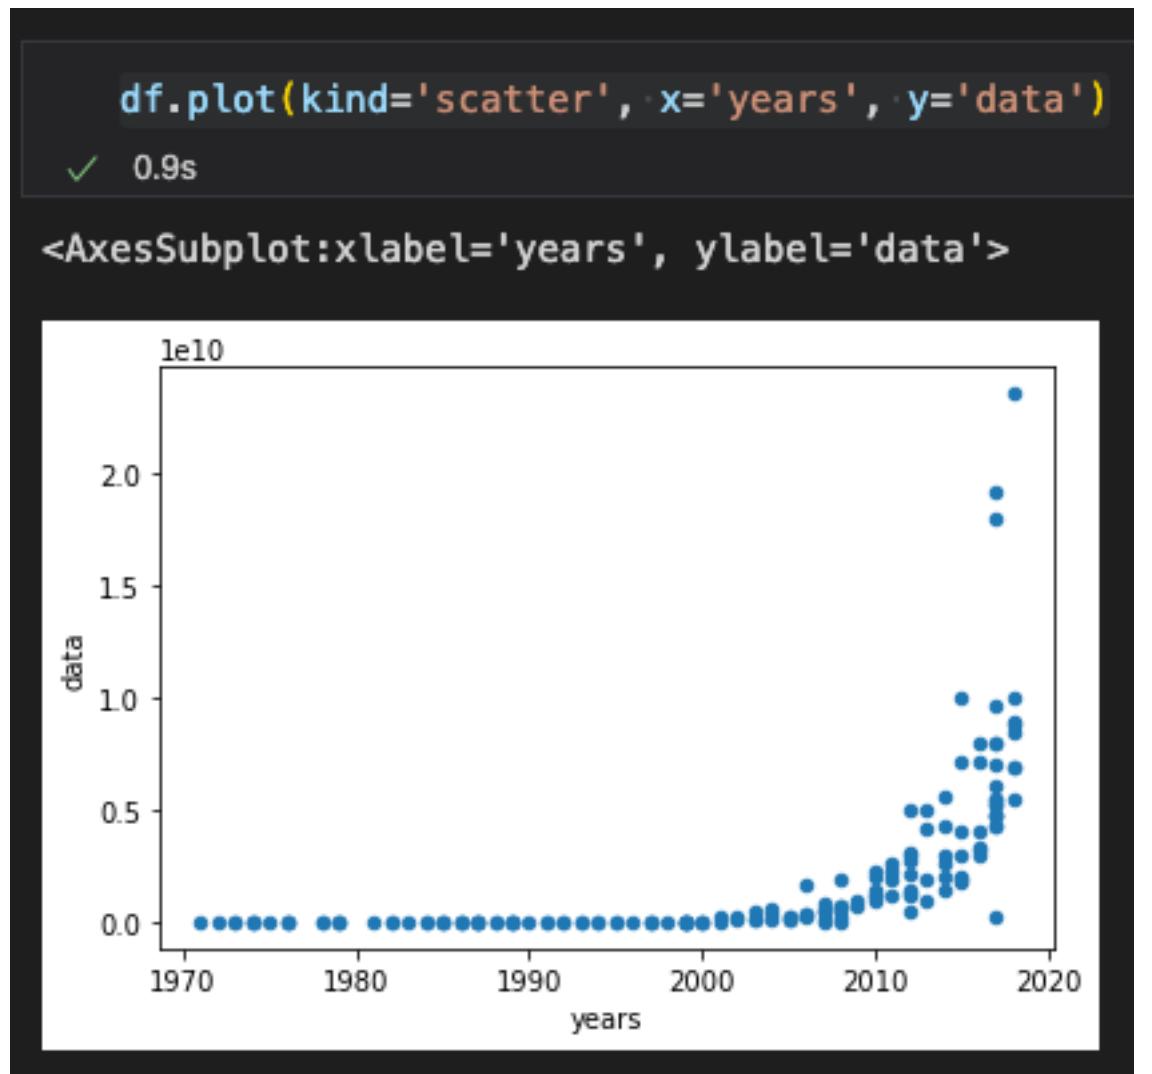
HOME WORK 2

```
df = pd.read_csv('https://raw.githubusercontent.com/lazyprogrammer/machine_learning_examples/master/tf2.0/moore.csv', header=None)
df= df.rename(columns={0: "years", 1: "data"})

df
✓ 0.2s
```

| | years | data |
|-----|-------|--------------|
| 0 | 1971 | 2300 |
| 1 | 1972 | 3500 |
| 2 | 1973 | 2500 |
| 3 | 1973 | 2500 |
| 4 | 1974 | 4100 |
| ... | ... | ... |
| 157 | 2017 | 180000000000 |
| 158 | 2017 | 192000000000 |
| 159 | 2018 | 8876000000 |
| 160 | 2018 | 236000000000 |
| 161 | 2018 | 90000000000 |

162 rows x 2 columns



HOME WORK 2

$$y = A_1 \sin(f_1 t + \theta_1) + b_1 + A_2 \cos(f_2 t + \theta_2) + b_2$$

Parameters = $A_1, f_1, \theta_1, b_1, A_2, f_2, \theta_2, b_2$

```
t = np.arange(0, 10, 0.005)

#random data on x-axis
A1 = 5*np.random.rand()
f1 = 10*np.random.rand()+5
b1 = 20*np.random.rand()

ceta1 = 2*np.random.rand()*np.pi

signal1 = A1 * np.sin(f1 * t + ceta1) + b1

#random data on x-axis
A2 = 5*np.random.rand()
f2 = 10*np.random.rand()+20
b2 = 20*np.random.rand()
ceta2 = 2*np.random.rand()*np.pi

signal2 = A2 * np.cos(f2 * t + ceta2) + b2

y = signal1 + signal2

print(f"A1 = {A1}, f1 = {f1} , ceta1 = {ceta1} b1= {b1}")
print(f"A2 = {A2}, f1 = {f2} , ceta2 = {ceta2} b2= {b2}")

✓ 0.0s

A1 = 4.757400431132537, f1 = 14.982025297569715 , ceta1 = 4.532944447976331 b1= 4.804226285370636
A2 = 1.448664560946444, f1 = 24.567685969748283 , ceta2 = 2.764982500149891 b2= 19.881588370782723
```

