

DEEP LEARNING FOR COMPUTER VISION

Week1



Dr. Tuchsanai. PloySuwan

Course Learning Outcomes (CLOs)

Upon completion of this course, students are expected to be able to

| CLO ID | CLO Description |
|--------|--|
| CLO-1 | ผู้เรียนเข้าใจหลักการพื้นฐานของงาน Deep learning สำหรับงาน Computer Vision |
| CLO-2 | ผู้เรียนเข้าใจหลักการพื้นฐานในการใช้ Convolutional neural network (CNN) สำหรับงาน Image Classifications |
| CLO-3 | ผู้เรียนสามารถอิมเพลเม้นต์และเลือกใช้ Modern deep learning architectures ที่เหมาะสมสำหรับงาน Computer Vision ชนิดต่าง เช่น VGG, ResNet, and Efficientnet |
| CLO-4 | ผู้เรียนเข้าใจหลักการการเลือกใช้งาน Loss function, activation functions และ stochastic optimization ได้อย่างเหมาะสม |
| CLO-5 | ผู้เรียนเข้าใจหลักการในการใช้ Transfer learning และการ fine tuning ที่เหมาะสม |
| CLO-6 | ผู้เรียนเข้าใจหลักการพื้นฐานในการใช้ Deep learning สำหรับการทำ Image Segmentation ในงาน Computer Vision |
| CLO-7 | ผู้เรียนสามารถเข้าใจหลักการของ Object detection and Object Tracking ในงาน Computer Vision |
| CLO-8 | ผู้เรียนสามารถเข้าใจหลักการของ Generative Adversarial Network |

| ลำดับที่ | หัวข้อ | วิธีการเรียน | ผลการเรียนรู้ | หมายเหตุ |
|----------|--|--------------|---------------------|-----------|
| 1 | The basic concept of deep learning and their applications in computer vision | บรรยาย | CLO-1 | Chapter 1 |
| 2 | Introduction to Images and Videos Basics with Python and OpenCV | บรรยาย | CLO-1 | Chapter 1 |
| 3 | Advance in OpenCV | บรรยาย | CLO-1 | Chapter 1 |
| 4 | Introduction to Pytorch | บรรยาย | CLO-2, CLO-3, CLO-4 | Chapter 2 |
| 5 | Fundamental Neural network | บรรยาย | CLO-2 | Chapter 2 |
| 6 | Stochastic optimization methods | บรรยาย | CLO-4 | Chapter 2 |
| 7 | Convolutional neural network for image classification | บรรยาย | CLO-2, CLO-3 | Chapter 3 |
| 9 | Modern architectures such as VGG, ResNet, and Efficientnet | บรรยาย | CLO-3, | Chapter 3 |
| 10 | Transfer learning and fine tuning | บรรยาย | CLO-5 | Chapter 3 |
| 11 | Deep learning for Segmentation for images | บรรยาย | CLO-6 | Chapter 4 |
| 12 | Deep learning Object detection and Object Tracking | บรรยาย | CLO-7 | Chapter 4 |
| 13 | GAN architecture (Generative Adversarial Network) for photo-realistic images | บรรยาย | CLO-8 | Chapter 5 |
| 14 | Generative image from text CLIP model | บรรยาย | CLO-8 | Chapter 5 |
| 15 | Discussions | บรรยาย | | |

APPROACHING (ALMOST) ANY MACHINE LEARNING PROBLEM



ABHISHEK THAKUR

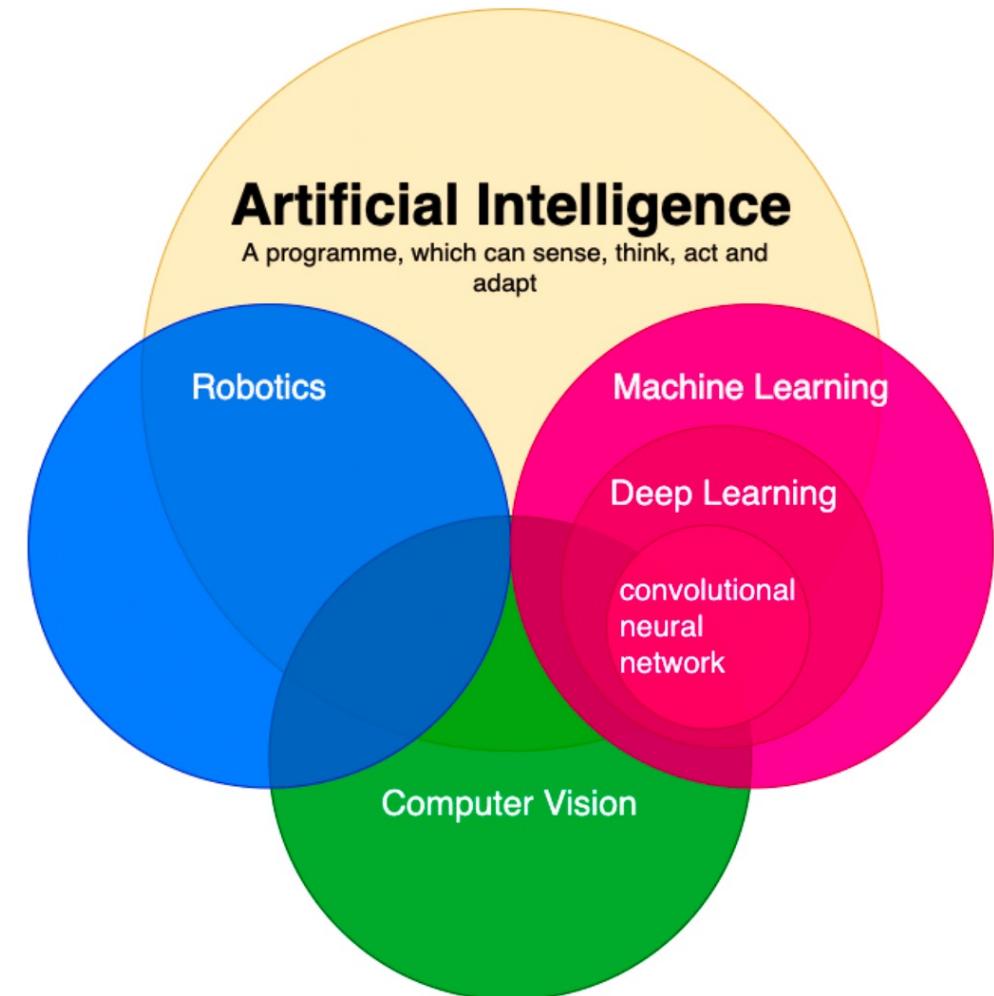
BOOKS

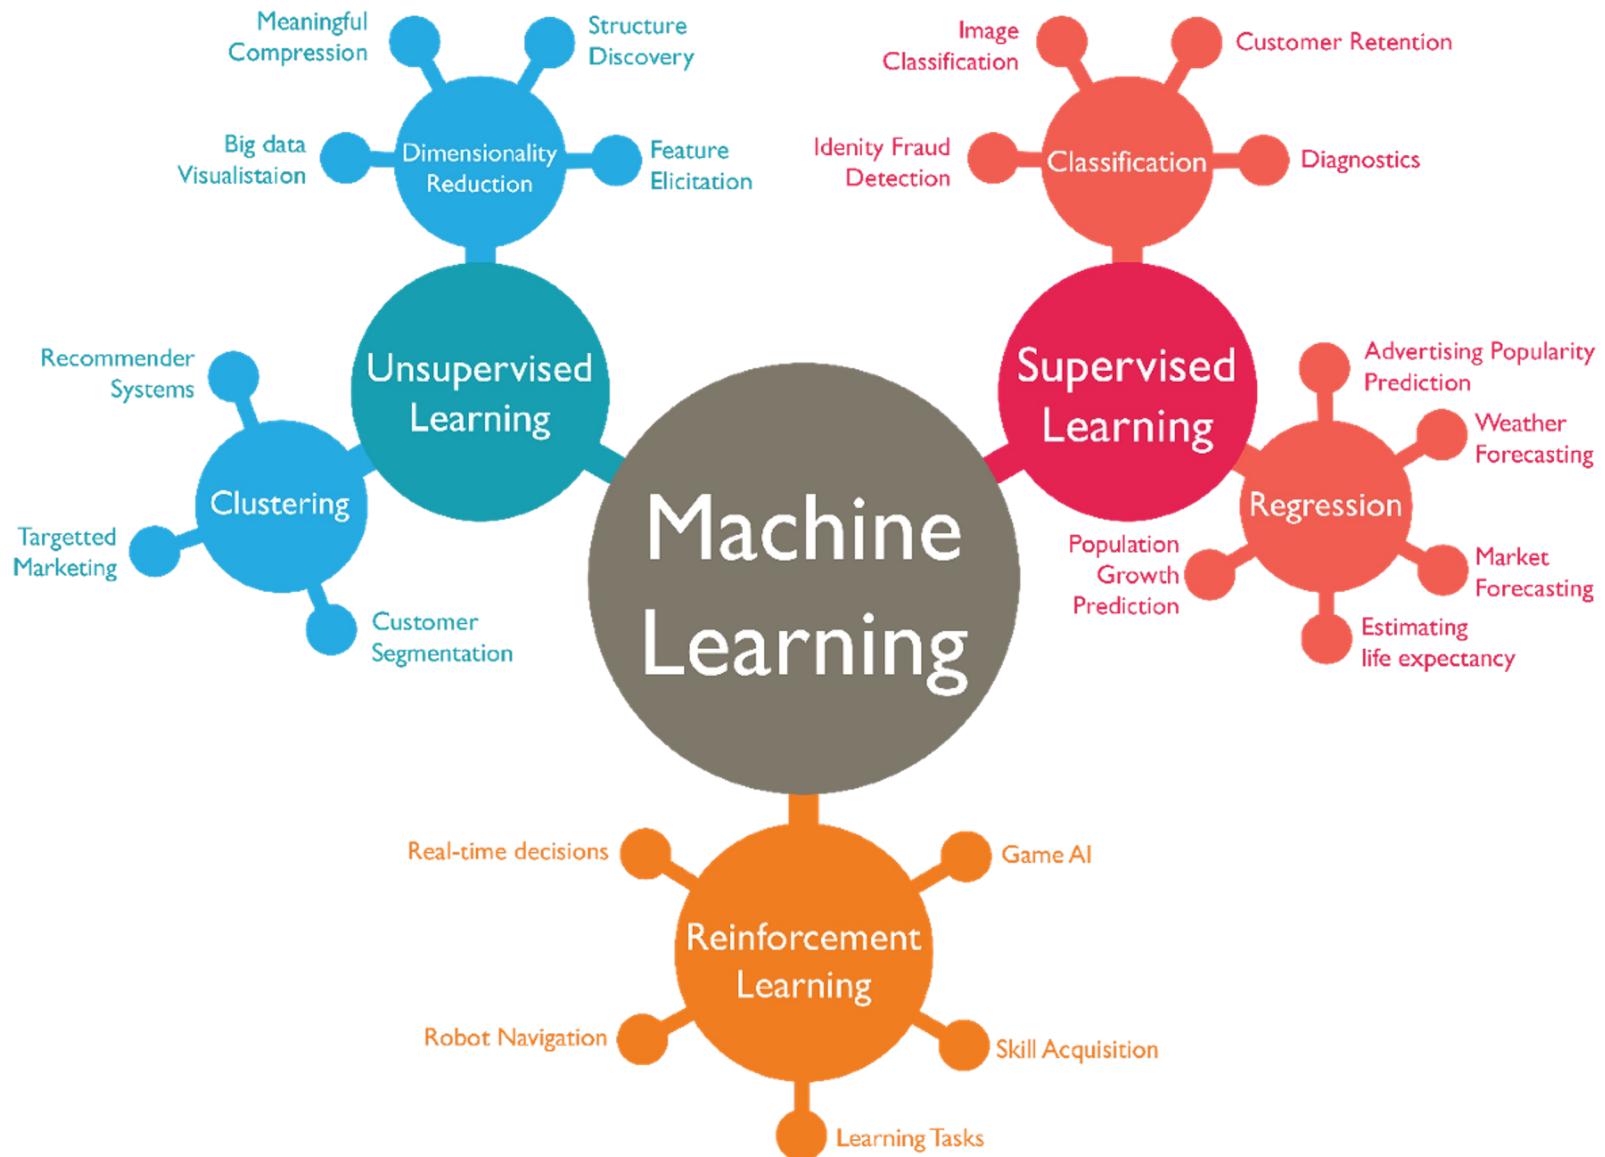
<https://github.com/abhishekrthakur/approachingalmost>

Is Computer Vision Artificial Intelligence?

- Relation between Artificial Intelligence, Machine Learning and Deep Learning, Computer Vision.

https://www.researchgate.net/figure/Relation-between-Artificial-Intelligence-Machine-Learning-and-Deep-Learning-Computer_fig1_342978934

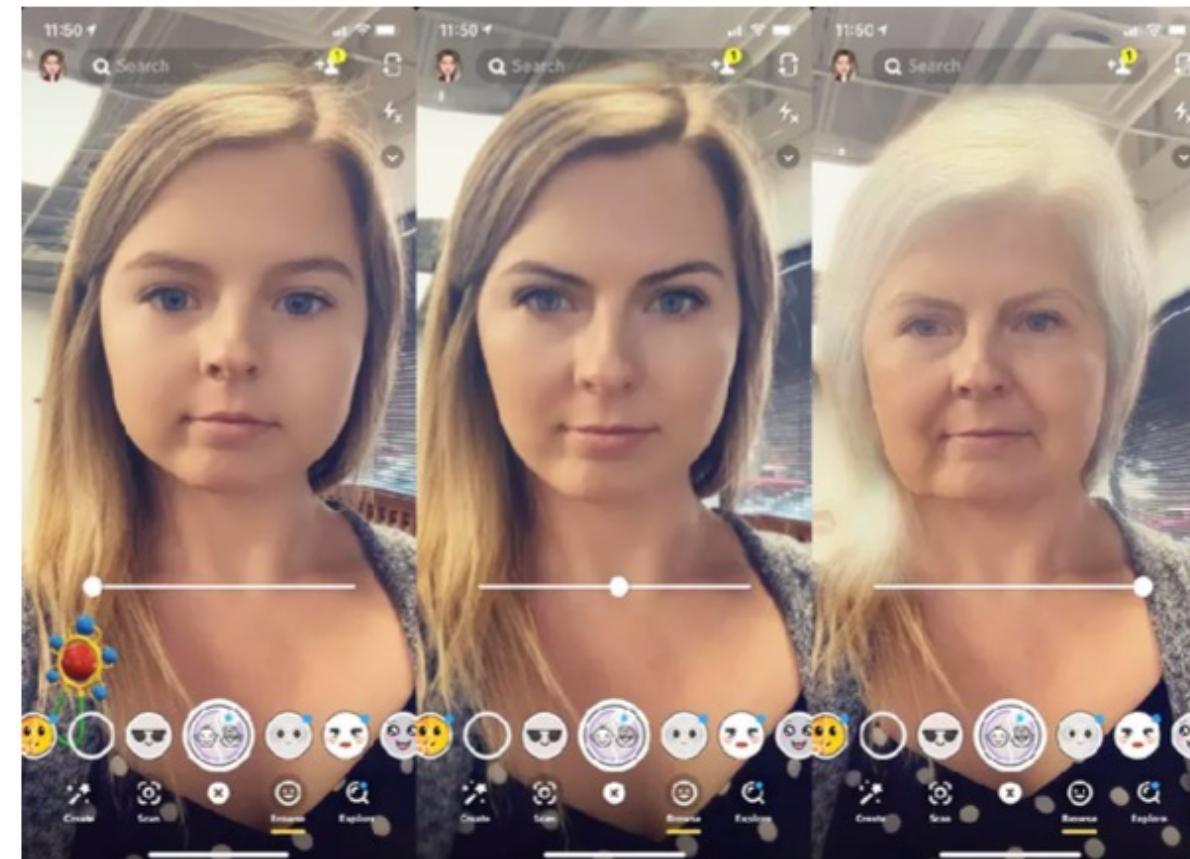




So what can Computer Vision do?

You might be familiar with these...

- Snapchat and Instagram filters
- Optical Character Recognition (OCR)
- Licence Plate Reading
- Self-driving cars
- Sporting Analysis
- Facial Recognition

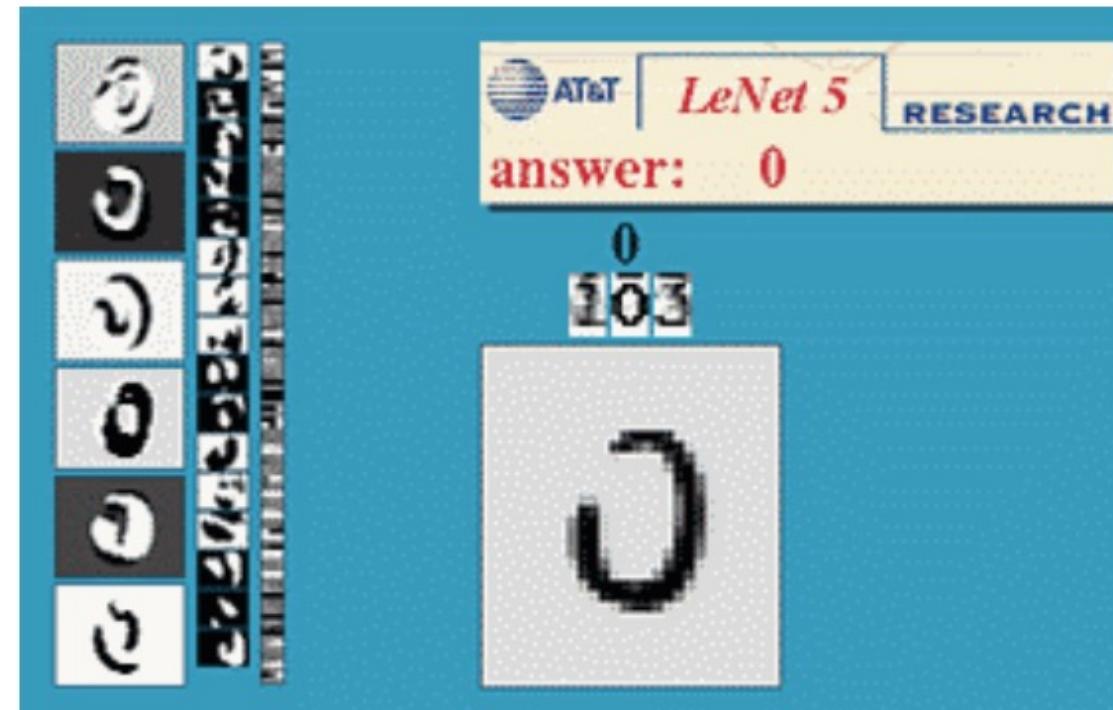


Source - Cnet - Snapchat's Time Machine AR lens creepily shows what you'll look like old

So what can Computer Vision do?

You might be familiar with these...

- Snapchat and Instagram filters
- Optical Character Recognition (OCR)
- Licence Plate Reading
- Self-driving cars
- Sporting Analysis
- Facial Recognition



Source -AT&T's LeNet OCR for Handwritten Digits

SO WHAT CAN COMPUTER VISION DO?

- Snapchat and Instagram filters
- Optical Character Recognition (OCR)
- **Licence Plate Reading**
- Self-driving cars
- Sporting Analysis
- Facial Recognition

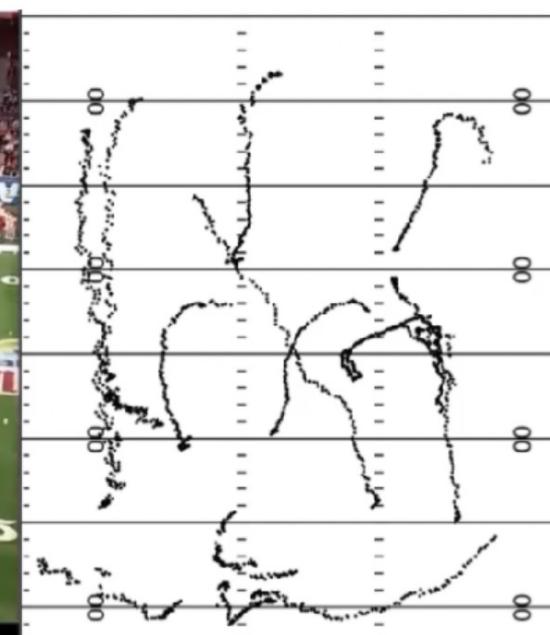
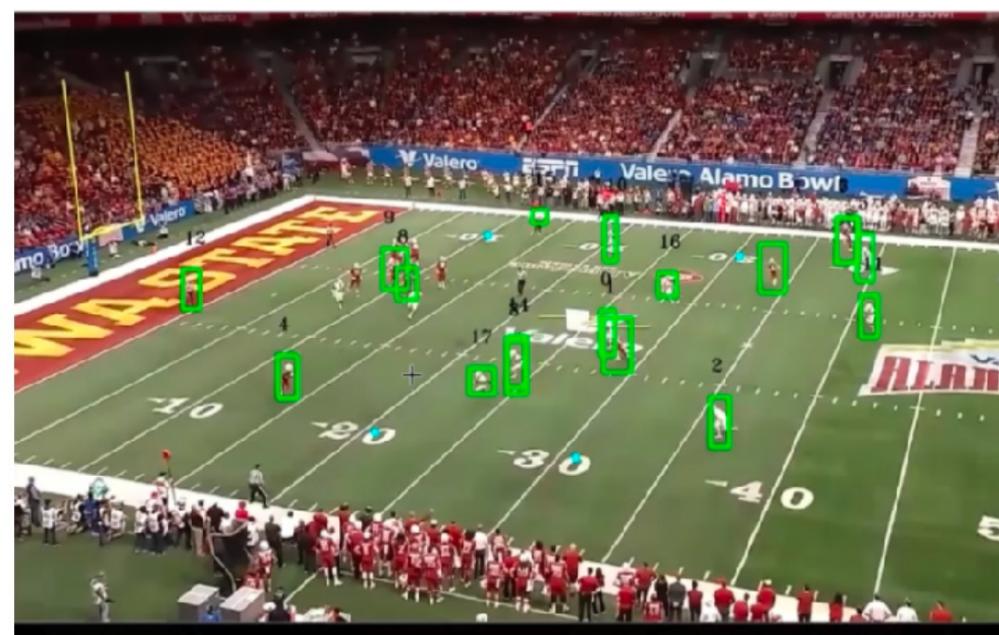
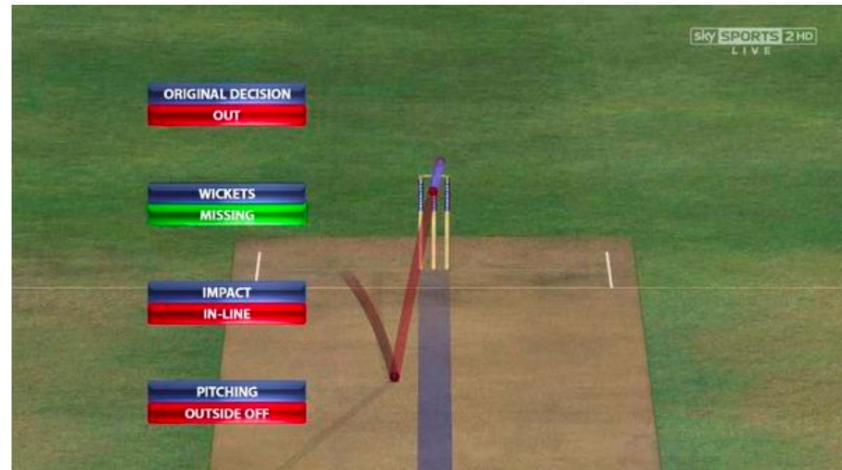


So what can Computer Vision do?

You might be familiar with these...

- Snapchat and Instagram filters
- Optical Character Recognition (OCR)
- Licence Plate Reading
- Self-driving cars
- **Sporting Analysis**
- Facial Recognition

HawkEye in Cricket



Source -Roboflow - AI Coach

Classical Computer Vision?

- What is meant by **Classical Computer Vision?**
- It encompasses Computer Vision algorithms that **do not** involve Machine Learning
- Before the advent of Machine Learning and Deep Learning, Computer Vision was a deeply explored field and many useful algorithms were developed for things like **feature extraction, OCR, Segmentation and simple transformations.**
- **OpenCV** is the Classical Computer Vision library of choice!



Deep Learning Computer Vision

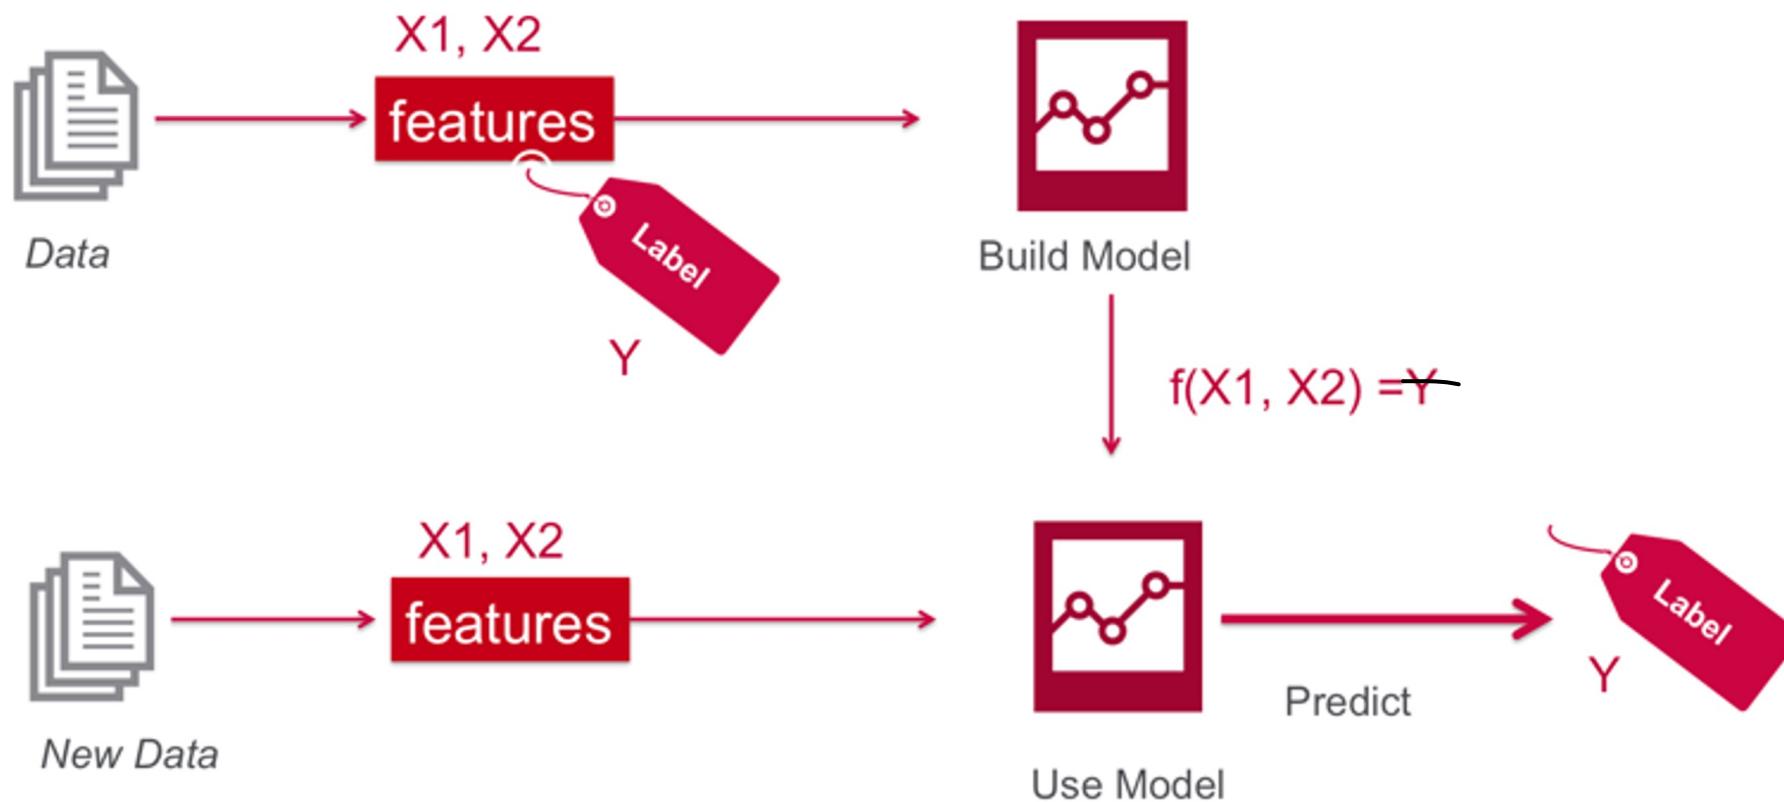
- Deep Learning was used in Computer Vision since the 1990s, however due to the computational requirements and intricate design, it remained on the sidelines for decades.
- Until the mid 2010s...which brought **two important building blocks** together.
 - **Accessible GPU processing** (NVIDIA's CUDA)



Deep Learning CV vs Classical CV

| Deep Learning | Classical Computer Vision |
|---|---|
| Adapts to new images well (assuming it's similar to the data it was trained on) | Small changes can have big negative impacts |
| Requires Models to be trained | Doesn't require training and can be used once coded |
| Model weights learn to adapt to varying image conditions | Relies on hardcode features and parameters |
| Requires GPU hardware (most times) | Can be run on CPU |

Supervised Learning



Classification v.s. Regression

Classification

- Predict which category an item belongs to based on labeled examples of known items
- Use in classify a new data

Regression

- Predict a continuous value target.
- Linear regression predicts a numeric value
- Logistic regression predicts a probability

Examples of classification

- Credit card fraud detection (fraud, not fraud)
- Credit card application (good credit, bad credit)
- Email spam detection (spam, not spam)
- Text sentiment analysis (happy, not happy)
- Predicting patient risk (high risk patient, low risk patient)
- Classifying a tumor as malignant or not

Examples of regression

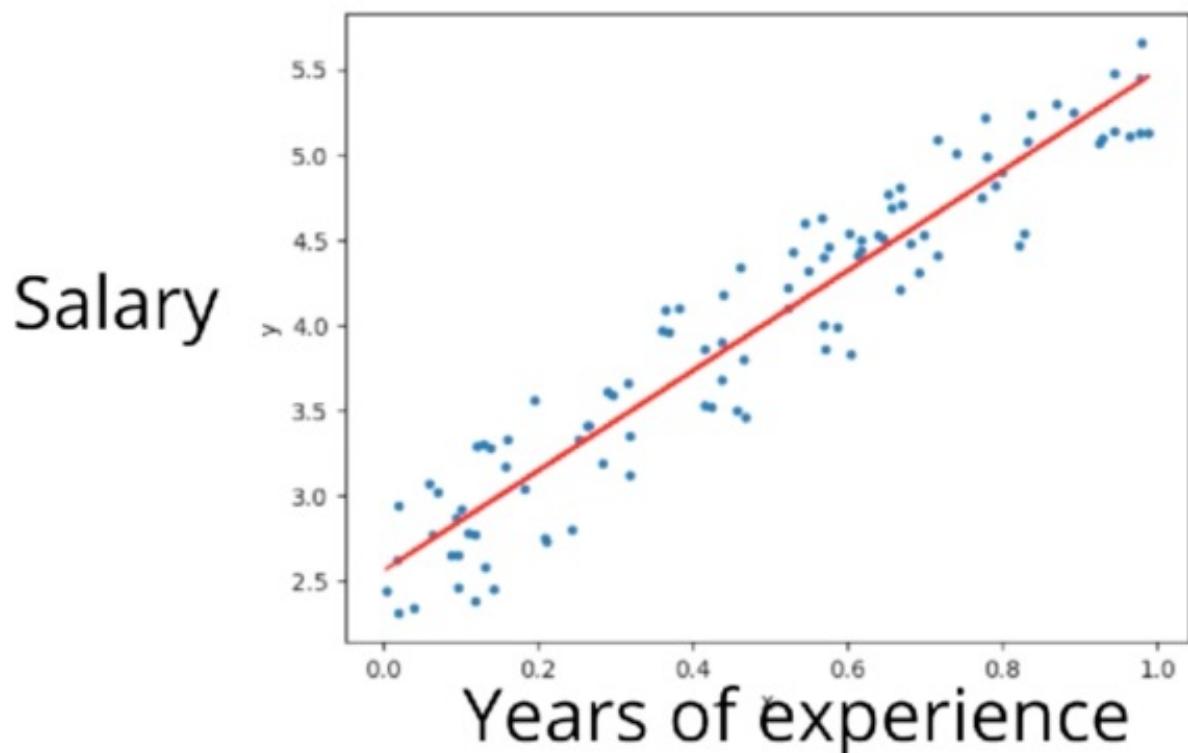
- Predicting the price of the house.
- Predicting age of a person
- Predicting the stock price for tomorrow.
- Predicting next year GDP.

Example: Regression

- Fit a line or curve
- This is why grumpy old statisticians like to say: “machine learning is just glorified curve-fitting”

$$\hat{y} = mx + b$$

(sometimes we also use “a” for the slope)



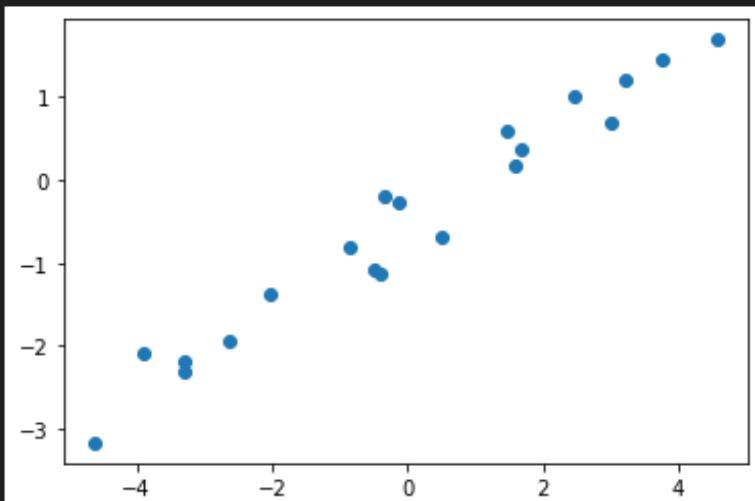
```
import matplotlib.pyplot as plt  
import numpy as np  
✓ 0.4s
```

```
# generate 20 data points  
N = 20  
  
#random data on x-axis  
x= np.random.rand(N)*10-5  
  
y = 0.5*x -1+ np.random.rand(N)  
✓ 0.2s
```

```
plt.scatter(x,y)
```

```
✓ 0.1s
```

```
<matplotlib.collections.PathCollection at 0x7fc1b147a410>
```



```
# In ML we want our data to be of shape:  
# (num_samples x num_dimensions)  
X = x.reshape(N, 1)  
Y = y.reshape(N, 1)
```

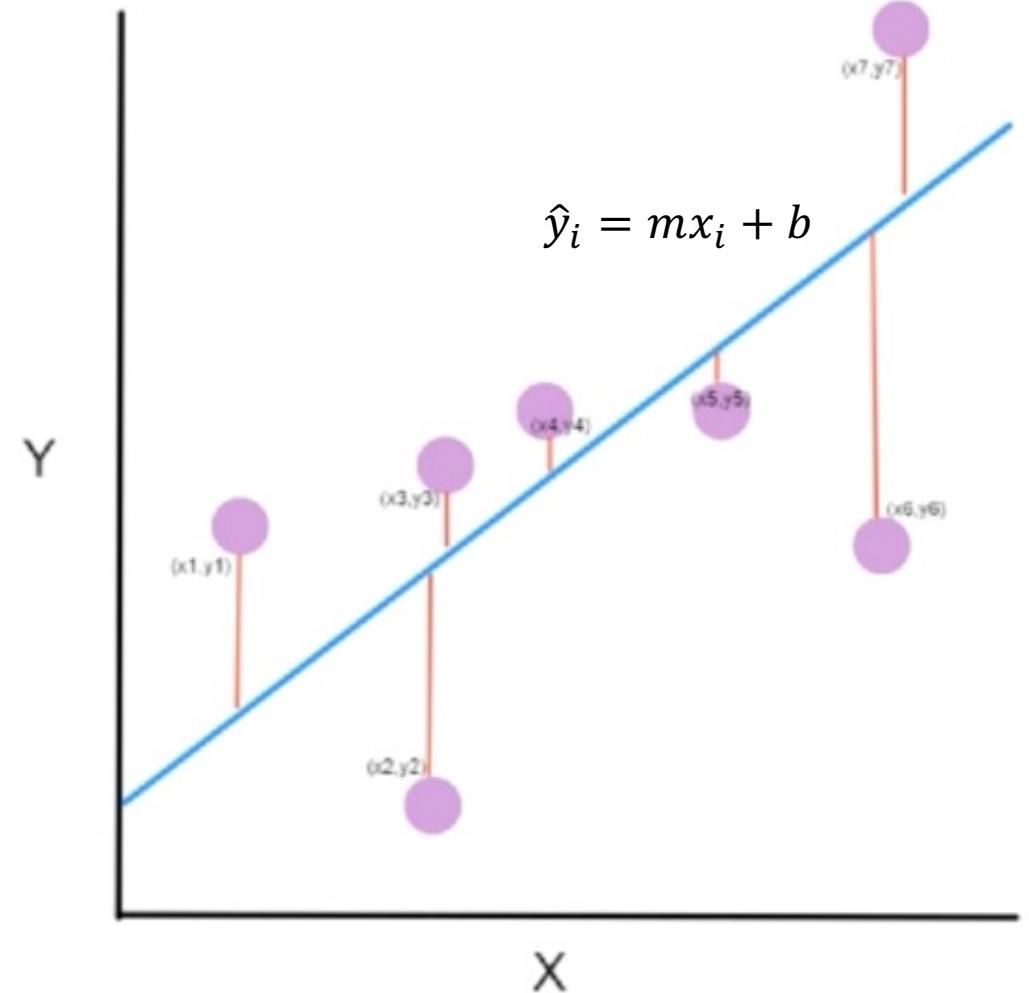
```
✓ 0.2s
```

The Loss

- Data = $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$
- N = number of samples in the dataset
- The line **cannot** perfectly pass through all the data points

MSE = Mean Squared Error

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$



Applying the MSE to find slope / intercept

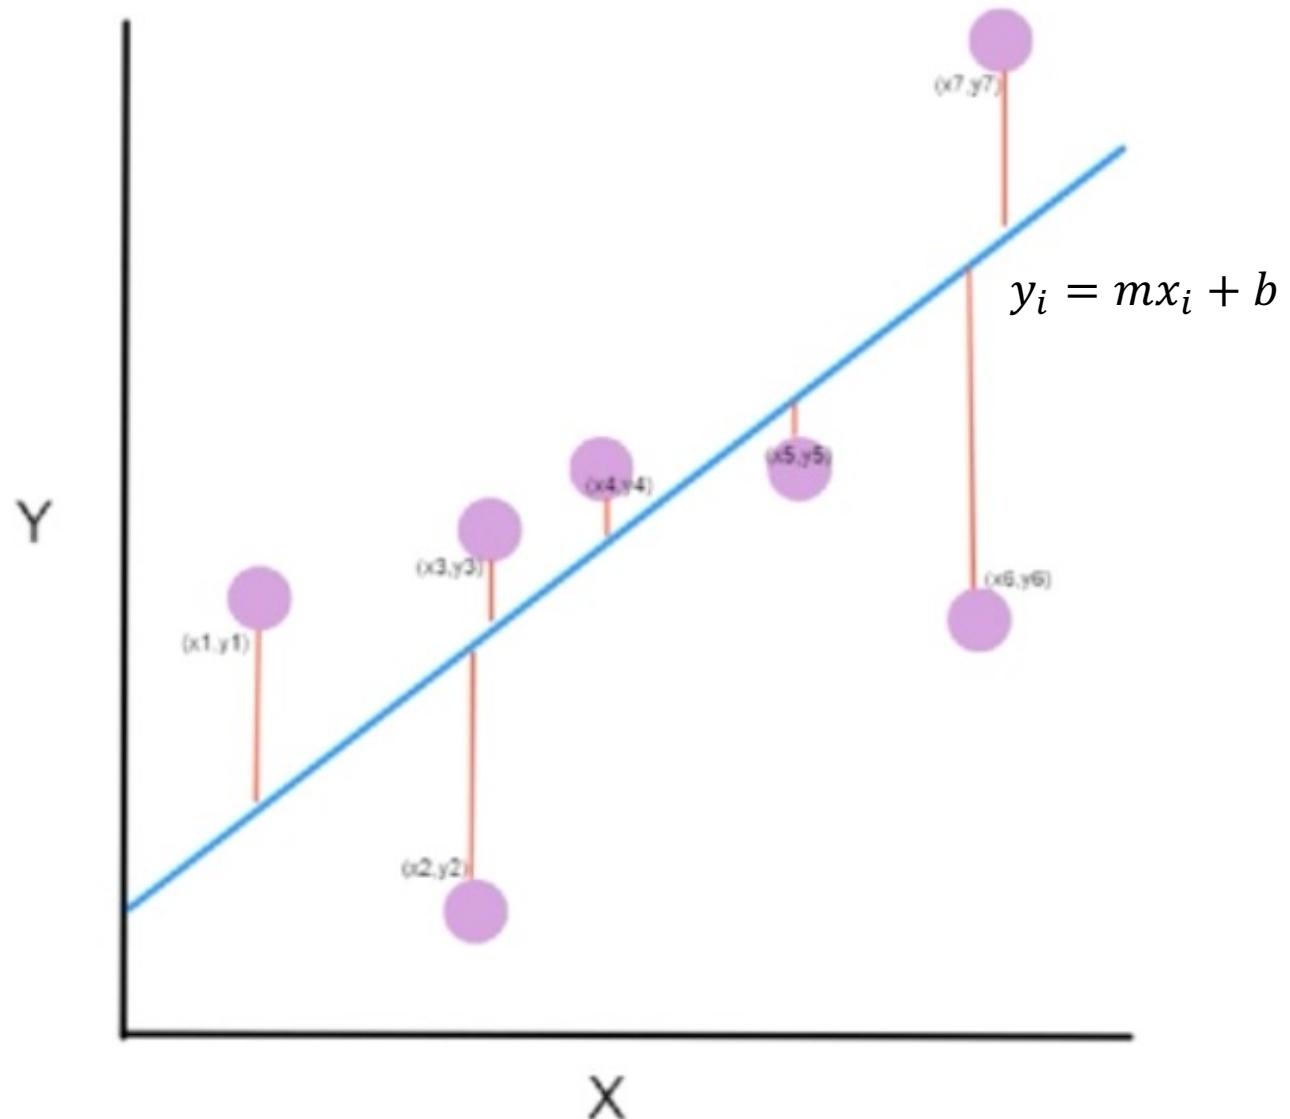
- Plug in the expression for the predictions (\hat{y}_i)
- Quiz: What are the **variables** in this expression?

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

Applying the MSE to find slope / intercept

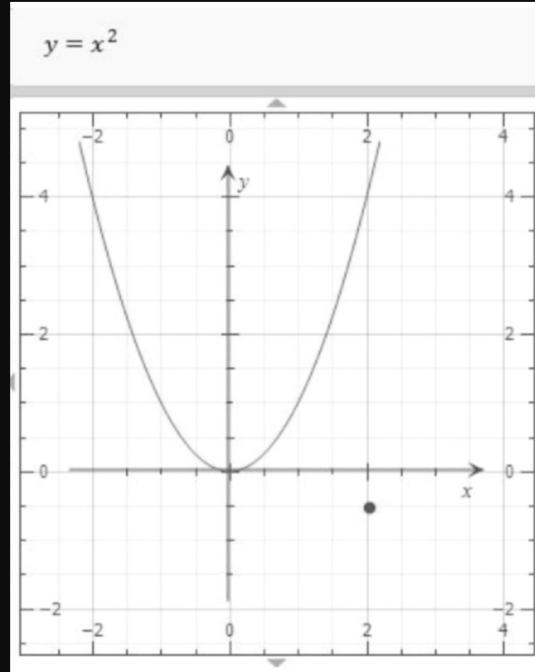
- Note: $L(\text{loss}) = \text{MSE}$
- We want to *minimize the loss* with respect to the parameters (m, b)

$$m^*, b^* = \arg \min_{m, b} L$$



First Solution

Minimize $f(x) = x^2$



- Use calculus!
- $df / dx = 2x = 0$
- Solve for x : $x = 0$

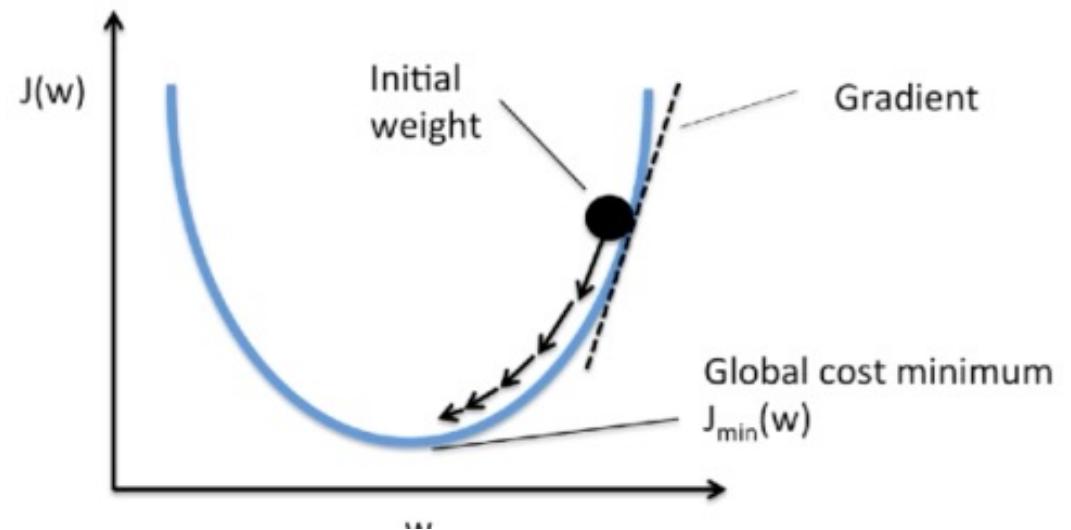
Second Solution : Gradient Descent

- Find the derivatives and set them to 0, solve for the parameters
- Yields 2 equations and 2 unknowns: can solve for (m, b)
- Try it as an exercise!: Should be in terms of $\{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$

$$\partial L/\partial m = 0, \partial L/\partial b = 0$$

Gradient Descent to find hidden parameters

```
# gradient descent  
# pseudocode  
  
 $\theta = (m, b) = \text{random}()$   
for i in range(n_epochs):  
     $\theta = \theta - \eta \nabla_{\theta} L$ 
```



$$L = MSE = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

$$\frac{\partial}{\partial m} = \frac{2}{n} \sum_{i=1}^n -x_i(y_i - (mx_i + b))$$

$$\frac{\partial}{\partial b} = \frac{2}{n} \sum_{i=1}^n -(y_i - (mx_i + b))$$

```
def Gradient(m,b) :  
    grad_m =0; grad_b = 0;  
  
    for i in range(N) :  
        grad_m += -X[i]*(Y[i]- (m*X[i]+b))  
        grad_b += - (Y[i]- (m*X[i]+b) )  
  
    return grad_m, grad_b  
  
Gradient(-1,-1)  
✓ 0.3s  
(array([-209.28382035]), array([-10.12894605]))
```

```
# gradient descent
# pseudocode

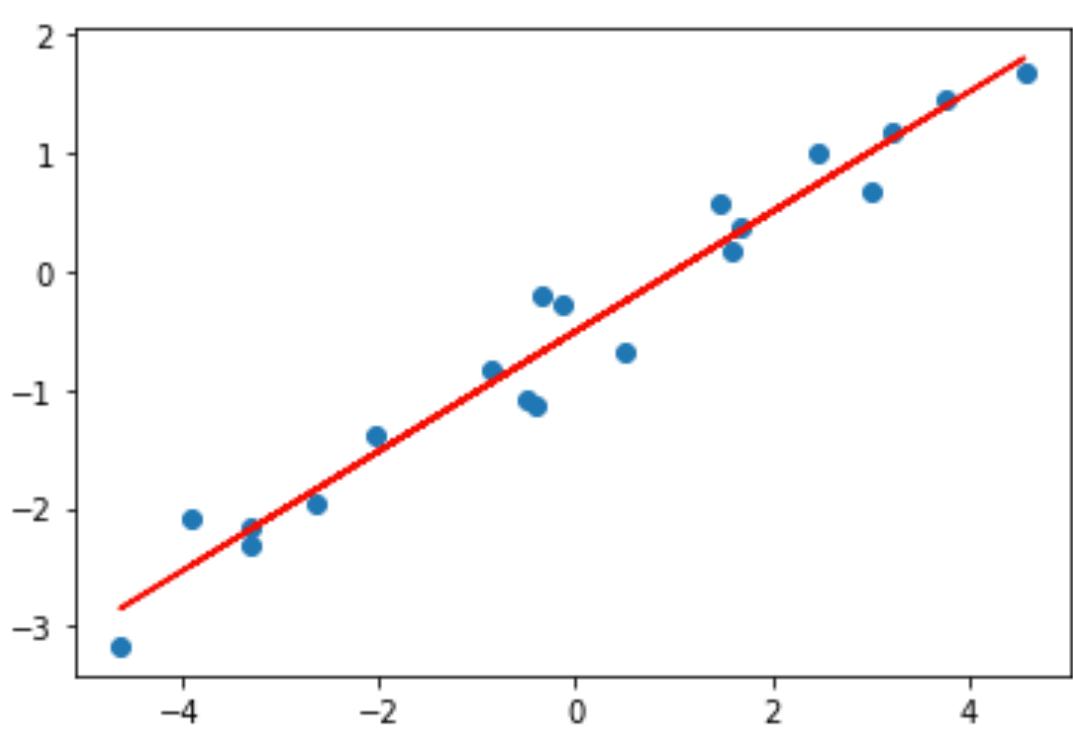
θ = (m, b) = random()
for i in range(n_epochs):
    θ = θ - η∇θL
```

```
n_epochs      = 1000
learning_rate = 0.01
theta         = np.random.rand(2,1)

for i in range(n_epochs):
    grad_m, grad_b = Gradient(theta[0], theta[1])
    theta[0] -= learning_rate*grad_m
    theta[1] -= learning_rate*grad_b

theta
✓ 0.1s
array([[ 0.50714187],
       [-0.51174366]])
```

```
predict = theta[0]*X + theta[1]
plt.plot(X, predict, 'r')
plt.scatter(x,y)
✓ 0.8s
```



SDG: Stochastic gradient descent

$$\theta = \theta - \eta \nabla_{\theta} \mathbb{L}$$

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

```
# Loss and optimizer
criterion = nn.MSELoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

Third Solution :Pytorch

```
# gradient descent
# pseudocode

θ = (m, b) = random()
for i in range(n_epochs):
    θ = θ - η∇_θ L
```

```
# Train the model
n_epochs = 30

for it in range(n_epochs):
    # zero the parameter gradients
    optimizer.zero_grad()

    # Forward pass  $\hat{y}_i = mx_i + b$ 
    outputs = model(inputs)
    loss = criterion(outputs, targets)

    # Backward and optimize
    loss.backward()
    optimizer.step()
```

```

1 # Train the model
2 n_epochs = 30
3 losses = []
4 for it in range(n_epochs):
5     # zero the parameter gradients
6     optimizer.zero_grad()
7
8     # Forward pass
9     outputs = model(inputs)
10    loss = criterion(outputs, targets)
11
12    # keep the loss so we can plot it later
13    losses.append(loss.item())
14
15    # Backward and optimize
16    loss.backward()
17    optimizer.step()
18
19    print(f'Epoch {it+1}/{n_epochs}, Loss: {loss.item():.4f}')

```

Epoch 1/30, Loss: 12.1352
Epoch 2/30, Loss: 2.8066
Epoch 3/30, Loss: 2.4554
Epoch 4/30, Loss: 2.1753
Epoch 5/30, Loss: 1.9483
Epoch 6/30, Loss: 1.7641
Epoch 7/30, Loss: 1.6148

```

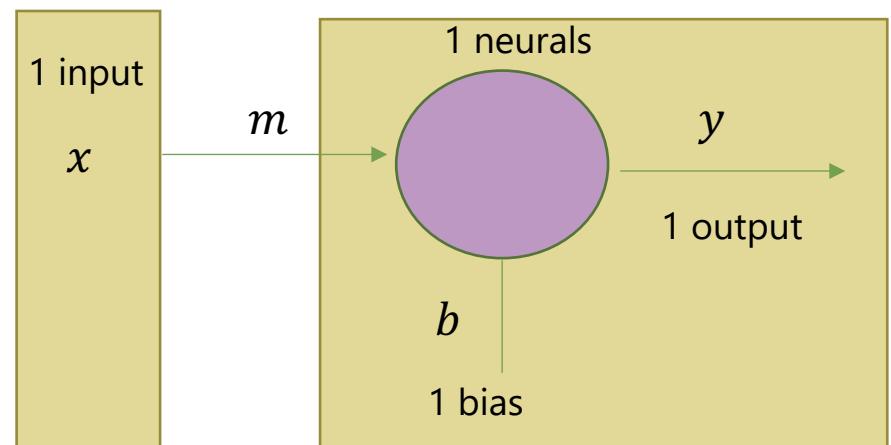
[ ] 1 # Create the linear regression model
2 model = nn.Linear(1, 1)

[ ] 1 # Loss and optimizer
2 criterion = nn.MSELoss()
3 optimizer = torch.optim.SGD(model.parameters(), lr=0.05)

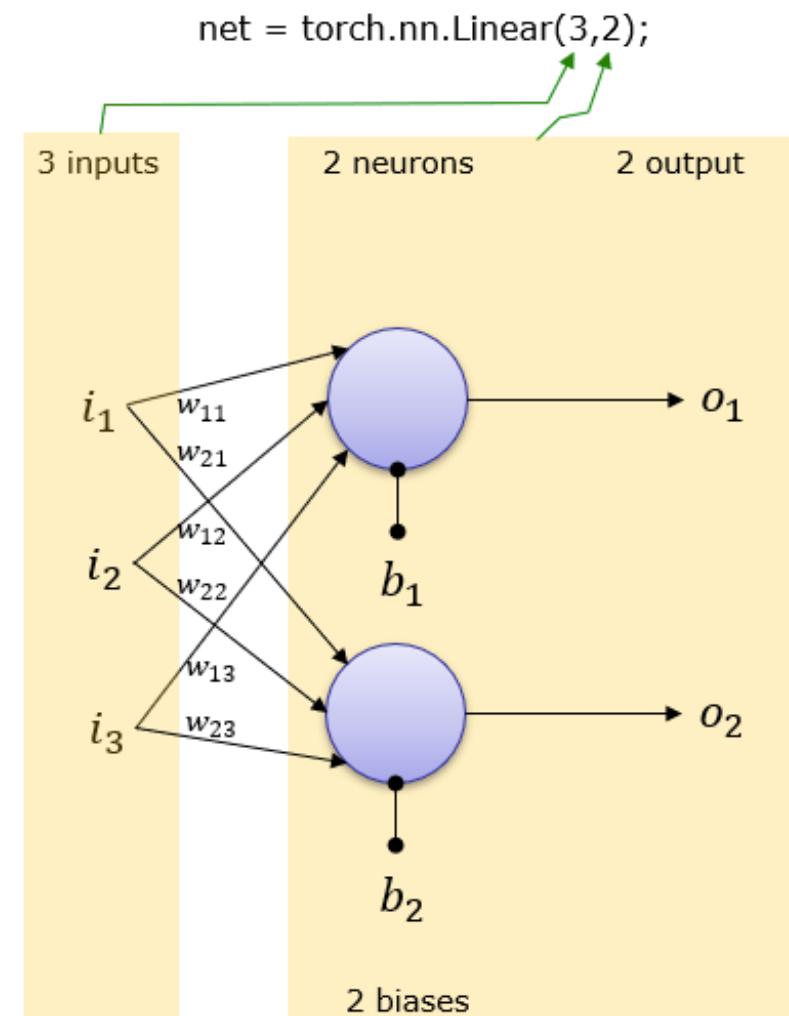
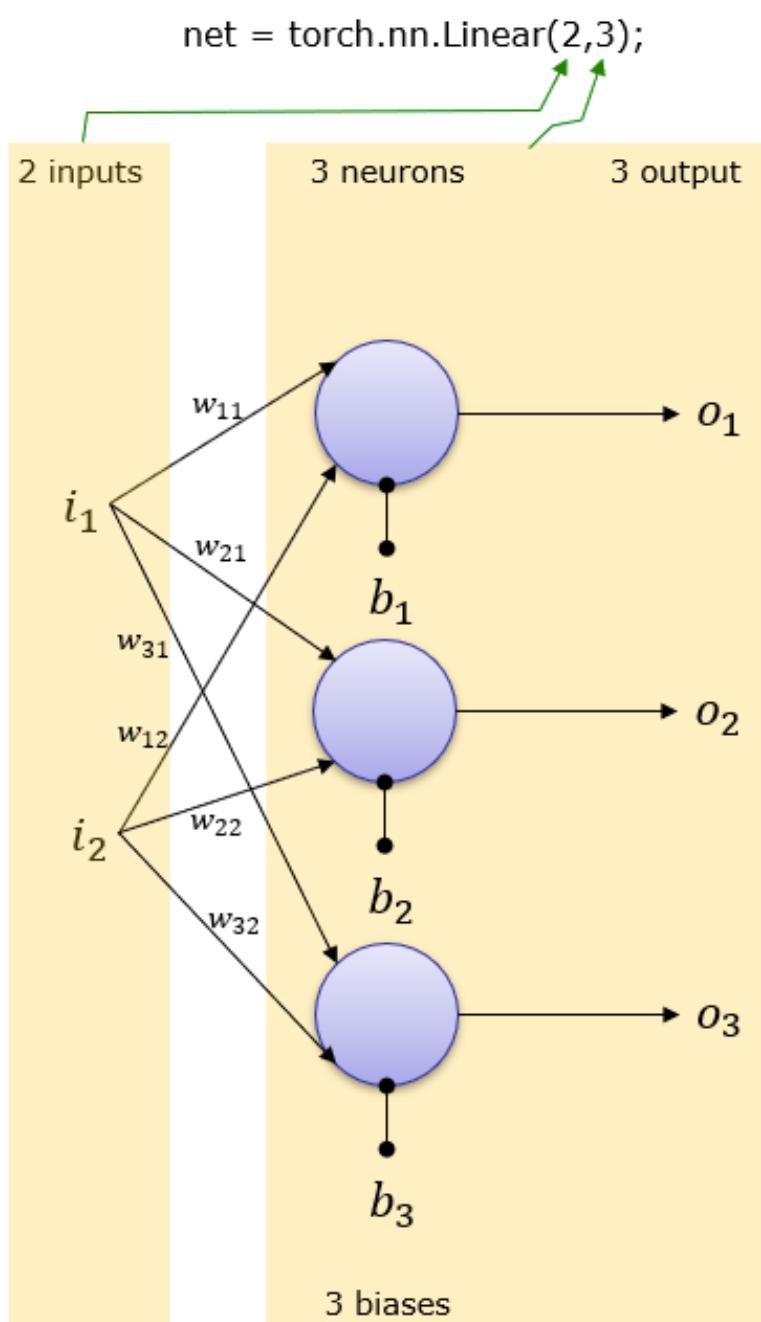
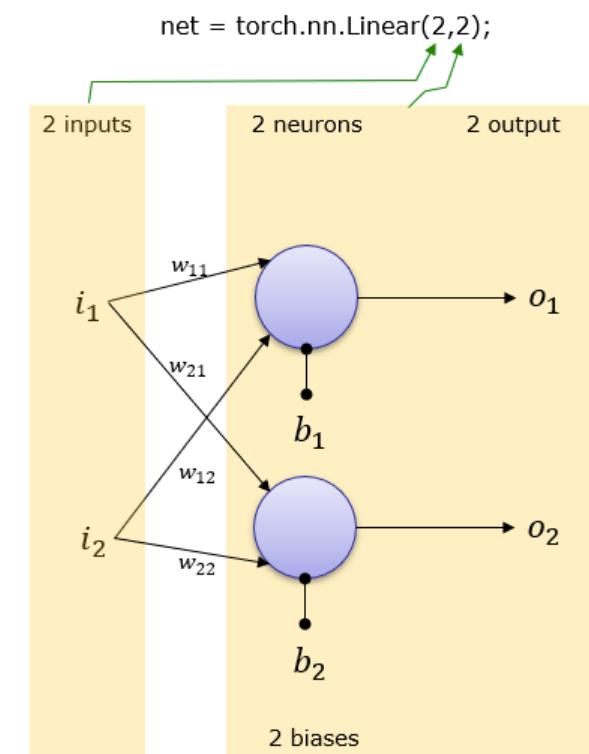
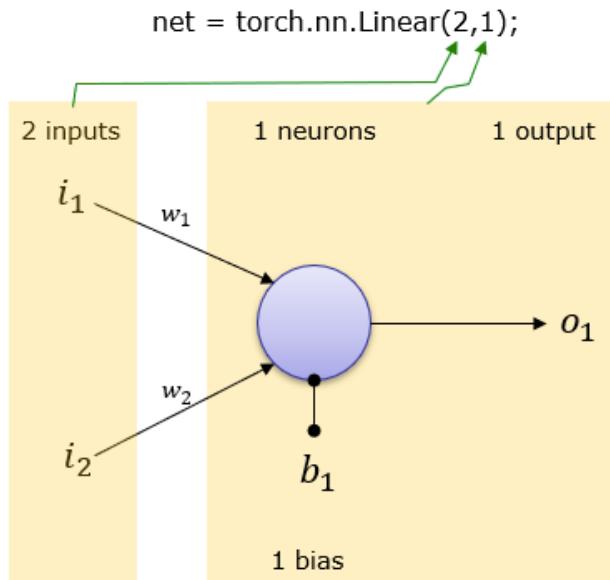
[ ] 1 # In ML we want our data to be of shape:
2 # (num_samples x num_dimensions)
3 X = X.reshape(N, 1)
4 Y = Y.reshape(N, 1)
5
6 # PyTorch uses float32 by default
7 # Numpy creates float64 by default
8 inputs = torch.from_numpy(X.astype(np.float32))
9 targets = torch.from_numpy(Y.astype(np.float32))

```

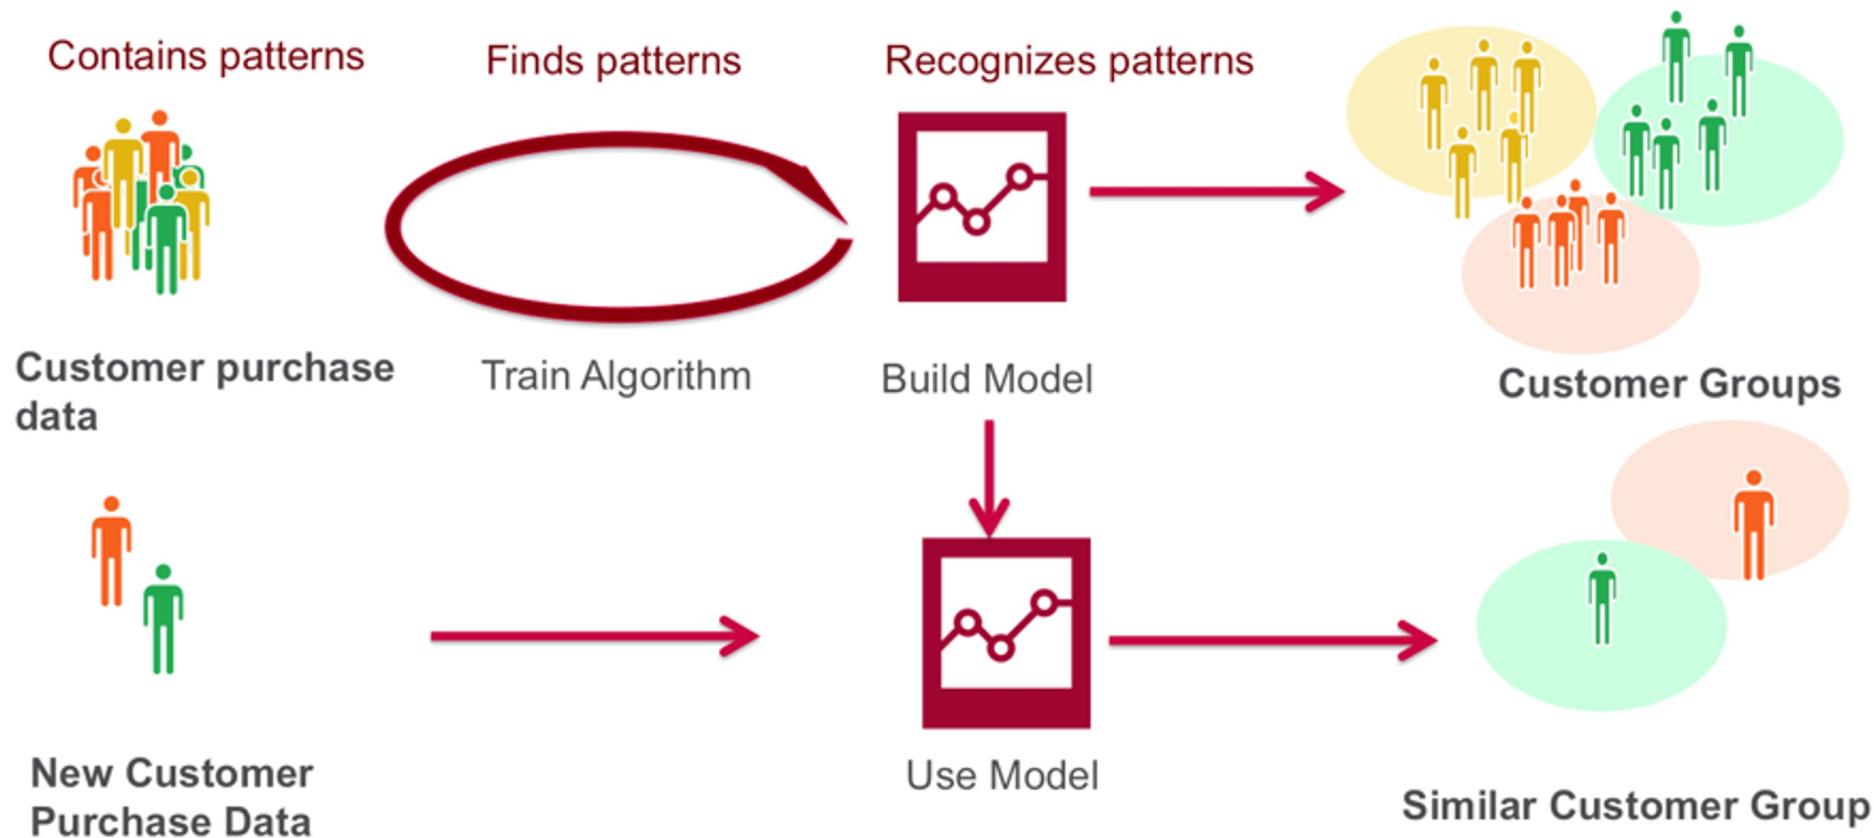
$$net = \text{torch}.Linear(1,1)$$



$$y = mx + b$$



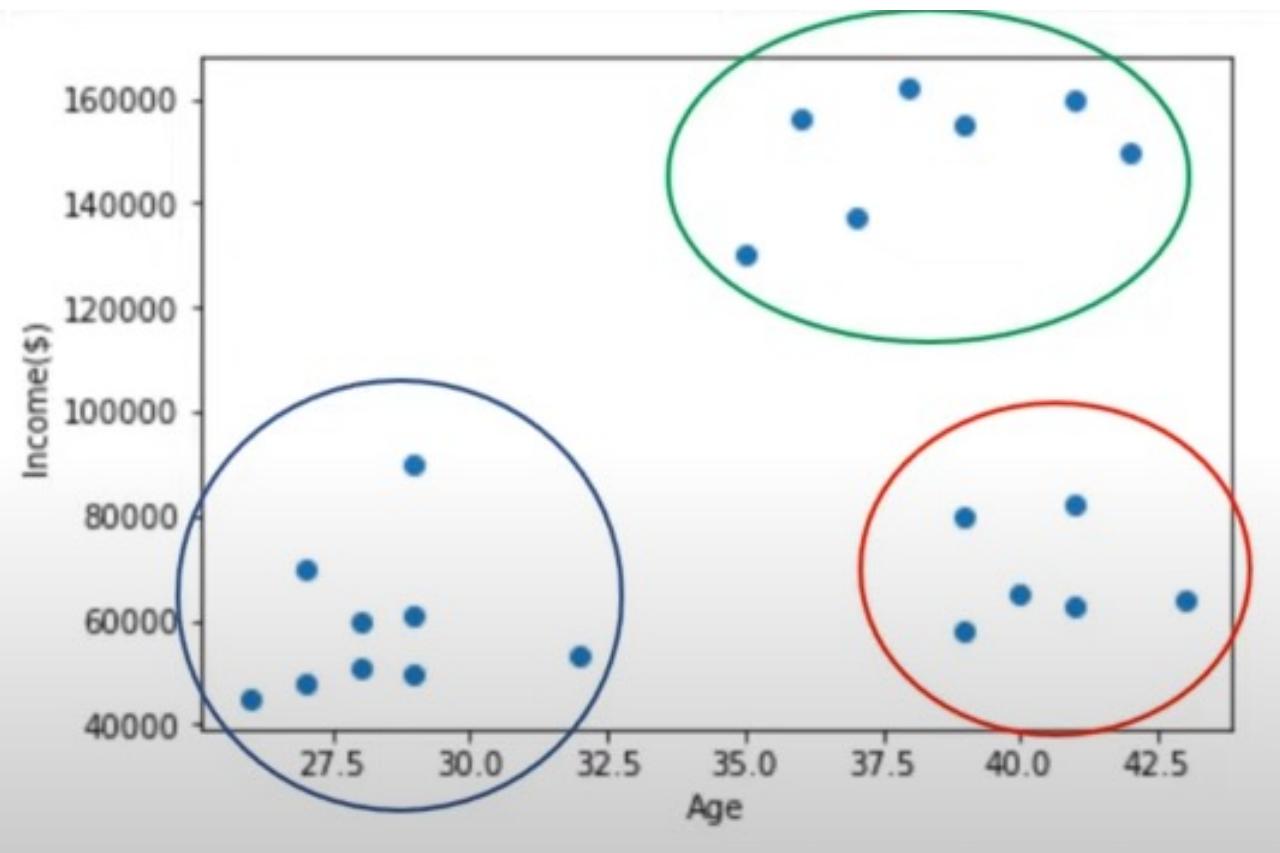
Unsupervised Learning | Self Supervised



Examples of clustering

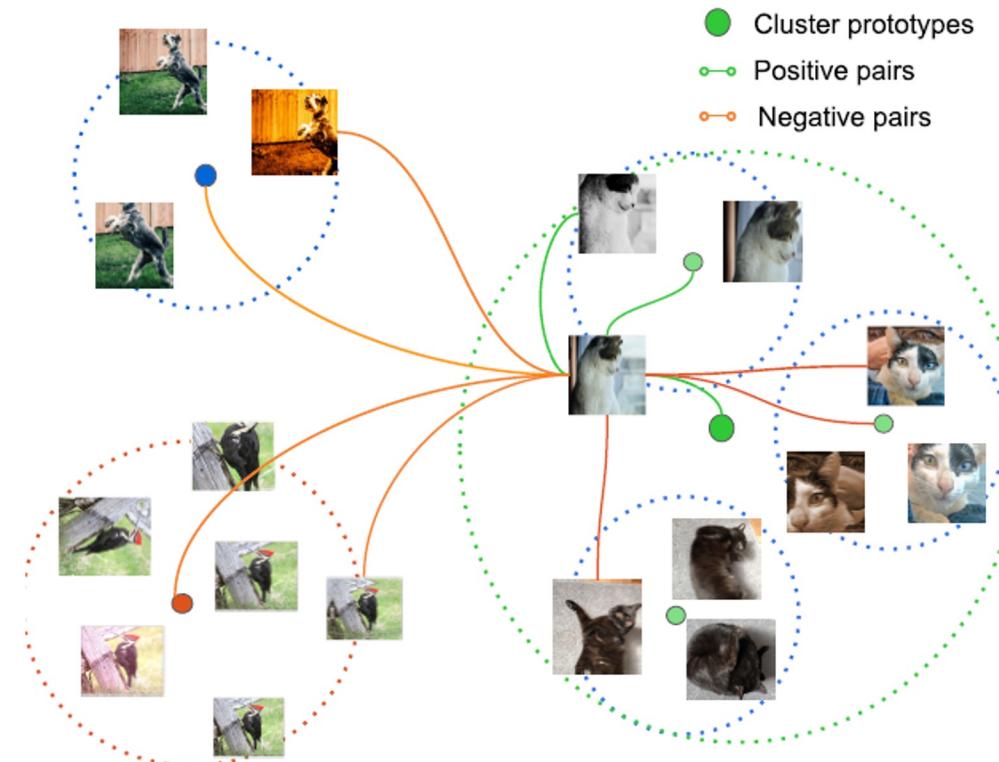
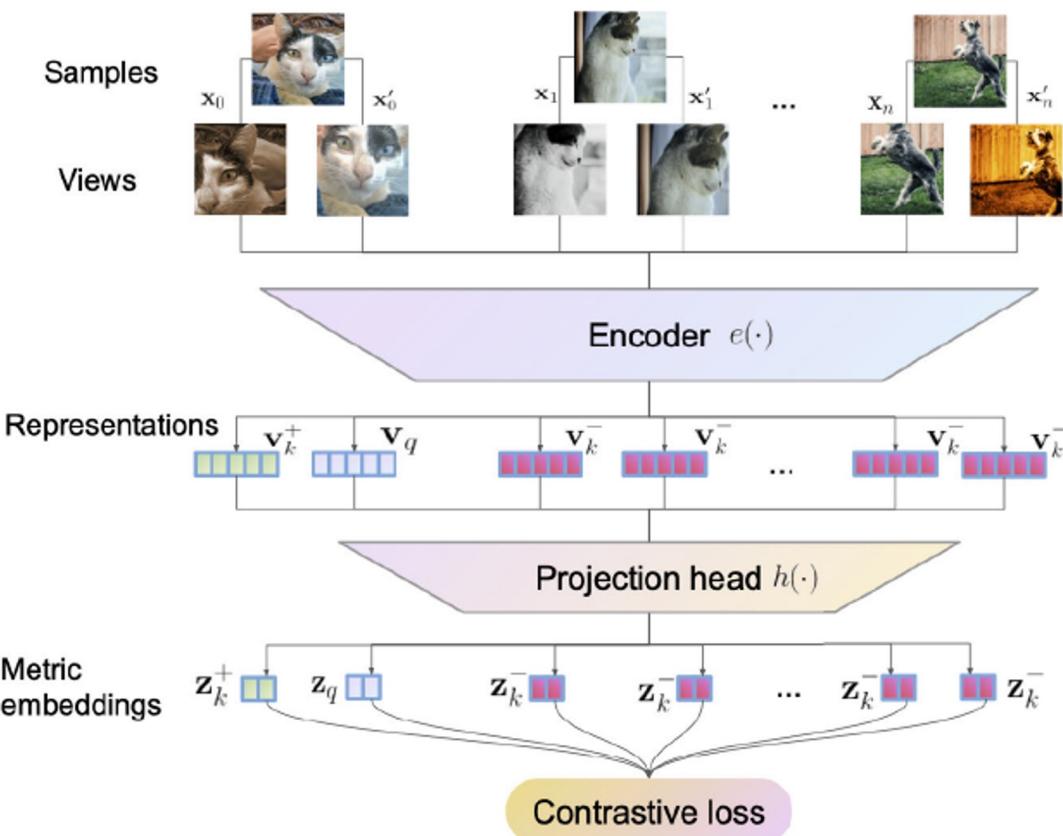
- Search results grouping
- Grouping similar customers
- Grouping similar patients
- Text categorization
- Network Security Anomaly detection

Unsupervised learning

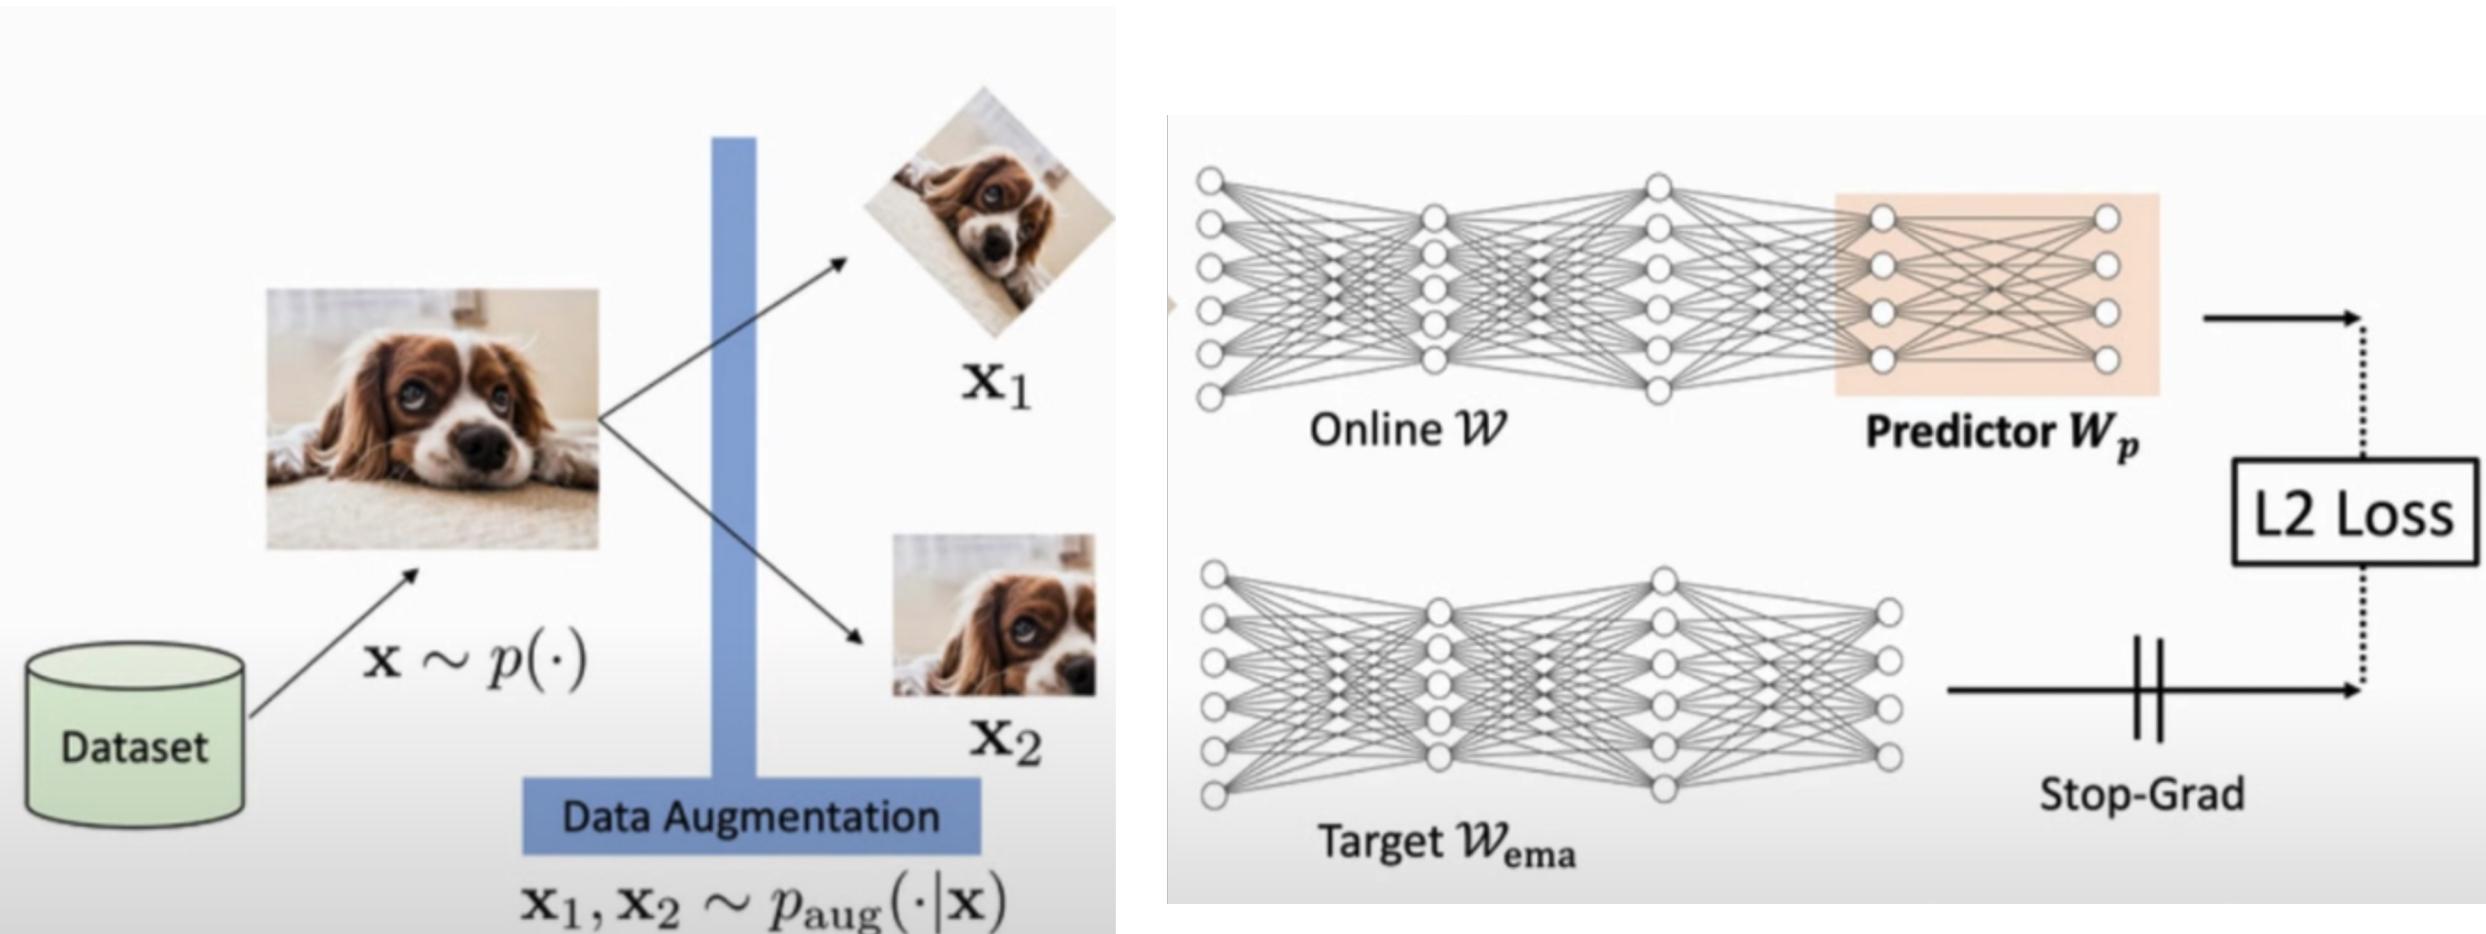


| | A | B | C |
|----|---------|-----|------------|
| 1 | Name | Age | Income(\$) |
| 2 | Rob | 27 | 70000 |
| 3 | Michael | 29 | 90000 |
| 4 | Mohan | 29 | 61000 |
| 5 | Ismail | 28 | 60000 |
| 6 | Kory | 42 | 150000 |
| 7 | Gautam | 39 | 155000 |
| 8 | David | 41 | 160000 |
| 9 | Andrea | 38 | 162000 |
| 10 | Brad | 36 | 156000 |

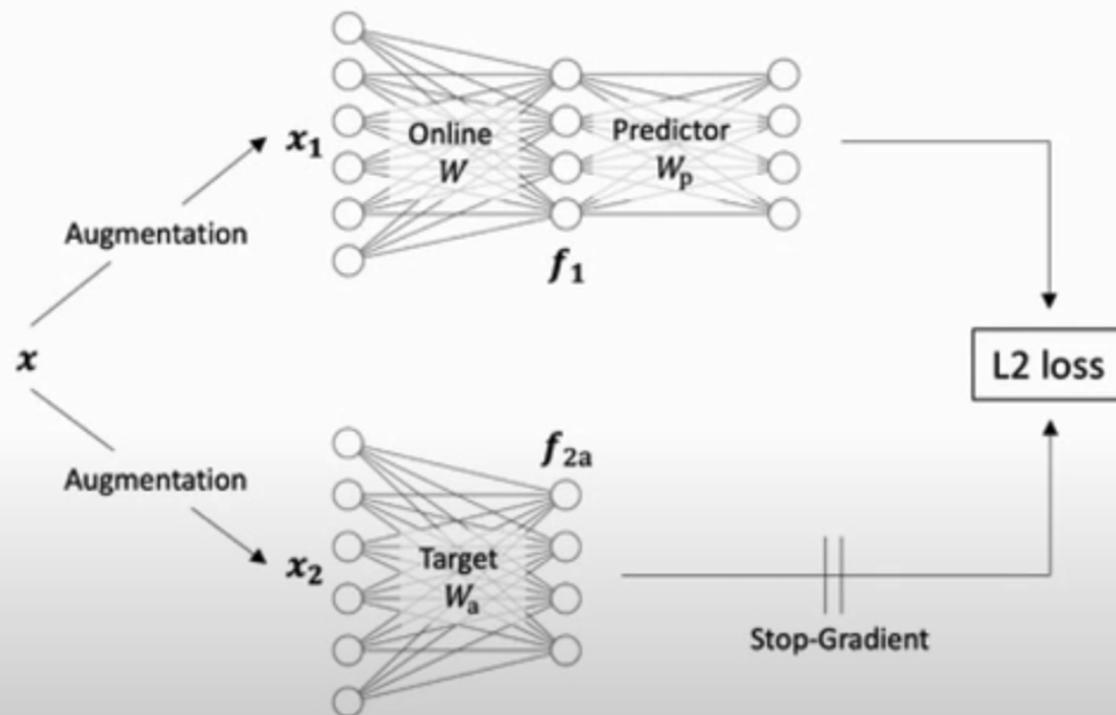
Self-Supervised



Easy review



A simple model



Objective:

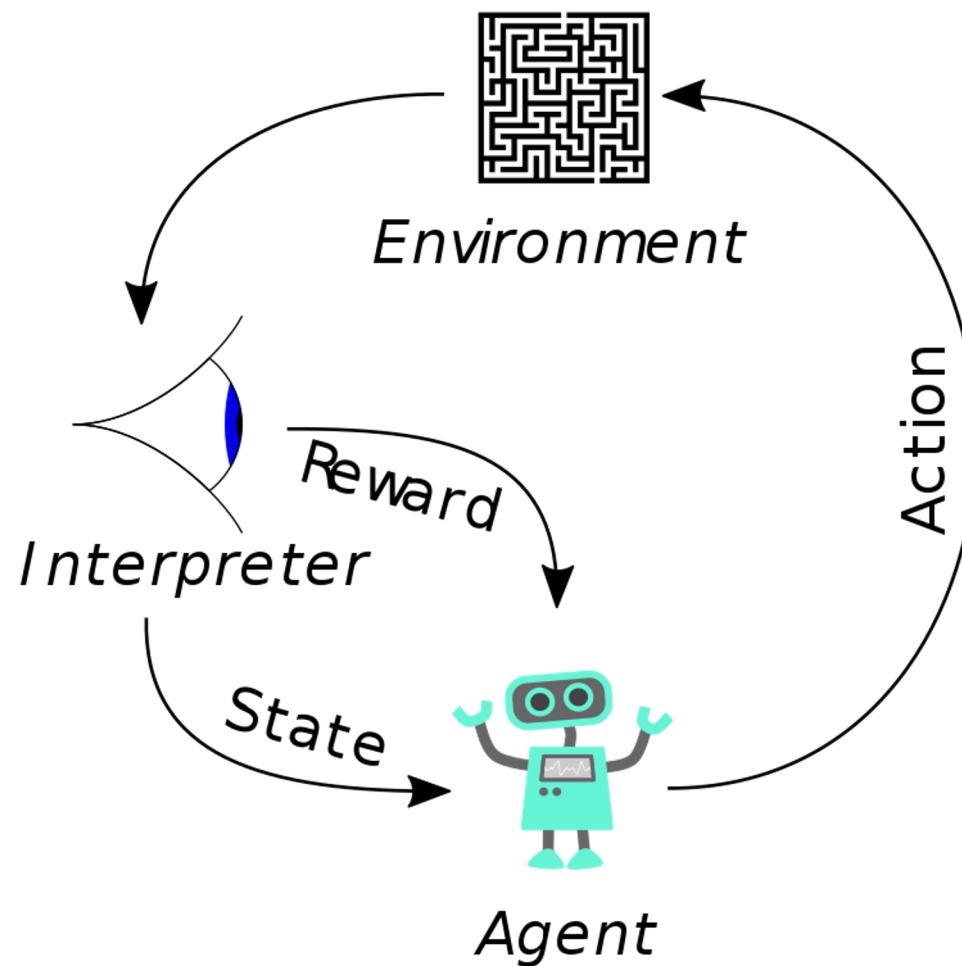
$$J(W, W_p) := \frac{1}{2} \mathbb{E}_{x_1, x_2} [\|W_p f_1 - \text{StopGrad}(f_{2a})\|_2^2]$$

Linear online network W

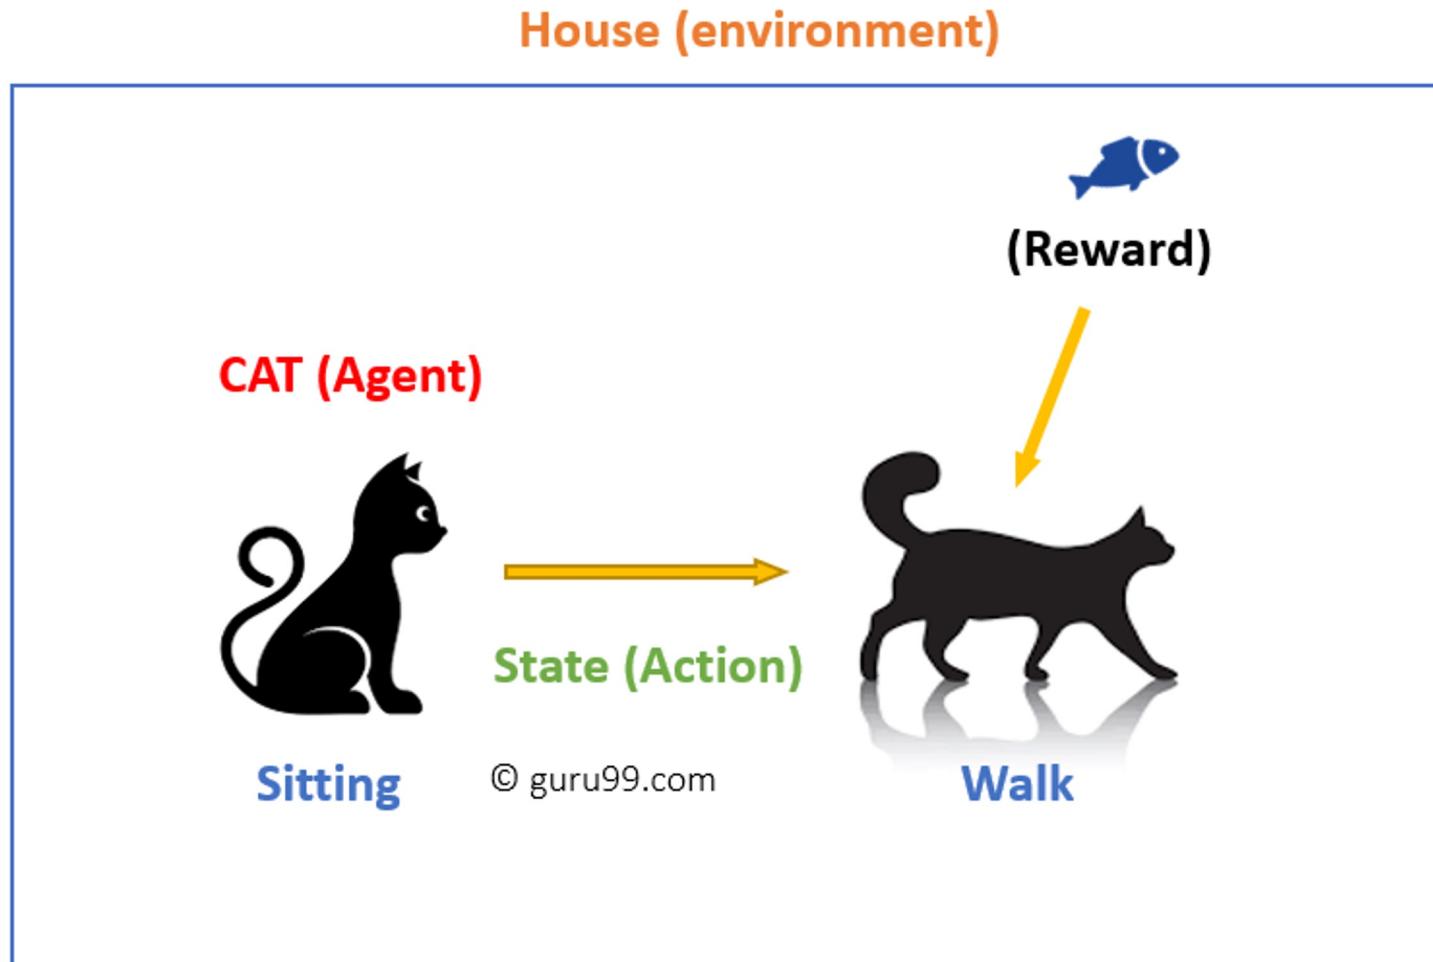
Linear target network W_a

Linear predictor W_p

Reinforcement learning



How Reinforcement Learning works?



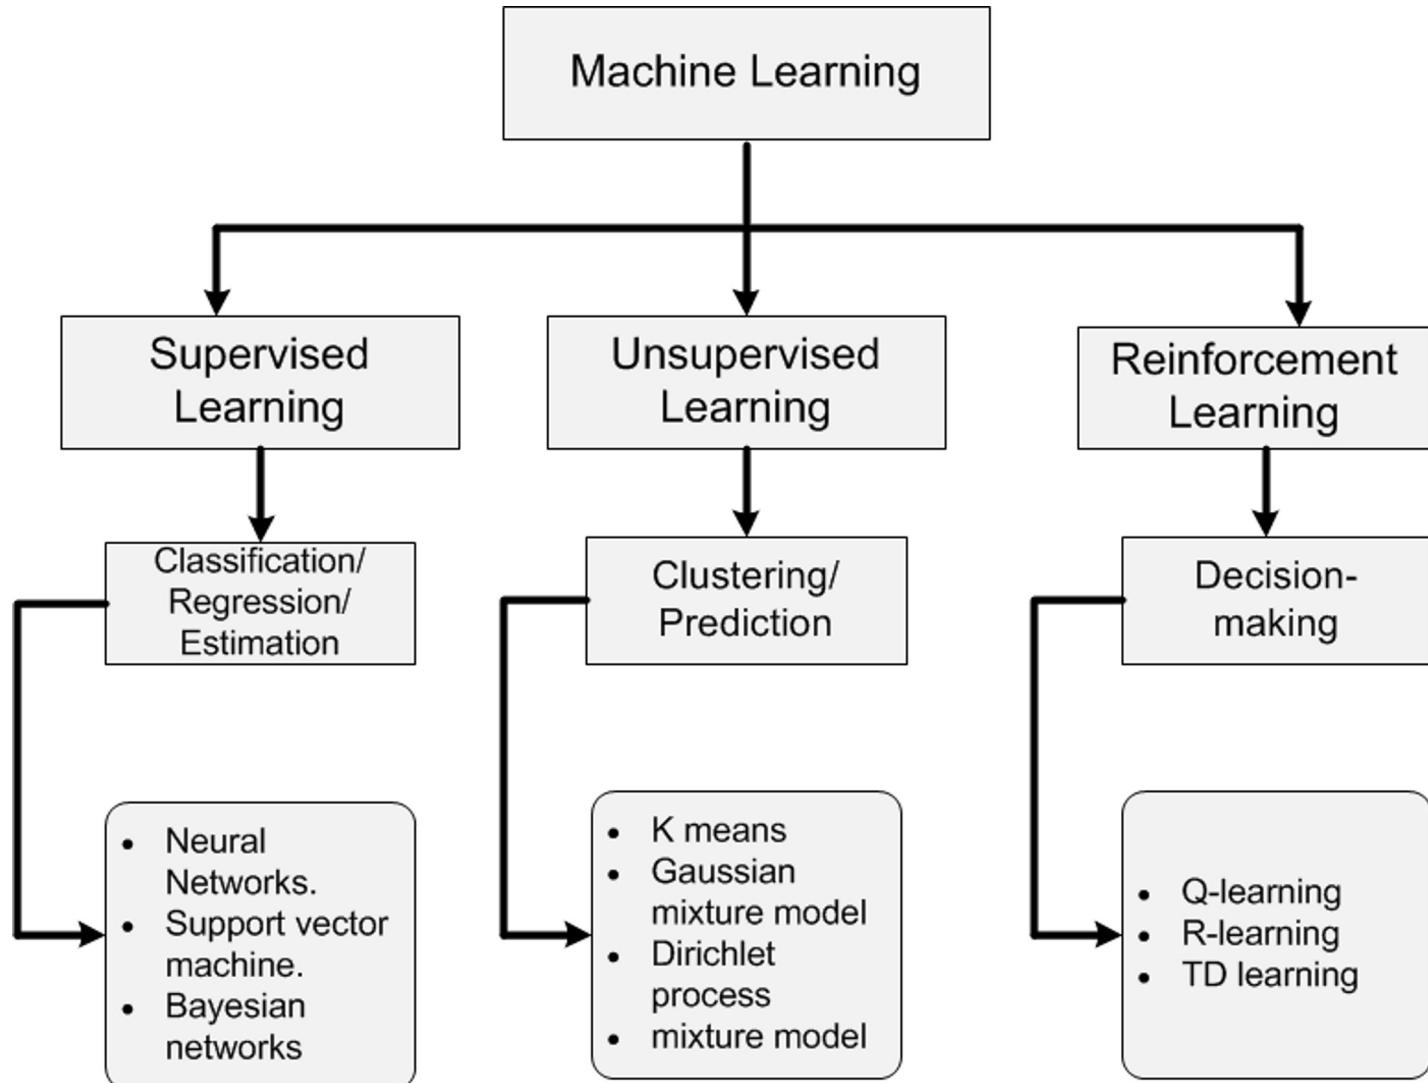
Source: Reinforcement Learning: What is, Algorithms, Applications,
Episode 0, 2023

How Reinforcement Learning works?

- Your cat is an agent that is exposed to the environment. In this case, it is your house. An example of a state could be your cat sitting, and you use a specific word in for cat to walk.
- Our agent reacts by performing an action transition from one "state" to another "state."
- For example, your cat goes from sitting to walking.
- The reaction of an agent is an action, and the policy is a method of selecting an action given a state in expectation of better outcomes.
- After the transition, they may get a reward or penalty in return.

Examples of Reinforcement Learning

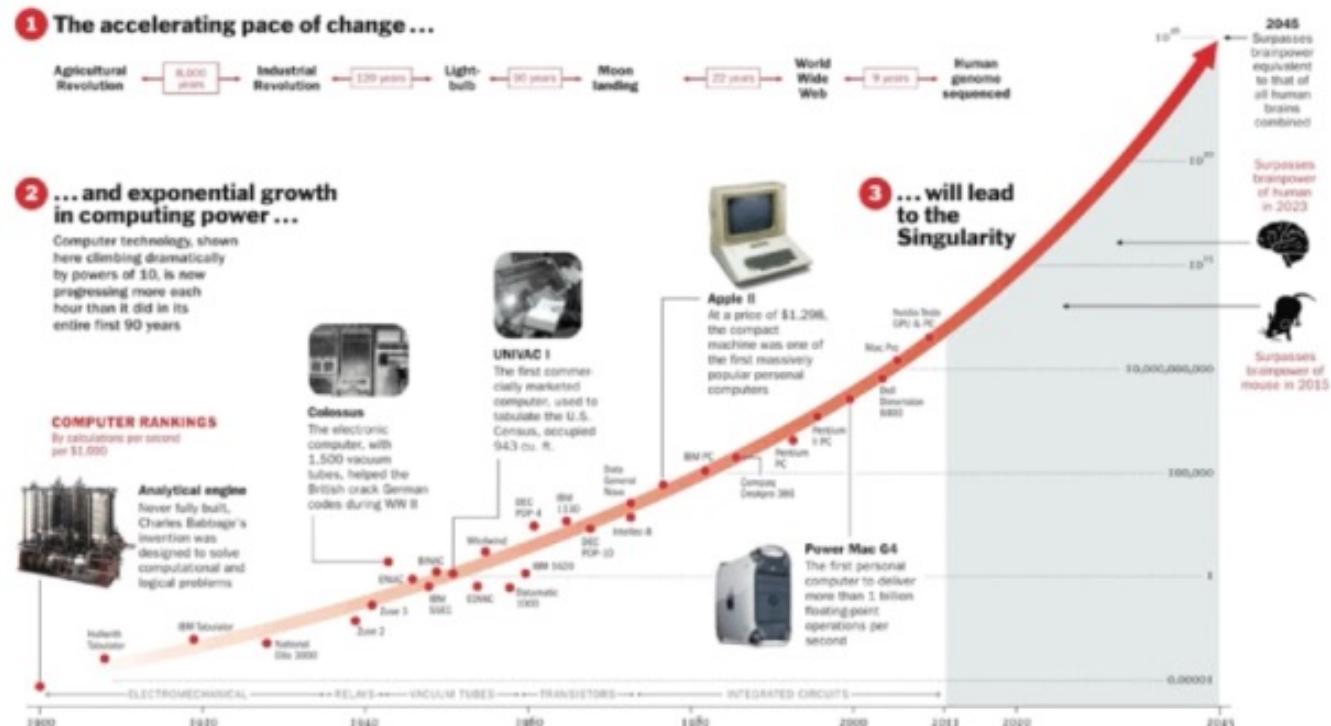
- Robotics for industrial automation.
- Business strategy planning
- Self-driving car
- It helps you to create training systems that provide custom instruction and materials according to the requirement of students.
- Aircraft control and robot motion control



Source: https://www.researchgate.net/figure/Comparison-of-different-types-of-machine-learning_fig6_325928183

Real-World Dataset: Moore's Law

- The number of transistors per square inch on integrated circuits doubles approximately every 2 years



Exponential Growth of Computing

Twentieth through twenty first century



Logarithmic Plot

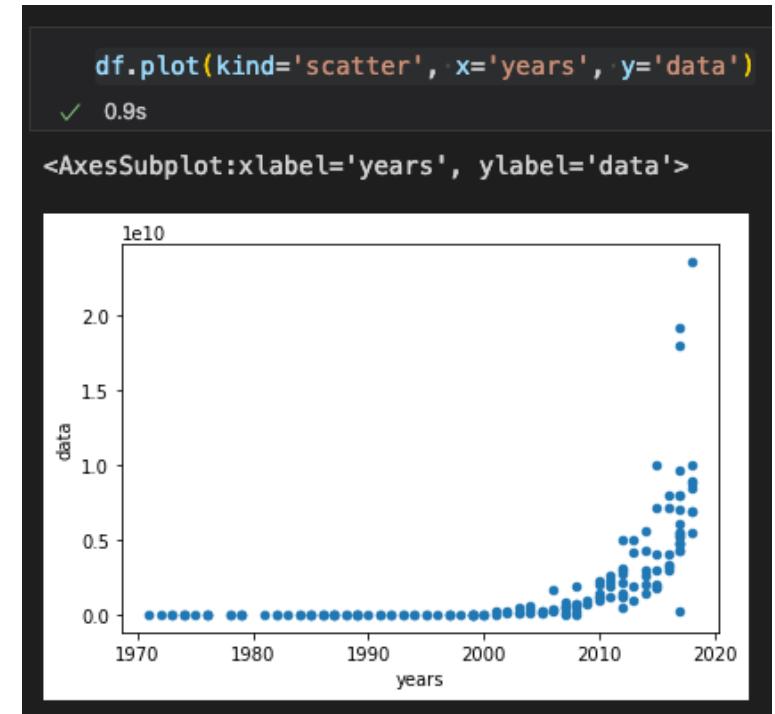


Home Work 1

```
df = pd.read_csv('https://raw.githubusercontent.com/lazyprogrammer/machine_learning_examples/master/tf2.0/moore.csv', header=None)
df= df.rename(columns={0: "years", 1: "data"})
df
0.2s
```

| | years | data |
|-----|-------|-------------|
| 0 | 1971 | 2300 |
| 1 | 1972 | 3500 |
| 2 | 1973 | 2500 |
| 3 | 1973 | 2500 |
| 4 | 1974 | 4100 |
| ... | ... | ... |
| 157 | 2017 | 18000000000 |
| 158 | 2017 | 19200000000 |
| 159 | 2018 | 8876000000 |
| 160 | 2018 | 23600000000 |
| 161 | 2018 | 9000000000 |

162 rows x 2 columns



https://github.com/Tuchsanai/DL-FOR-COMPUTER-VISION-2565_1/blob/main/week1/code/Moore_rules.ipynb

Home Work 2

$$\text{signal} = A \sin(t + \theta) + b$$

$$y = \text{signal} + \text{noise}$$

https://github.com/Tuchsanai/DL-FOR-COMPUTER-VISION/blob/main/week1/code/sinewith_noise.ipynb

```
t = np.arange(0, 20, 0.1)

#random data on x-axis
A = np.random.rand()
b = 20*np.random.rand()
ceta = 2*np.random.rand()*np.pi

signal = A*np.sin(t+ceta)+b
y = signal + 0.1*np.random.randn(len(t))

print(f"A = {A}, ceta = {ceta}")
✓ 0.2s
= 0.666992306605201, ceta = 0.26461134460282454
```

