DEEP LEARNING FOR COMPUTER VISION

Week7

Dr. Tuchsanai. PloySuwan

Autoencoders

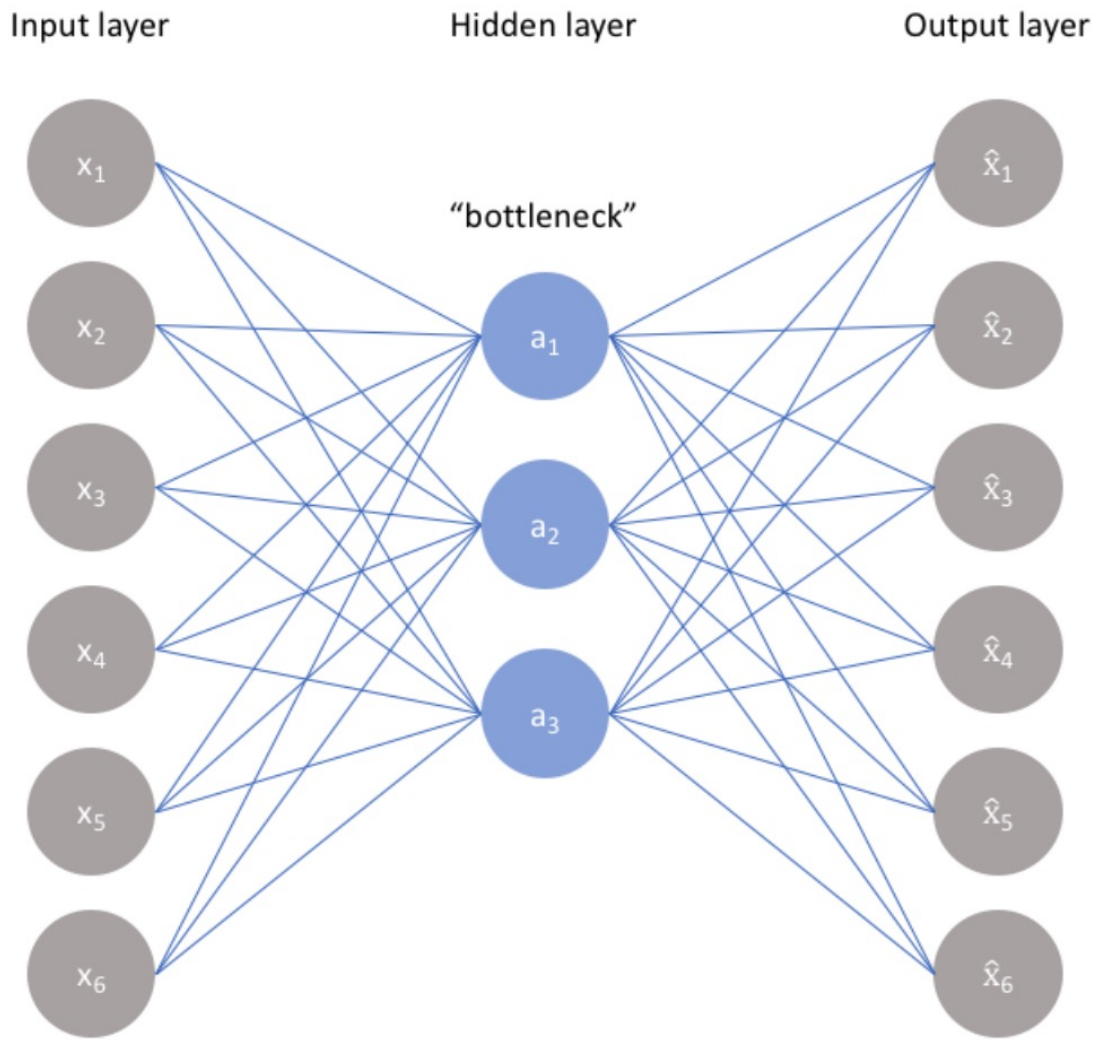
Input layer
Hidden layer
Output layer
"bottleneck"
x_1
x_2
x_3
x_4
x_5
x_6
a_1
a_2
a_3
$\hat{x}_1$
$\hat{x}_2$
$\hat{x}_3$
$\hat{x}_4$
$\hat{x}_5$
$\hat{x}_6$

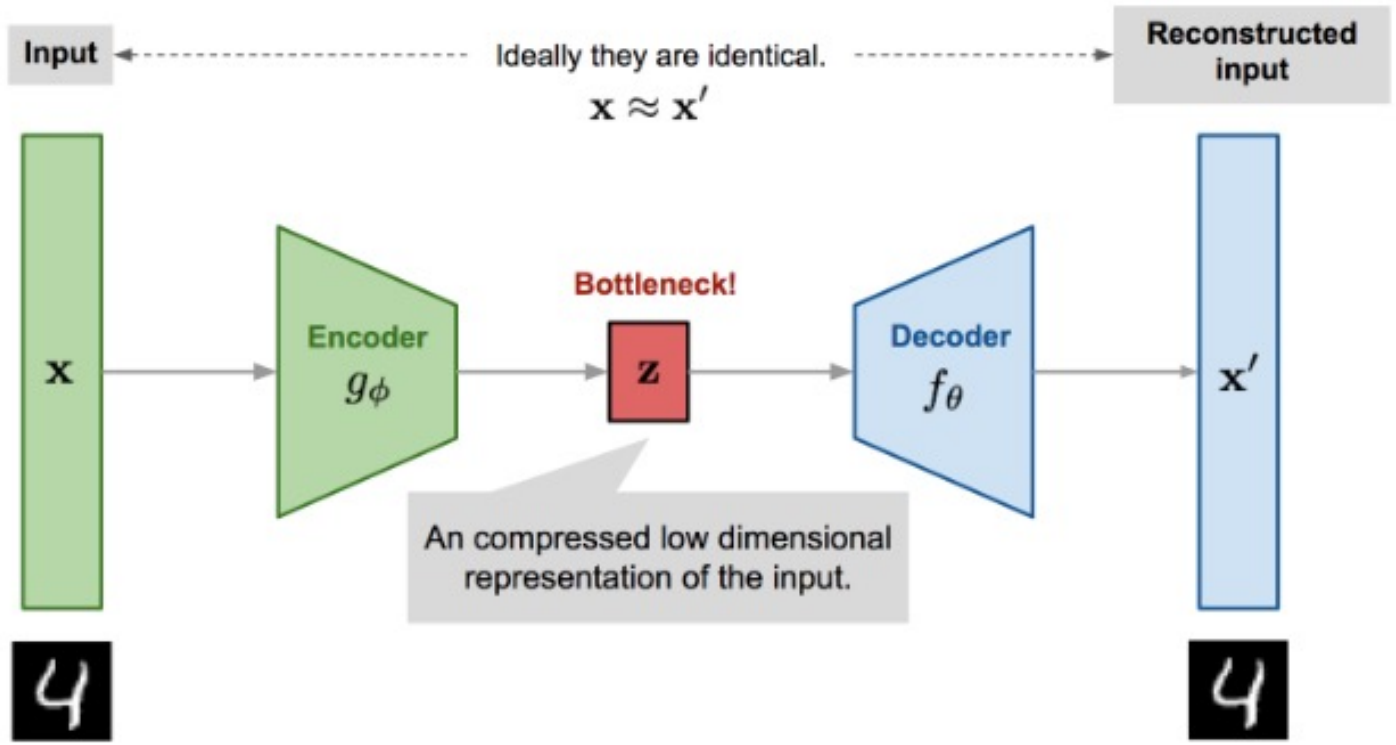

$$L_{\text{AE}}(\theta, \phi) = \frac{1}{n} \sum_{i=1}^{n} (\mathbf{x}^{(i)} - f_\theta(g_\phi(\mathbf{x}^{(i)})))^2$$

Data compression

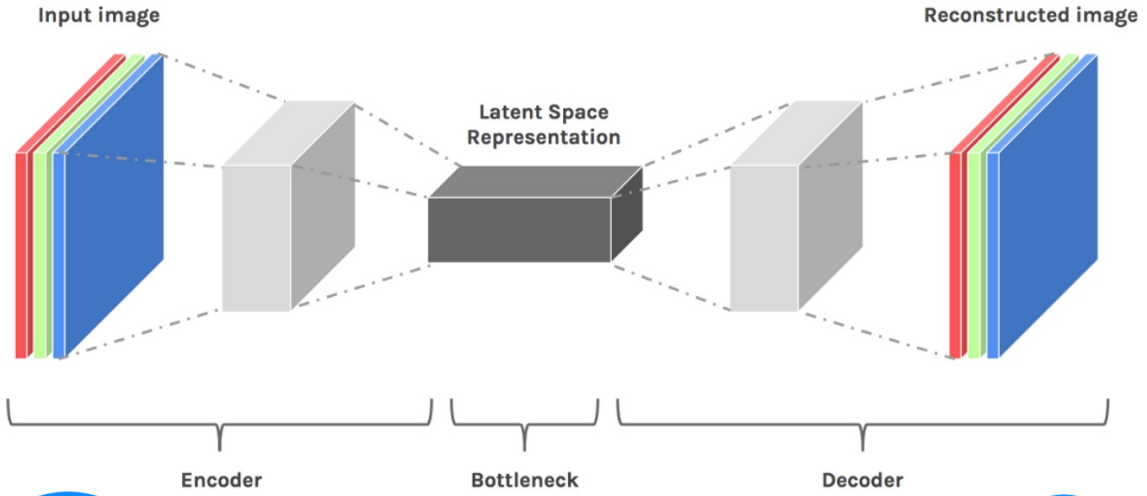

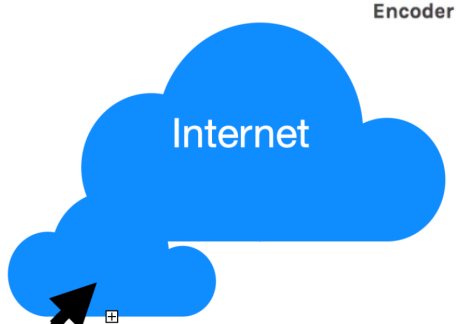

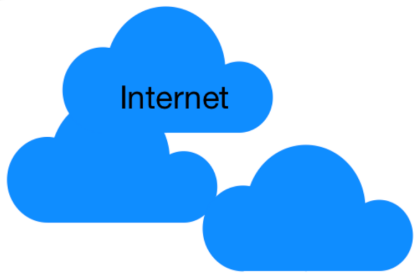

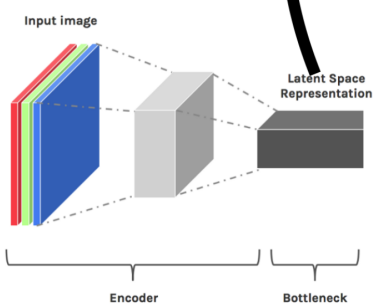

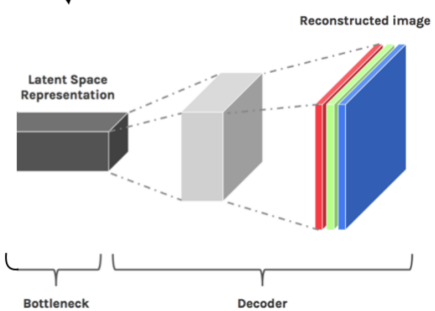

Denoising data

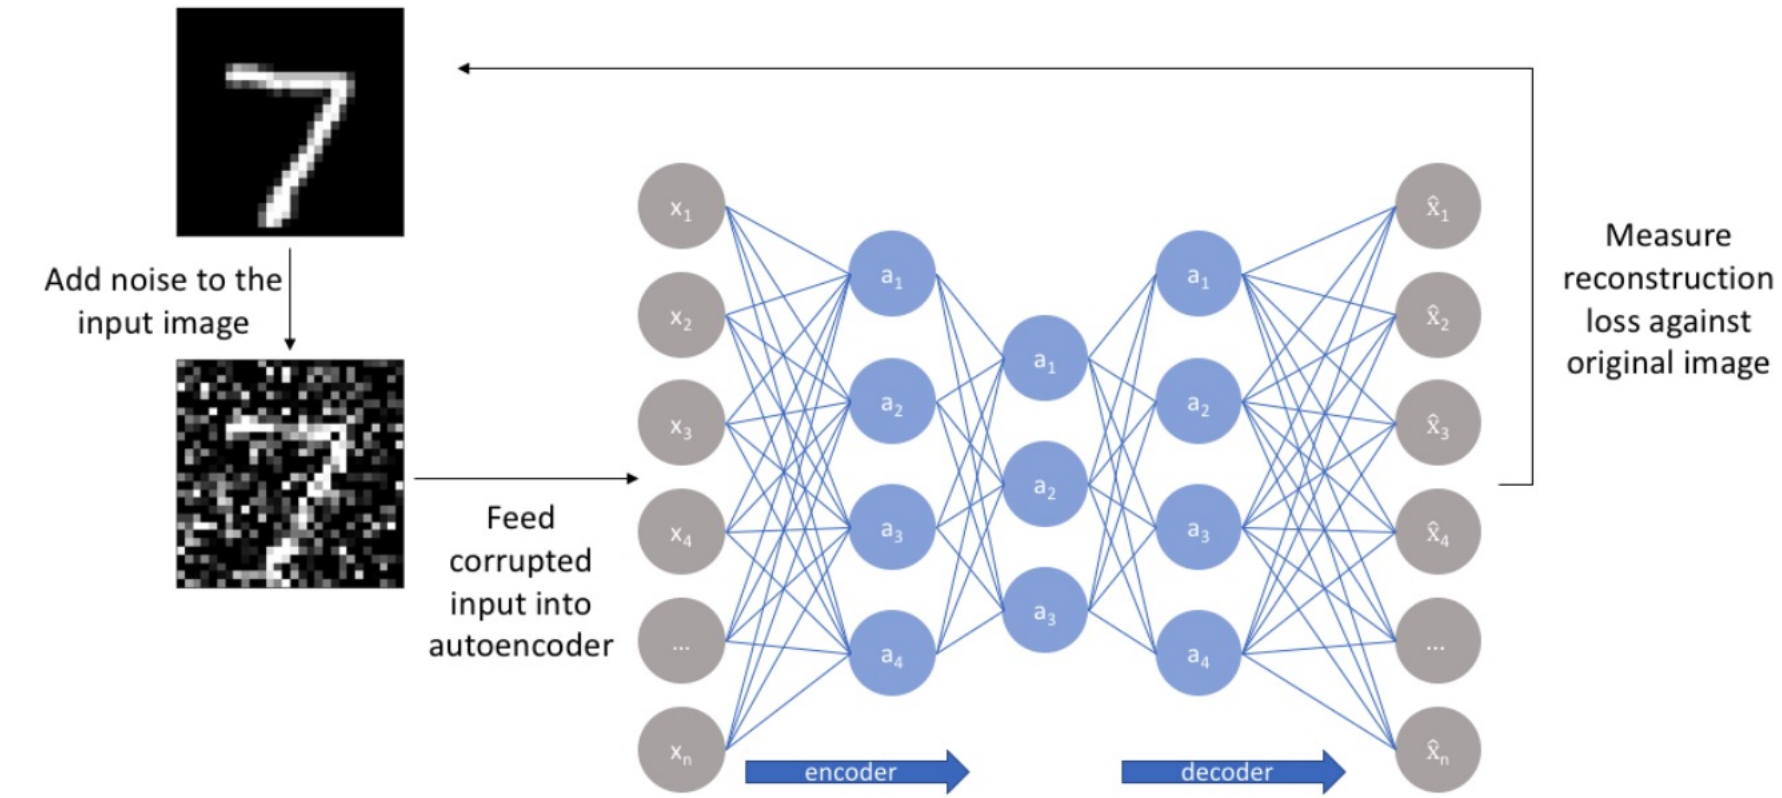

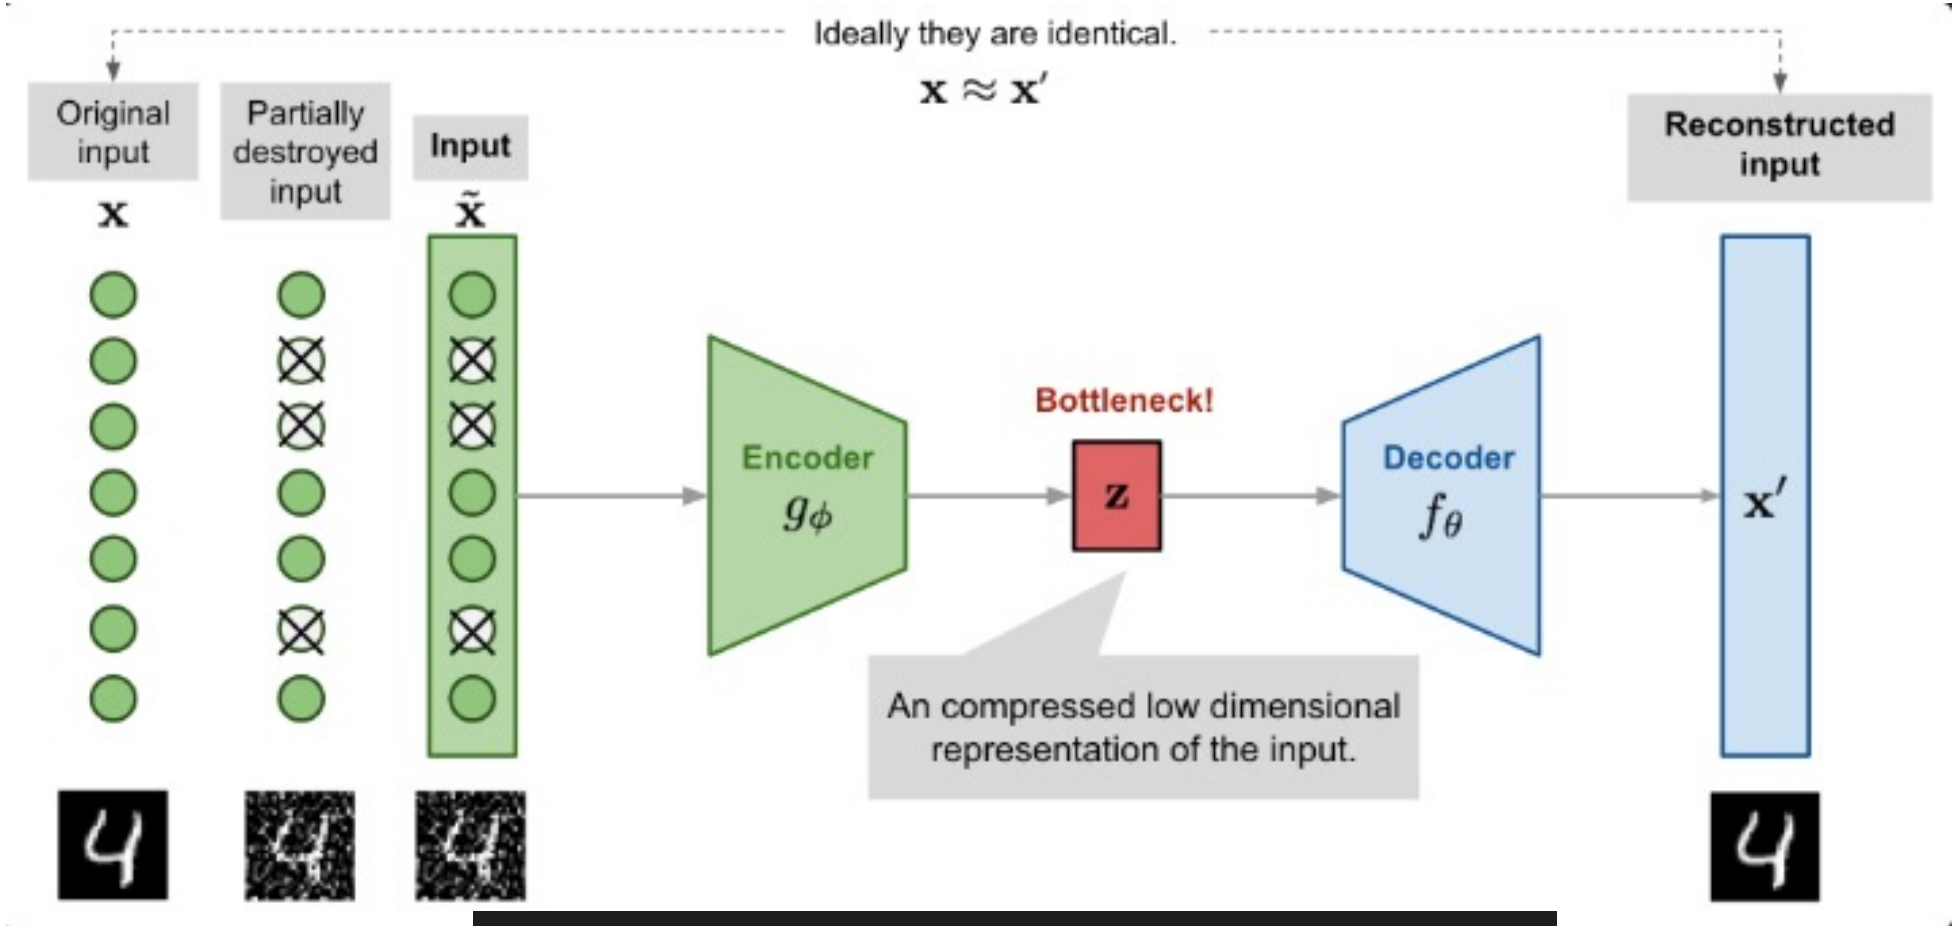

$$\tilde{\mathbf{x}}^{(i)} \sim \mathcal{M}_{\mathcal{D}}(\tilde{\mathbf{x}}^{(i)}|\mathbf{x}^{(i)})$$

$$L_{\text{DAE}}(\theta, \phi) = \frac{1}{n}\sum_{i=1}^{n}(\mathbf{x}^{(i)} - f_\theta(g_\phi(\tilde{\mathbf{x}}^{(i)})))^2$$

Introduction to Autoencoders for Anomaly Detection

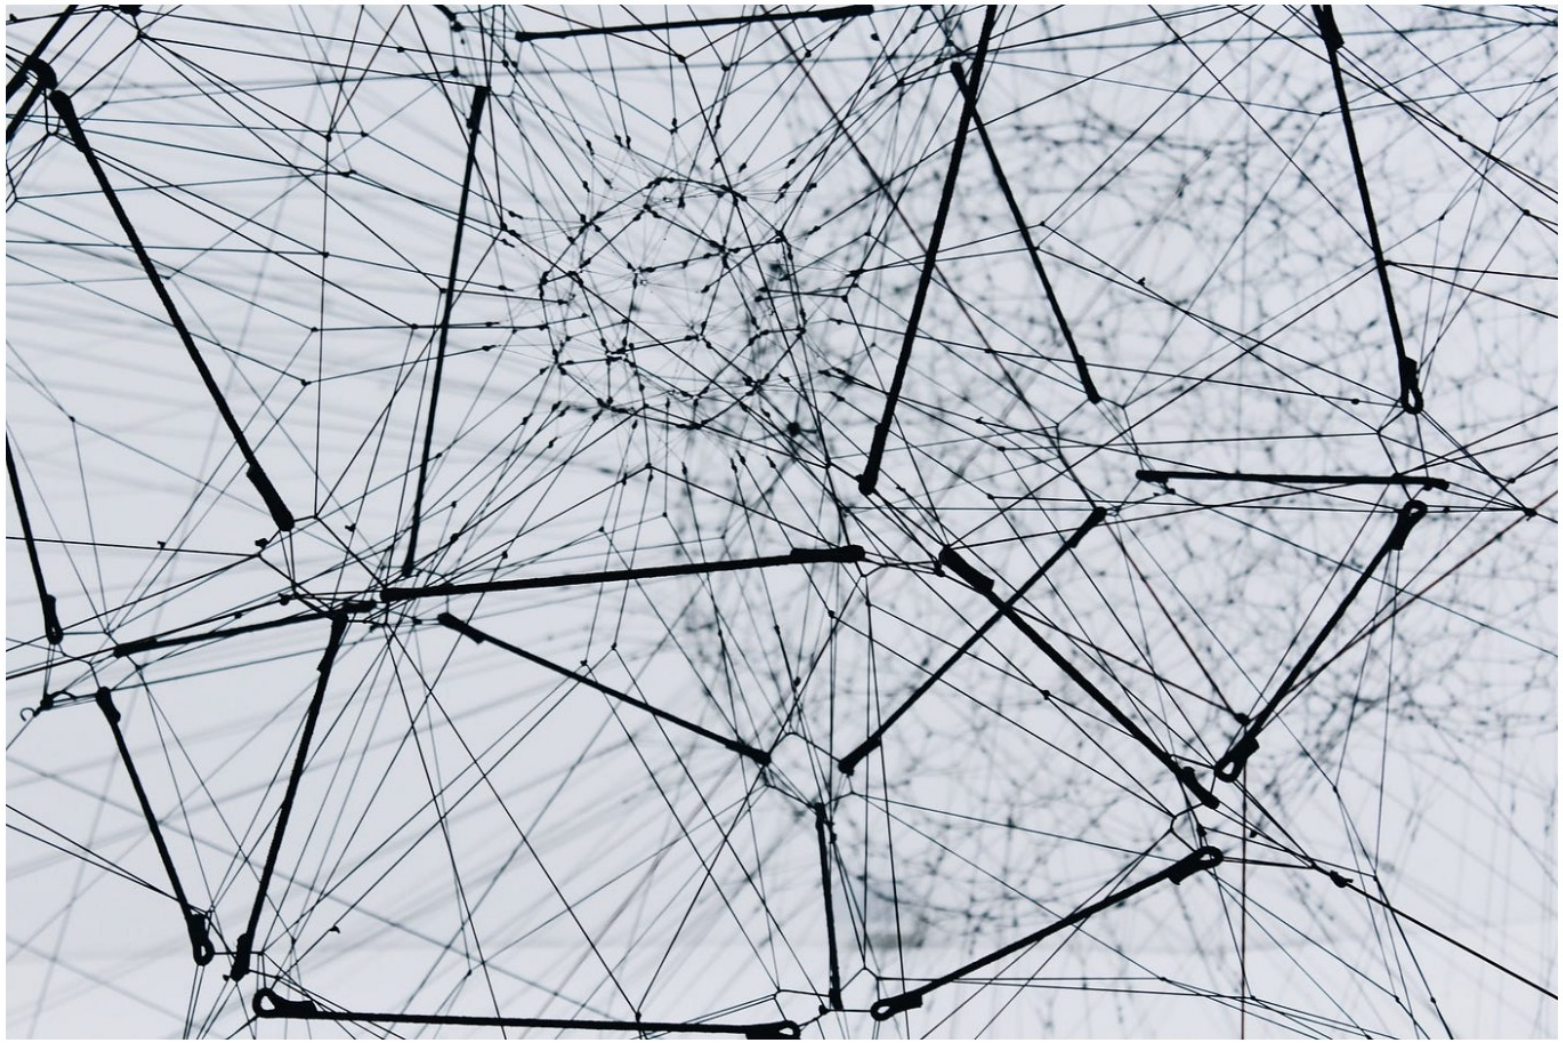

Photo by Alina Grubnyak / Unsplash

Introduction to Autoencoders for Anomaly Detection

https://medium.com/artificialis/introduction-to-autoencoders-for-anomaly-detection-a9897591cc72

Outlier Detection with Local Outlier Factor (LOF)

https://medium.com/@volkanillik/outlier-detection-with-local-outlier-factor-lof-8d1e29acabb4

Handbook of Anomaly Detection with Python Outlier Detection — (12) Autoencoders

https://towardsdatascience.com/anomaly-detection-with-autoencoder-b4cdce4866a6

Time Series Anomaly Detection: The Detective's Toolbox

https://medium.com/georgian-impact-blog/time-series-anomaly-detection-the-detectives-toolbox-9ef131dddaf9

Variational autoencoders

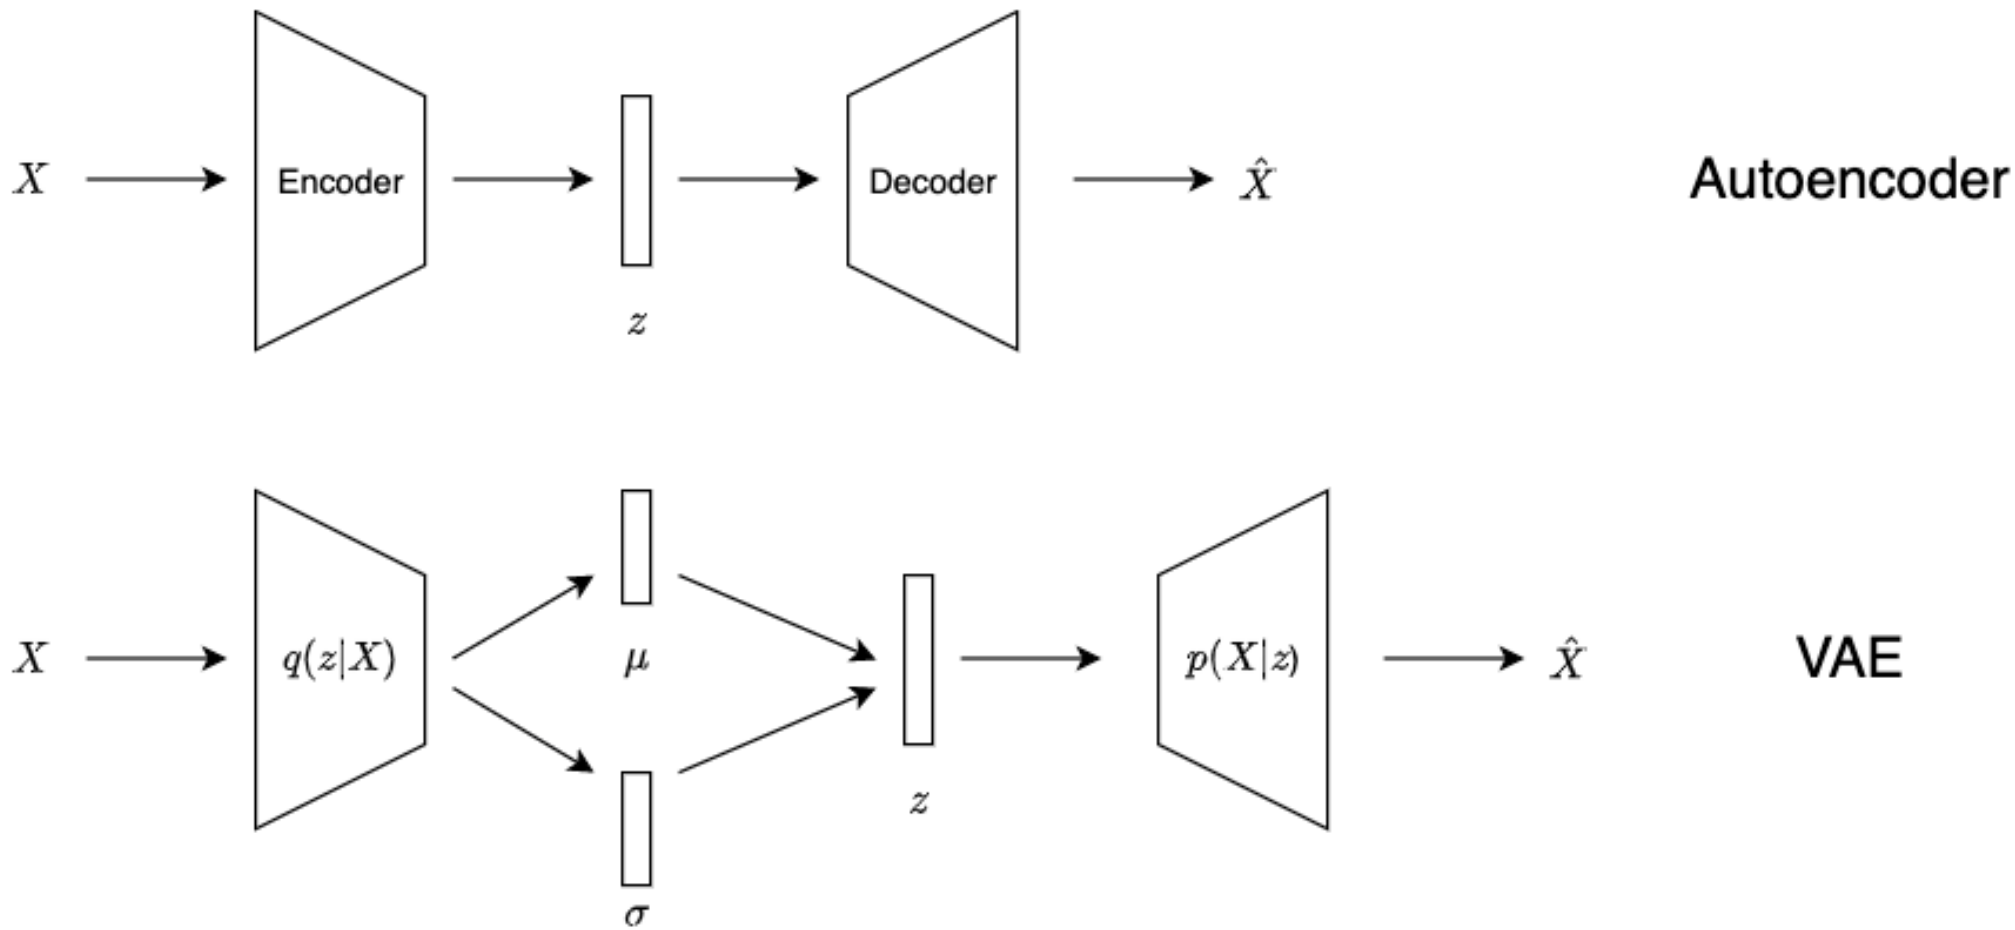
X
Encoder
z
Decoder
$\hat{X}$
Autoencoder
X
$q(z|X)$
μ
σ
z
$p(X|z)$
$\hat{X}$
VAE

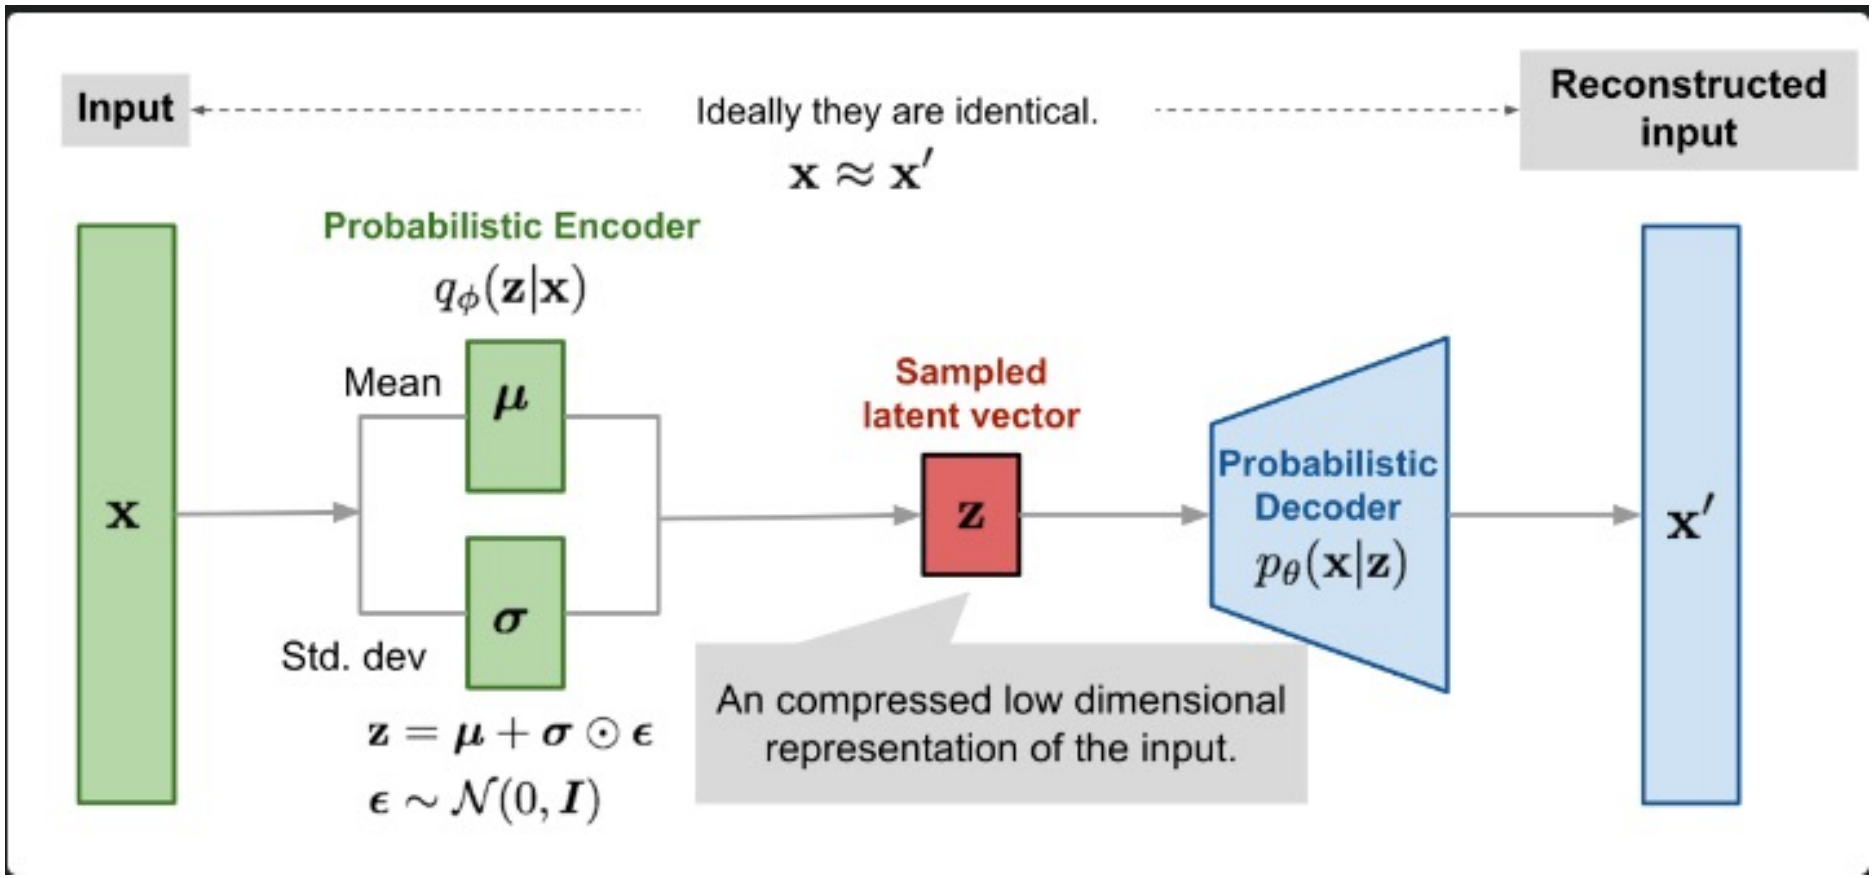
Input
Ideally they are identical.
$\mathbf{x} \approx \mathbf{x}'$
Reconstructed input
Probabilistic Encoder
$q_\phi(\mathbf{z}|\mathbf{x})$
Mean
$\boldsymbol{\mu}$
Std. dev
$\boldsymbol{\sigma}$
$\mathbf{z} = \boldsymbol{\mu} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon}$
$\boldsymbol{\epsilon} \sim \mathcal{N}(0, \boldsymbol{I})$
$\mathbf{x}$
Sampled latent vector
$\mathbf{z}$
An compressed low dimensional representation of the input.
Probabilistic Decoder
$p_\theta(\mathbf{x}|\mathbf{z})$
$\mathbf{x}'$

What is Variational AutoEncoder?

VAE is an autoencoder whose encodings distribution is regularised during the training in order to ensure that its latent space has good properties allowing us to generate some new data. A variational autoencoder (VAE) provides a probabilistic manner for describing an observation in latent space.

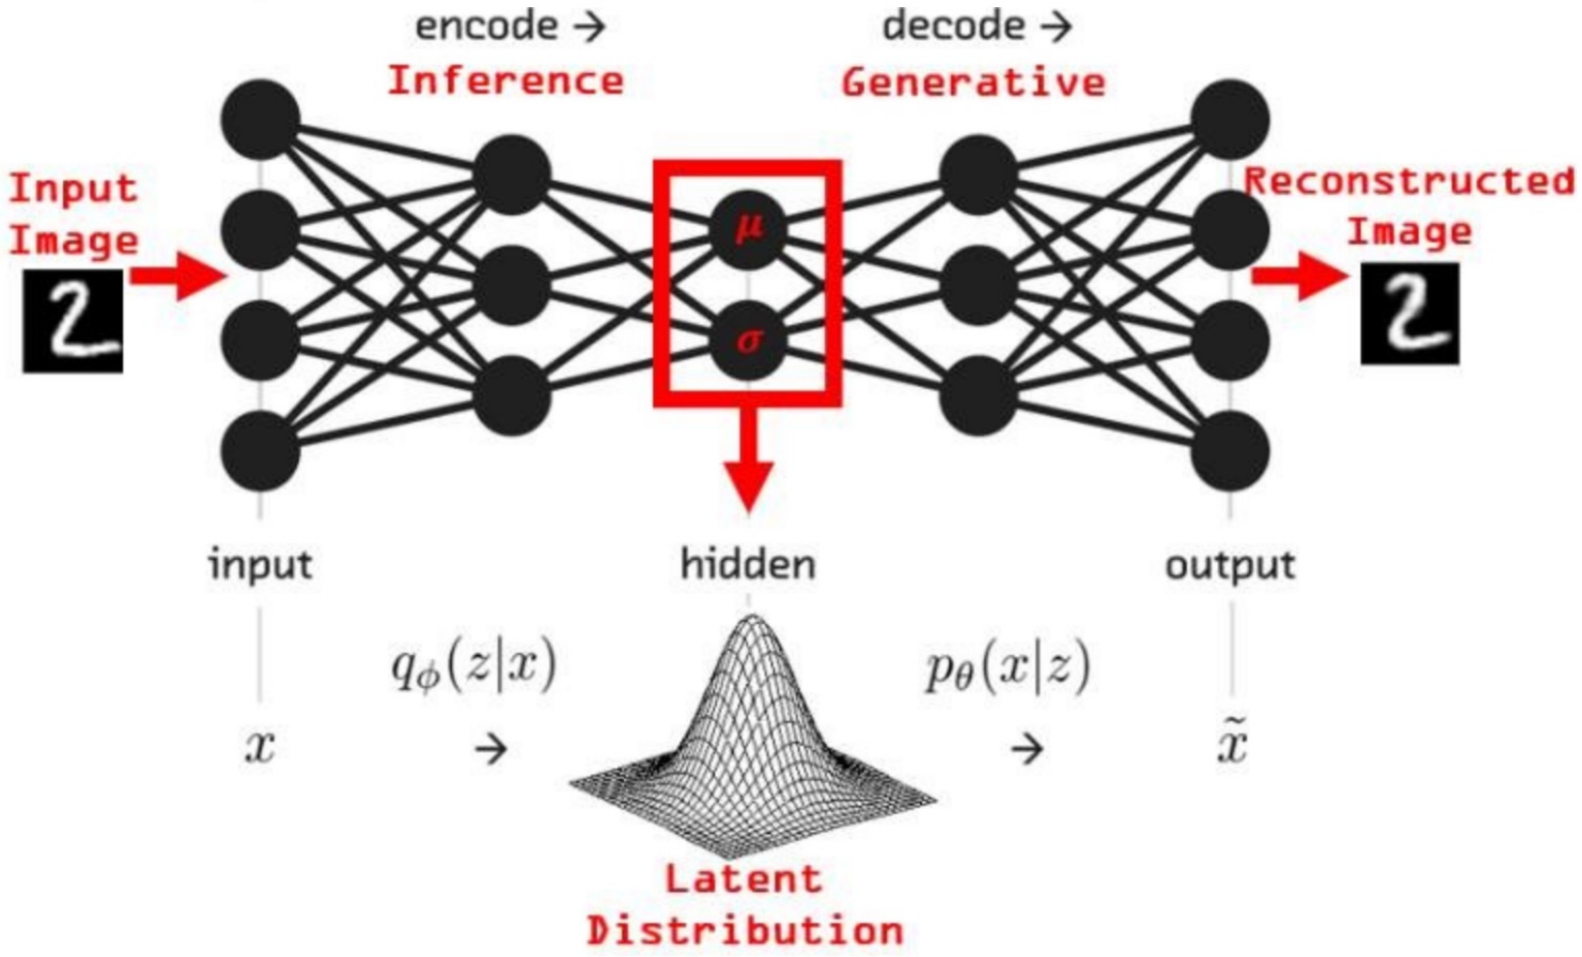

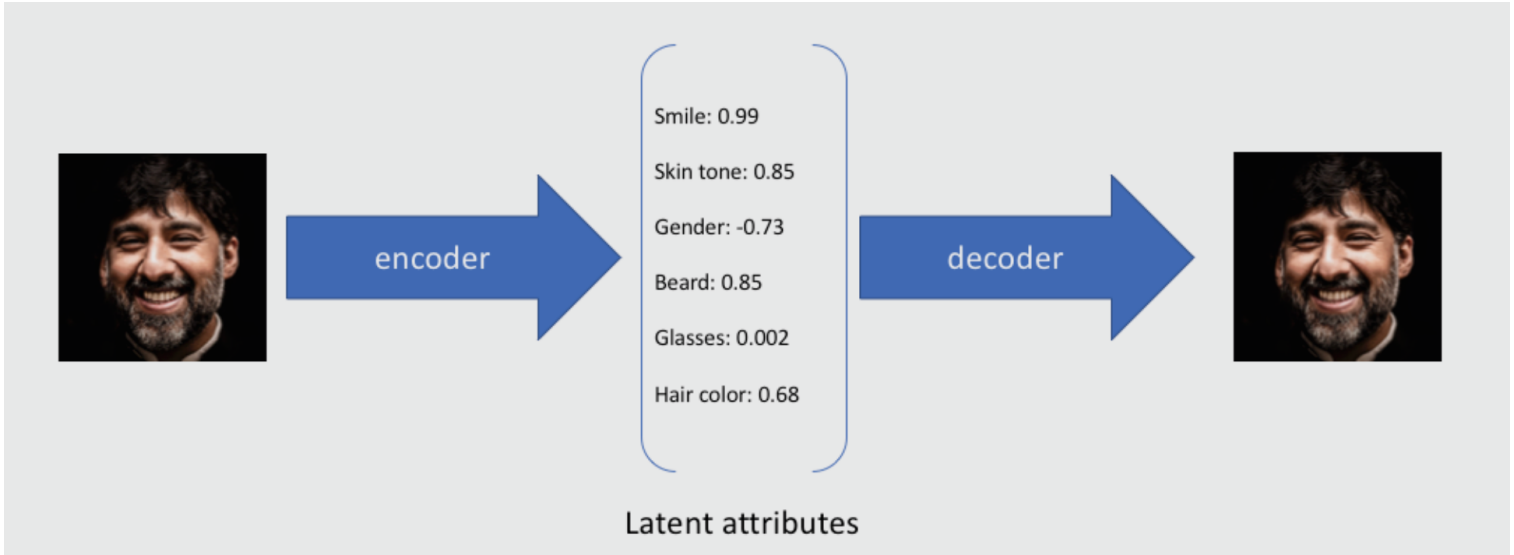
encoder
Smile: 0.99
Skin tone: 0.85
Gender: -0.73
Beard: 0.85
Glasses: 0.002
Hair color: 0.68
decoder
Latent attributes

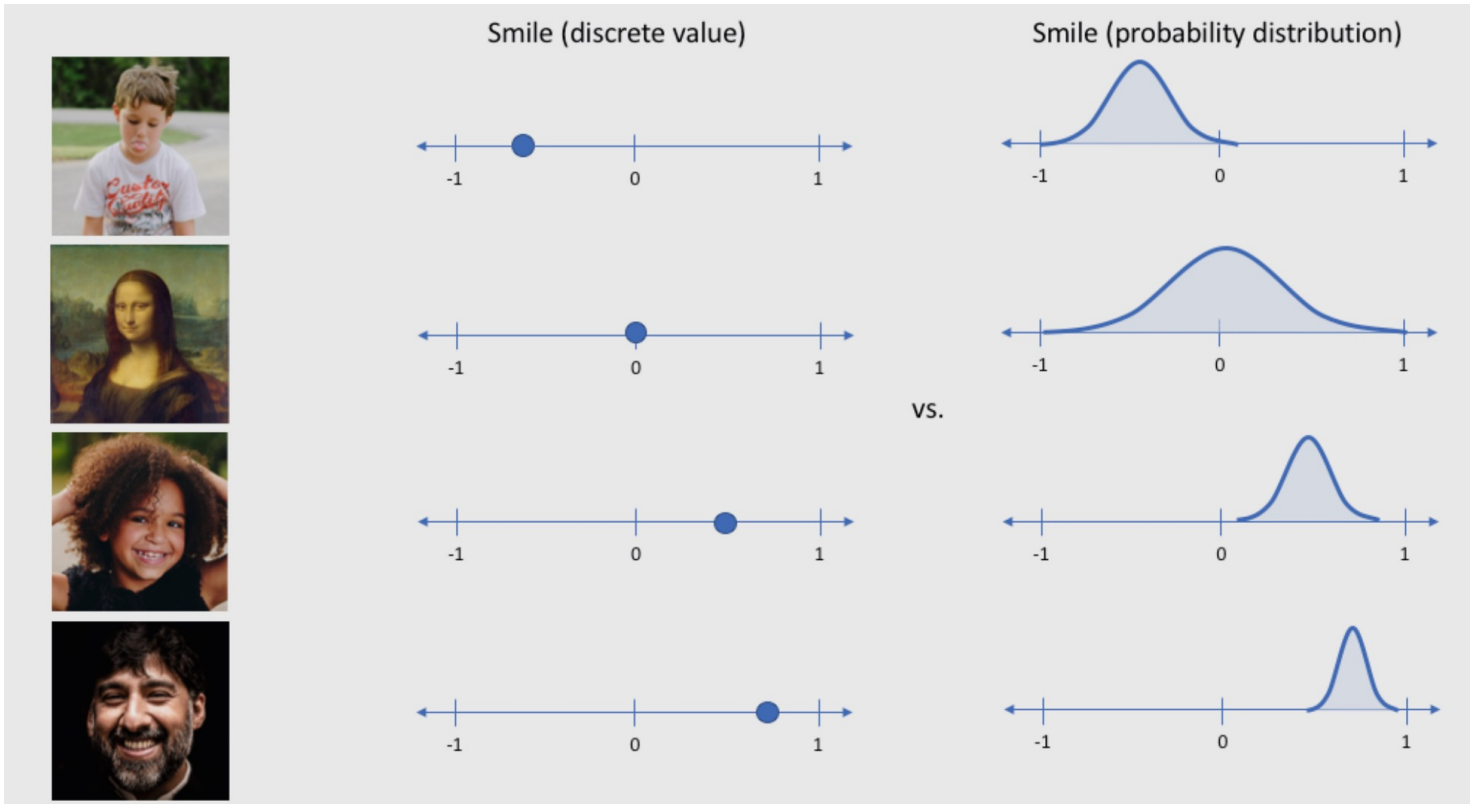
Smile (discrete value)
Smile (probability distribution)
-1
0
1
vs.

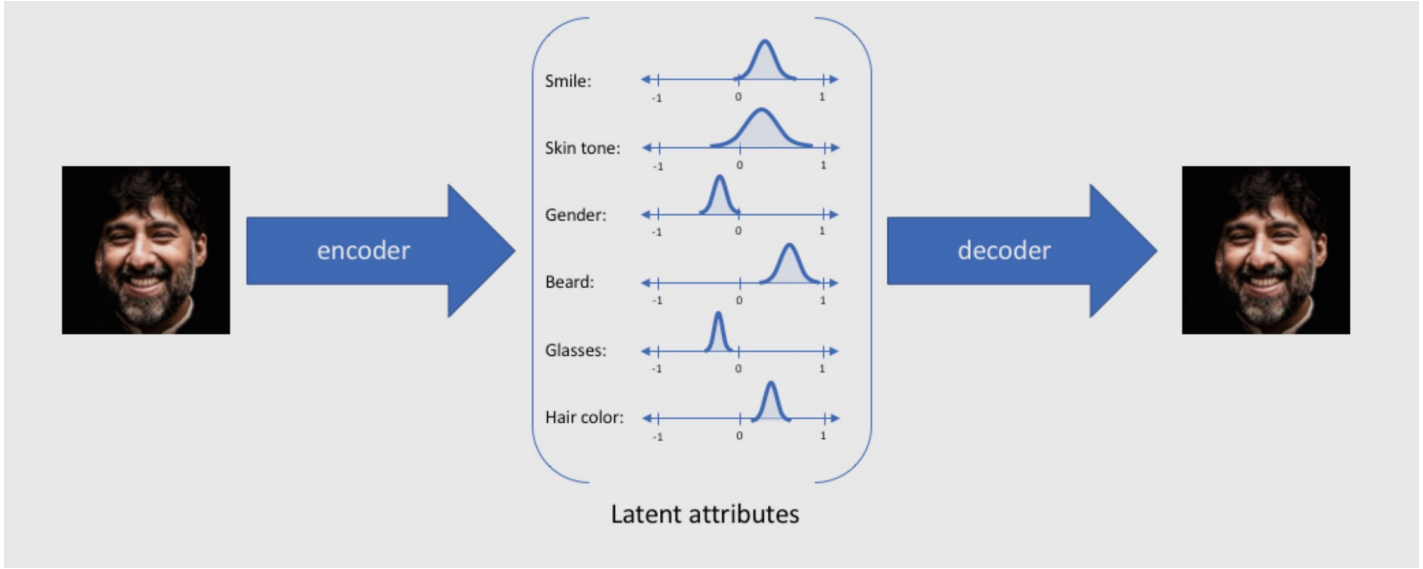
Smile:
Skin tone:
Gender:
Beard:
Glasses:
Hair color:
-1
0
1
encoder
decoder
Latent attributes

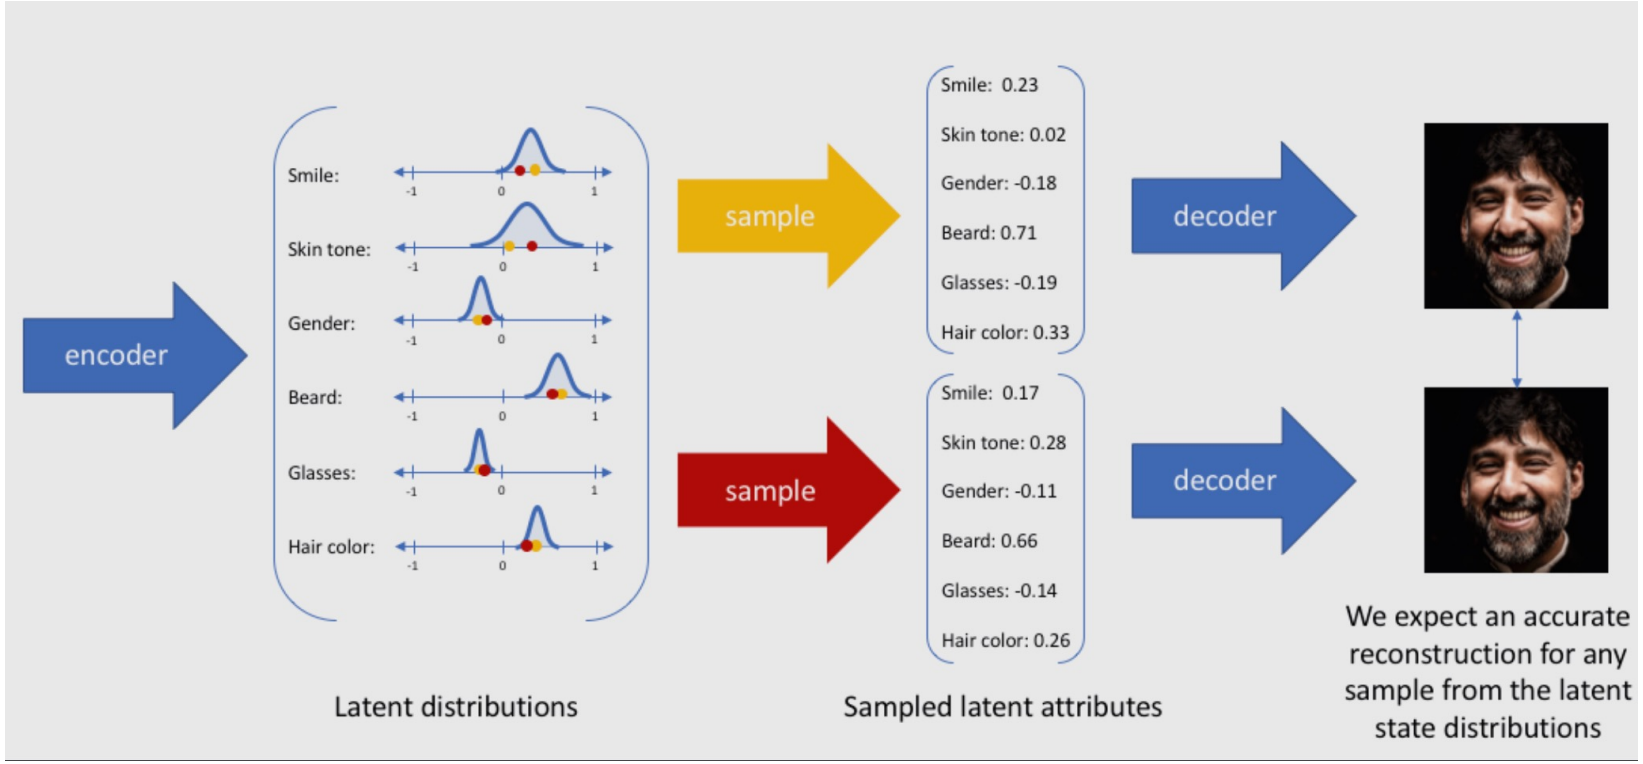
encoder
Smile:
Skin tone:
Gender:
Beard:
Glasses:
Hair color:
-1
0
1
sample
sample
Smile: 0.23
Skin tone: 0.02
Gender: -0.18
Beard: 0.71
Glasses: -0.19
Hair color: 0.33
Smile: 0.17
Skin tone: 0.28
Gender: -0.11
Beard: 0.66
Glasses: -0.14
Hair color: 0.26
decoder
decoder
We expect an accurate reconstruction for any sample from the latent state distributions
Latent distributions
Sampled latent attributes

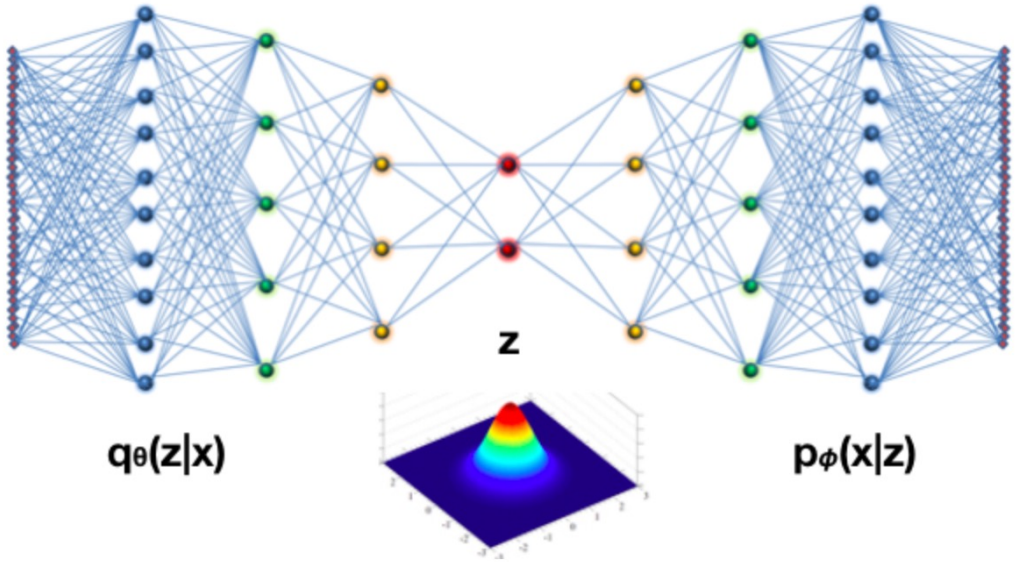

Figure 2: VAE

$$\min KL\left(q\left(z|x\right)||p\left(z|x\right)\right)$$

$$E_{q(z|x)}\log p\left(x|z\right) - KL\left(q\left(z|x\right)||p\left(z\right)\right)$$

The KL divergence between the approximate and the real posterior distributions is given by,

$$D_{KL}\left(q_\theta(z|x_i)||p(z|x_i)\right) = -\int q_\theta(z|x_i)\log\left(\frac{p(z|x_i)}{q_\theta(z|x_i)}\right)dz \geq 0$$

https://medium.com/dataseries/variational-autoencoder-with-pytorch-2d359cbf027b

Variational Autoencoder - dogs generation

https://www.kaggle.com/code/speedwagon/variational-autoencoder-dogs-generation

Mid term ?

Generate music with Variational AutoEncoder

https://www.kaggle.com/code/basu369victor/generate-music-with-variational-autoencoder